

END 2 END DATA SCIENCE

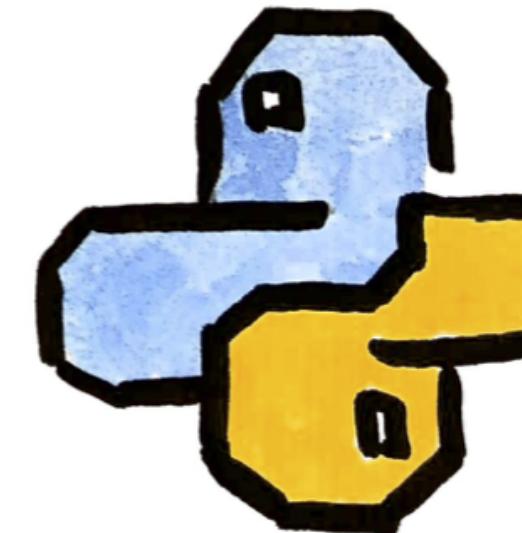
Mark Keinhörster &
Dr. Shirin Glander

Convolutional Neural Nets



- image classification
class = tree
- object detection
flower

Computer
Vision



About this Workshop

- Learn what neural nets are and what the big deal is with deep learning
- Interactively build a model that differentiates between fruits on images
- Get a glimpse of production-readiness
- Learn about Luigi pipelines and their main components
- Write your production ready pipeline
- Get an overview of luigis modules
- Get an overview about DVC and how to implement pipelines
- Use TensorFlow Serving to deploy your model

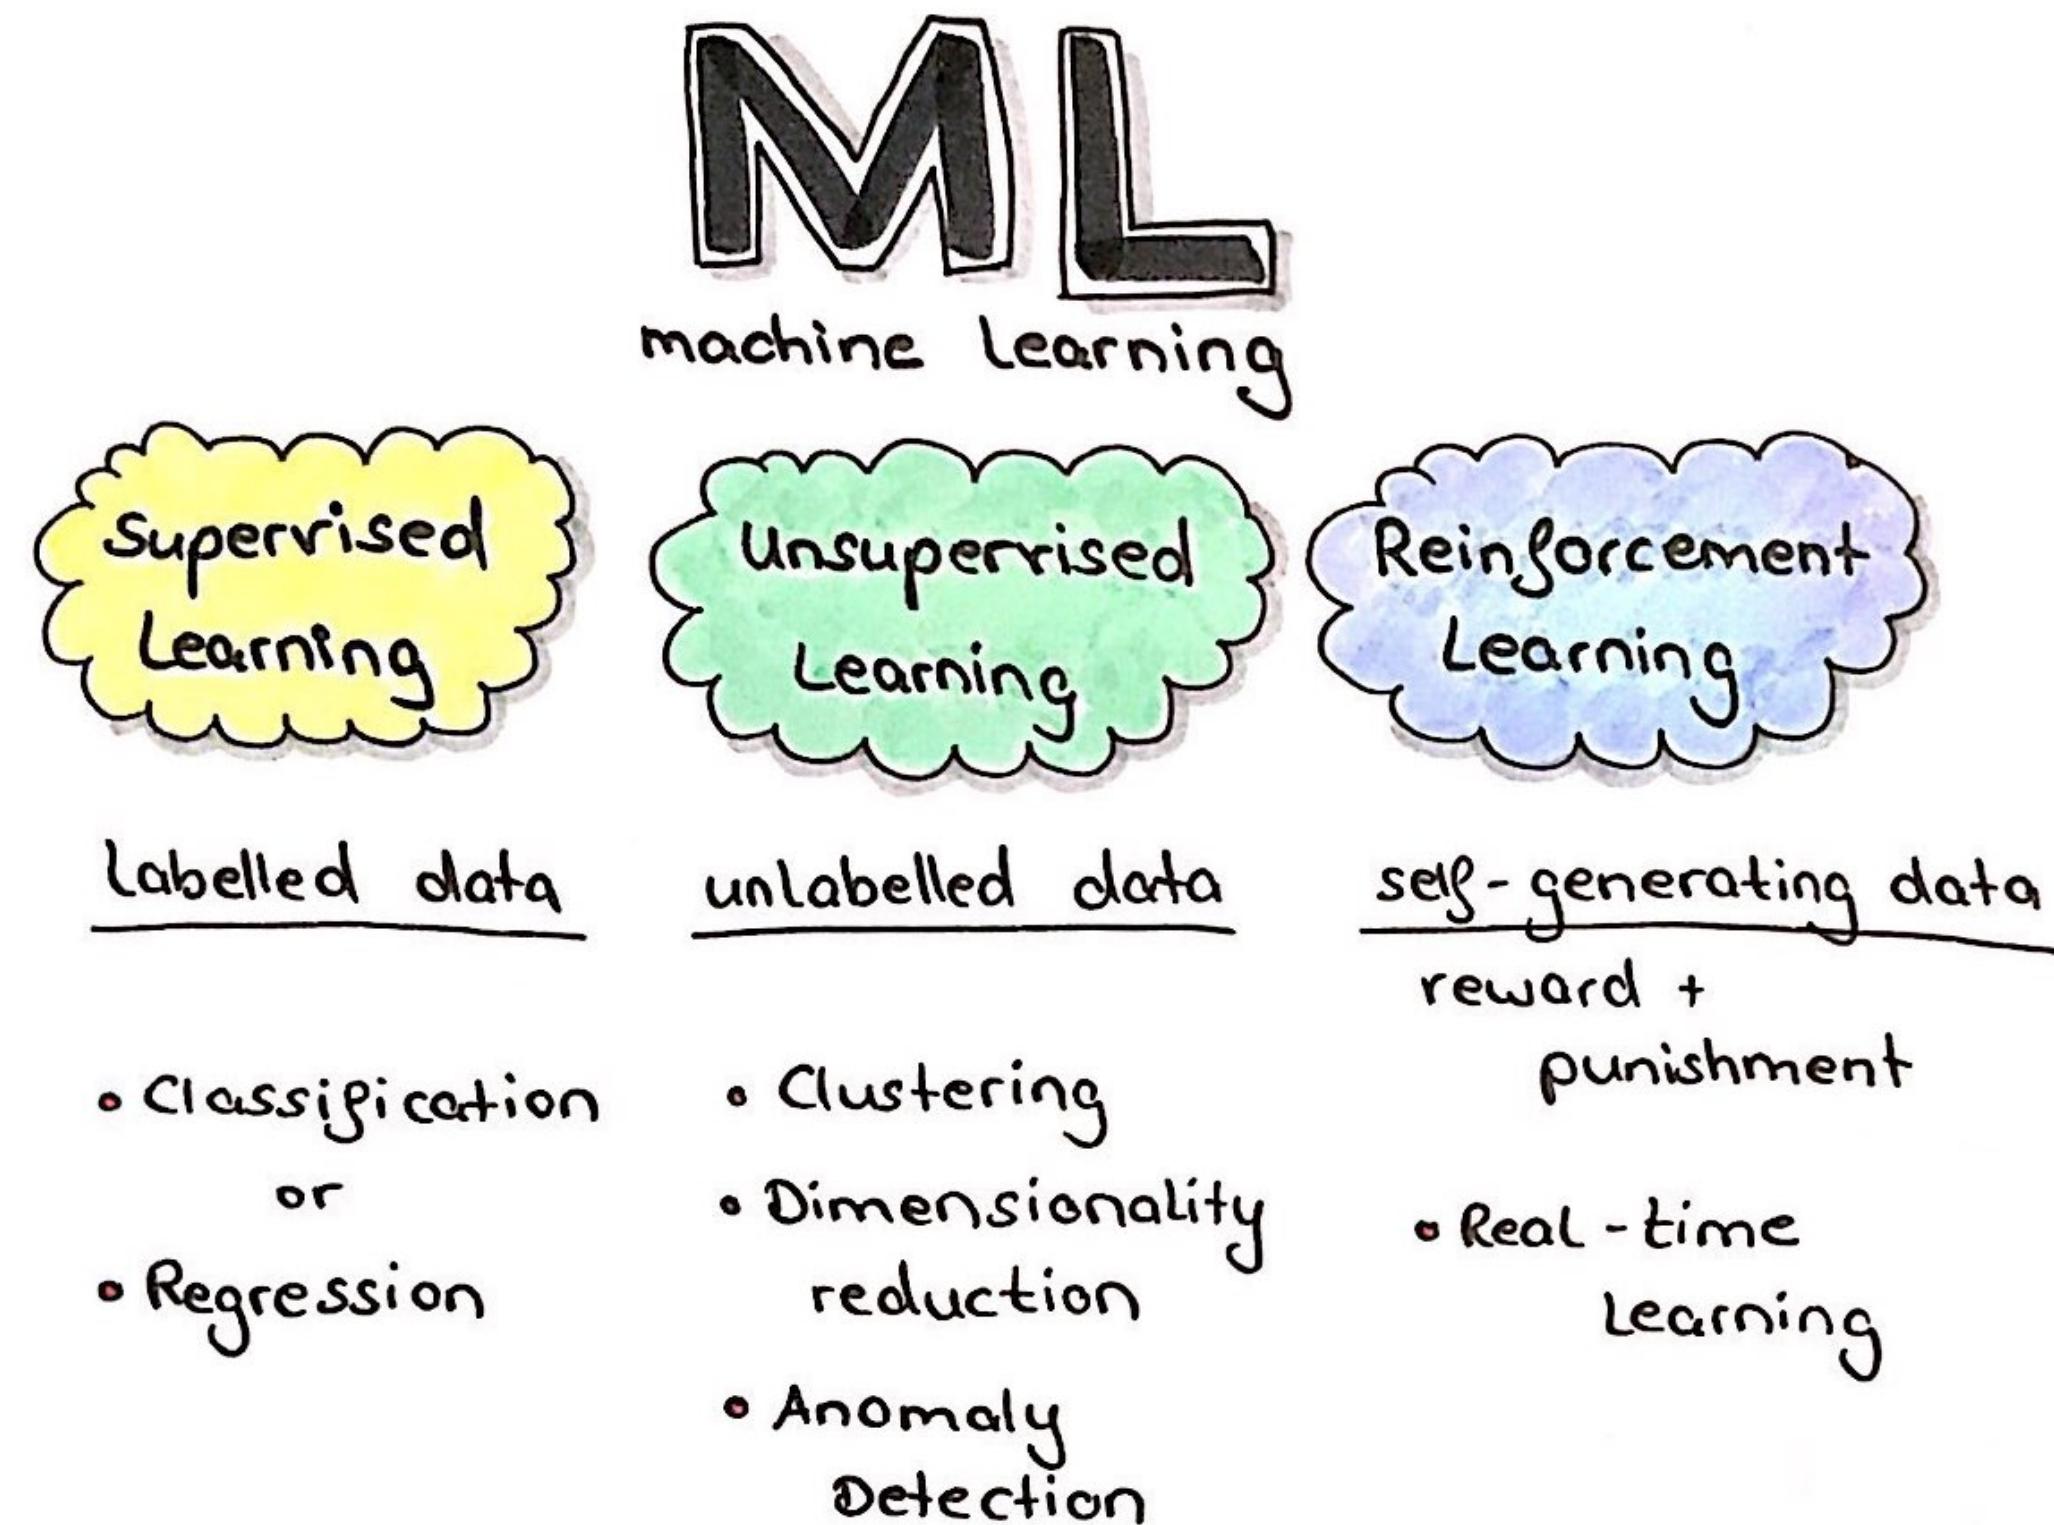
About us

MAC_HI_NE
LEARNI**G**

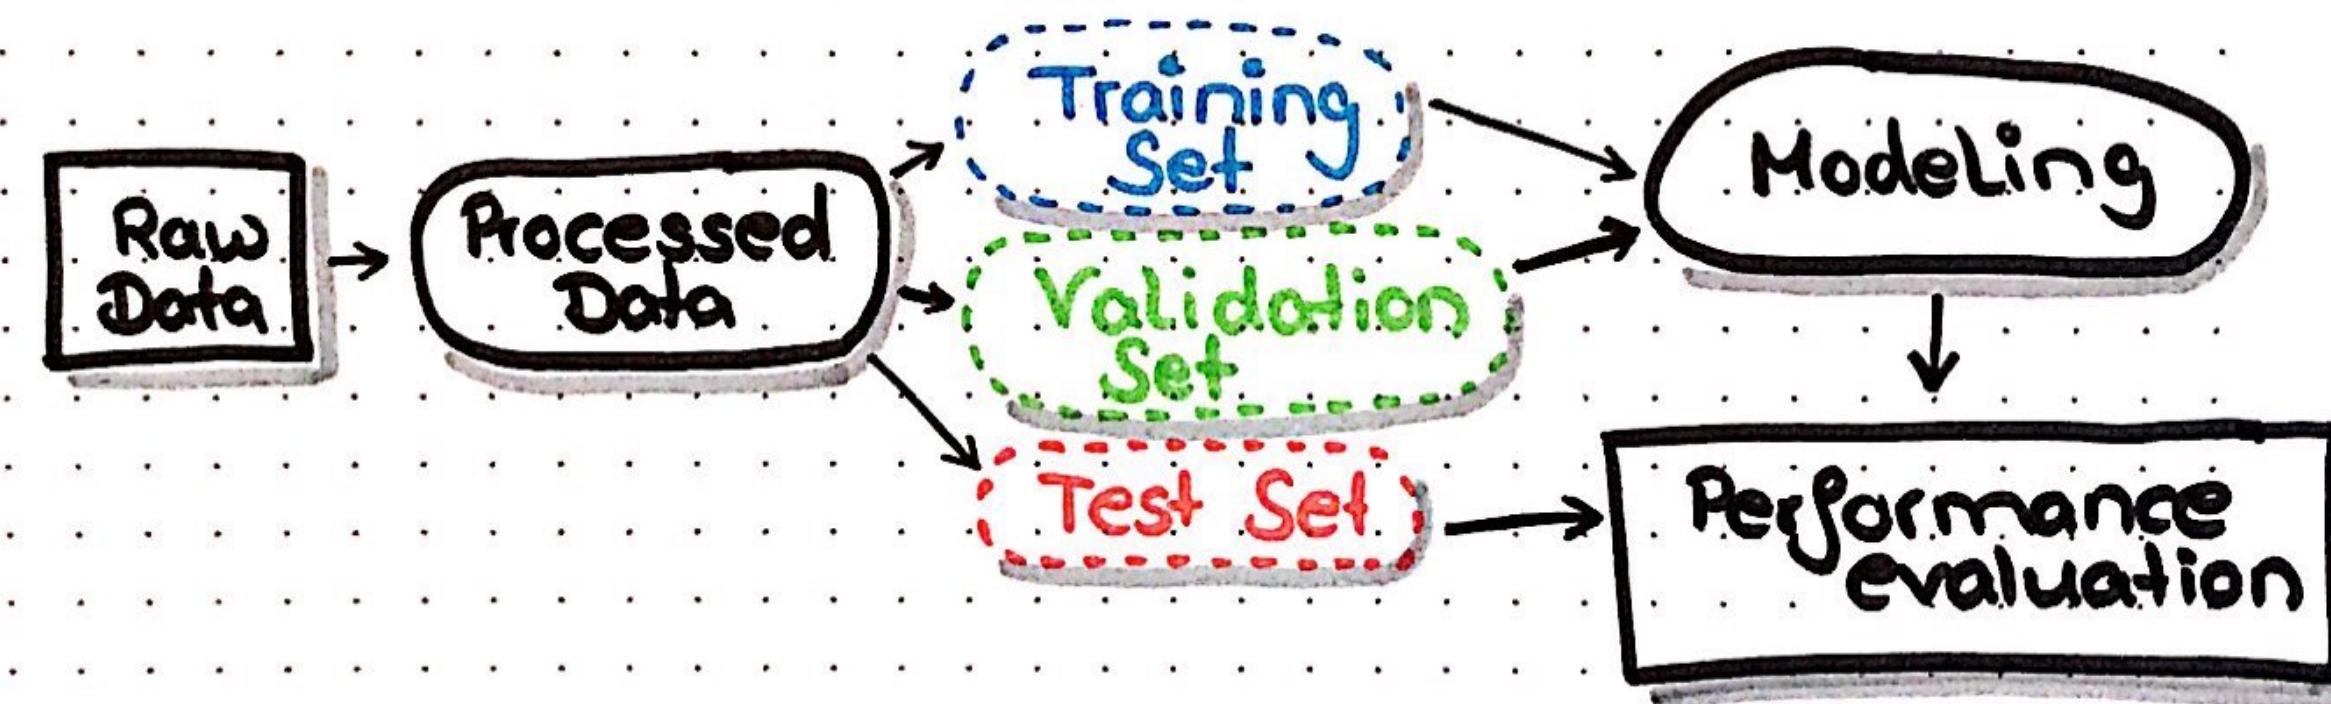
The image features the words "MA"CHINE" and "L"EARNG" in a bold, sans-serif font. The letters are filled with a detailed green circuit board pattern, complete with gold-colored tracks, black dots, and various electronic components. The letters are arranged in two rows: "MA"CHINE" on top and "L"EARNG" below it. The background is plain white.

<https://pixabay.com/de/a-ich-ai-anatomie-2729781/>

What is Machine Learning (ML)



A typical ML Workflow



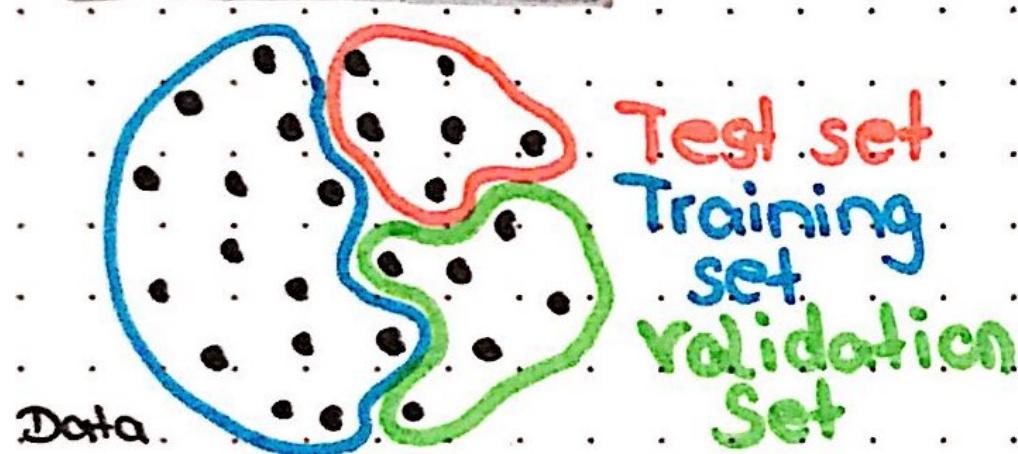
Why use validation and test data?



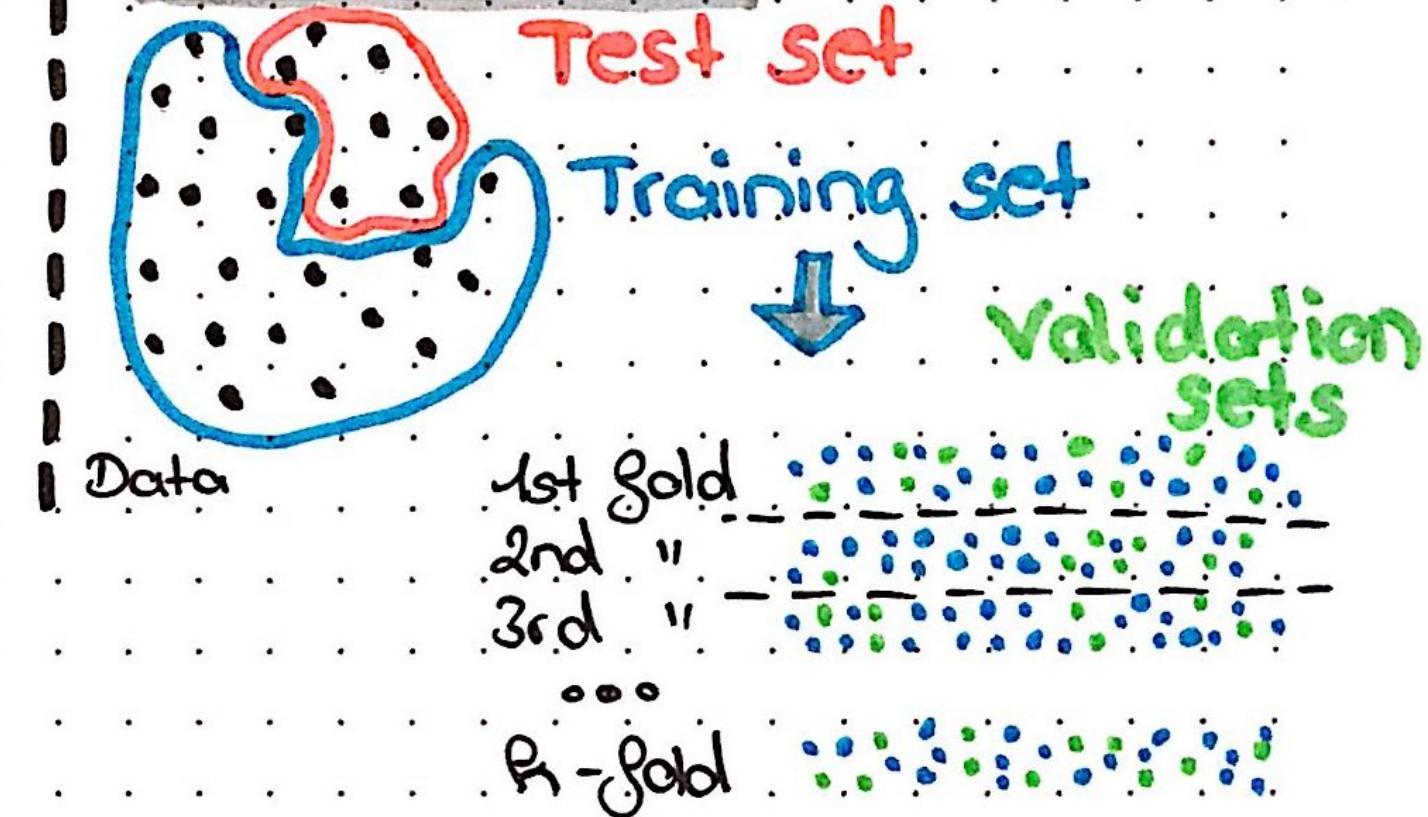
**THE BEST WAY TO
EXPLAIN OVERFITTING**

Validation methods

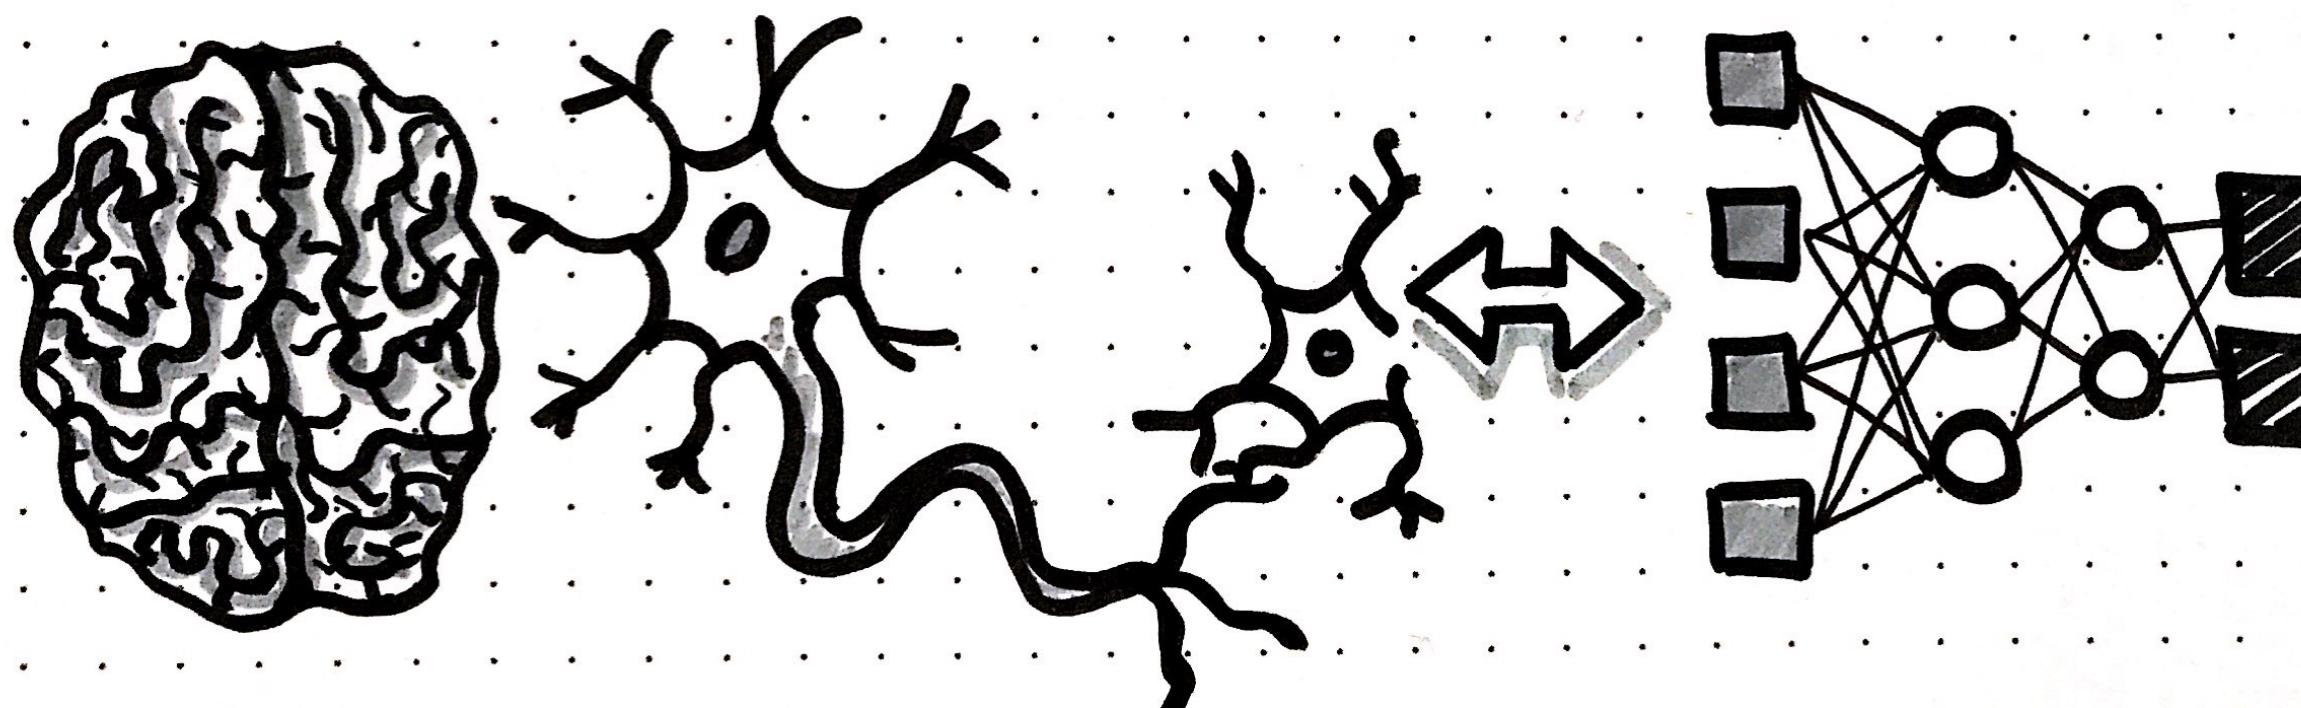
Hold-out validation



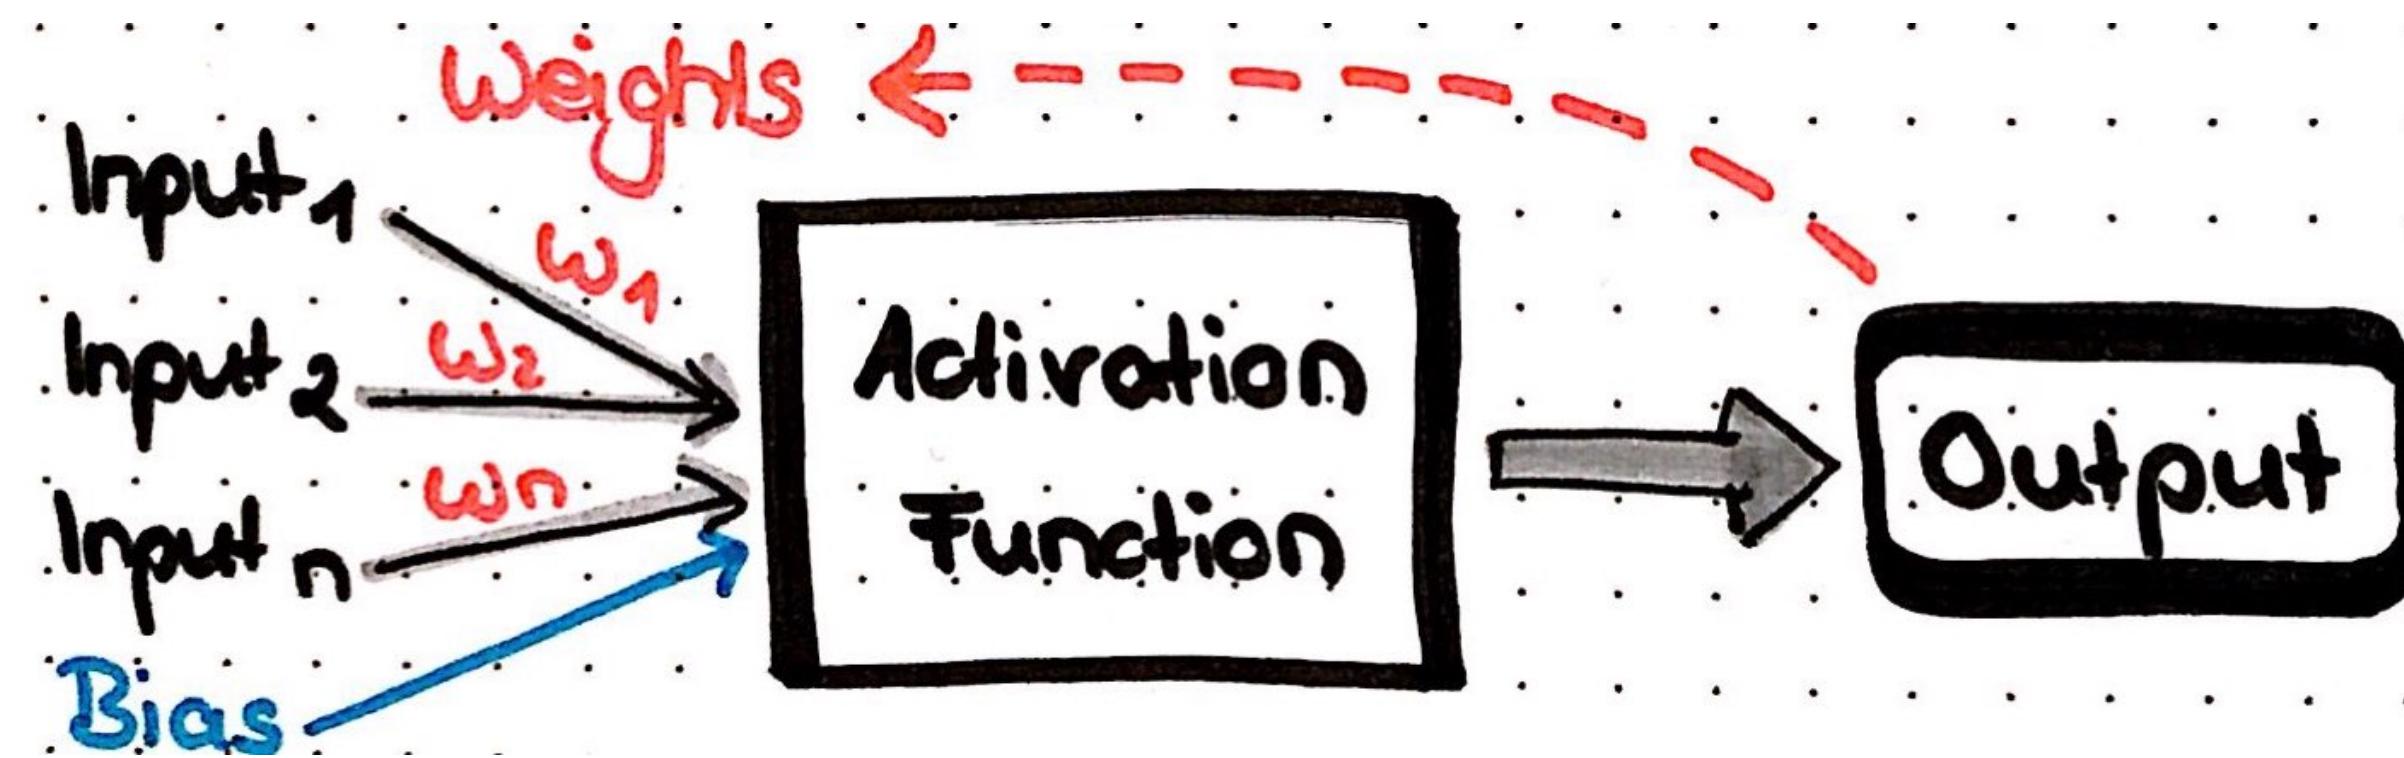
Cross-validation



What are neural nets?

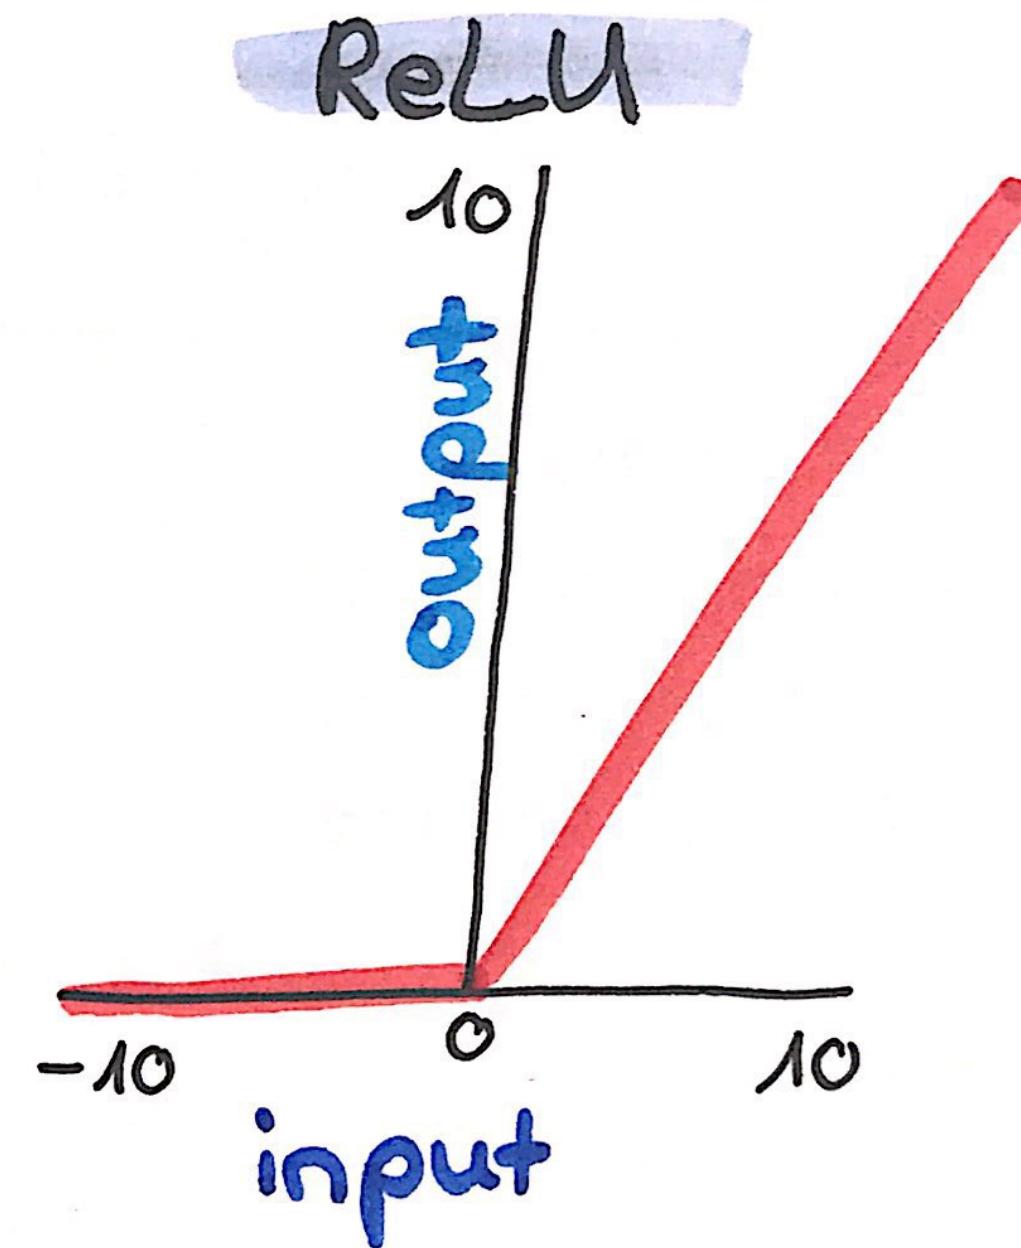
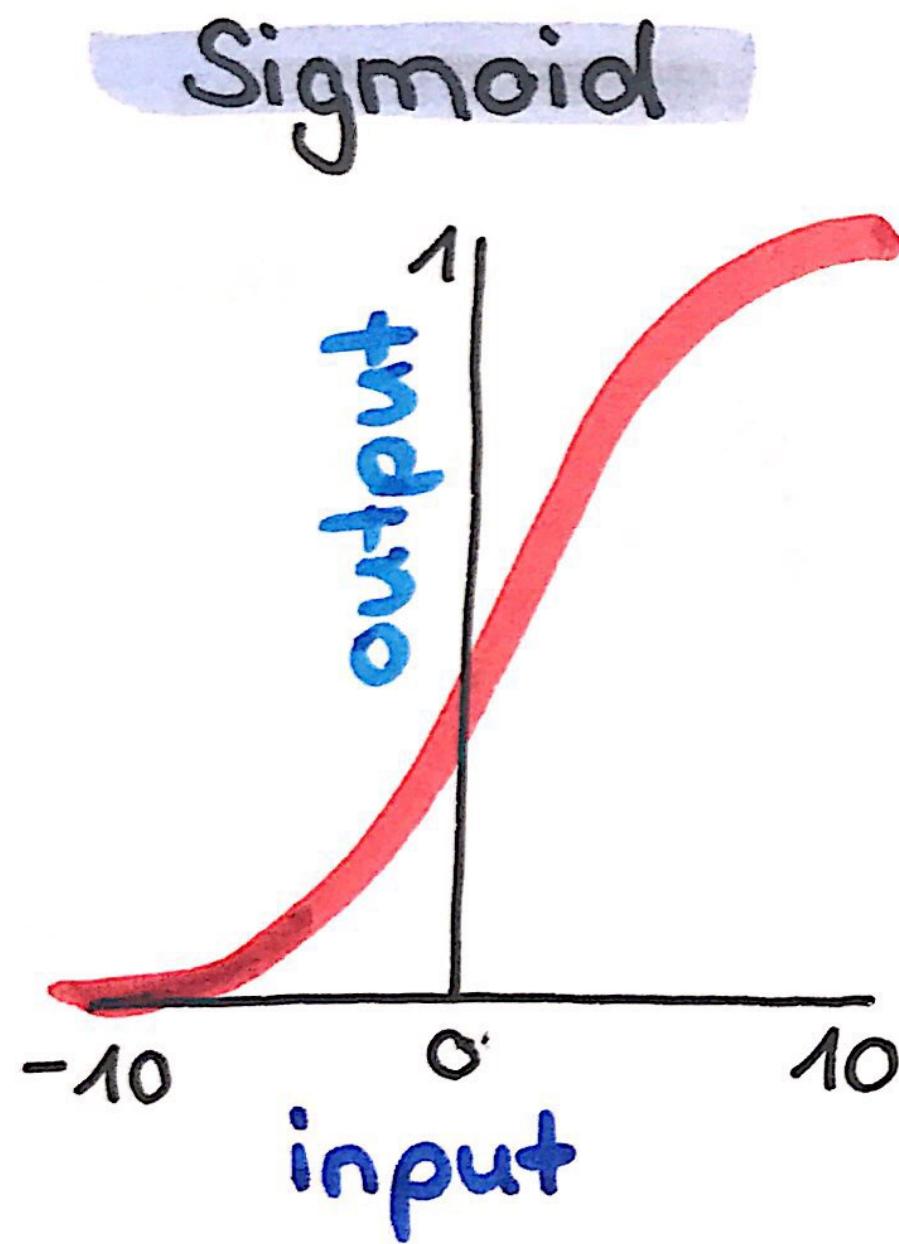


Perceptrons



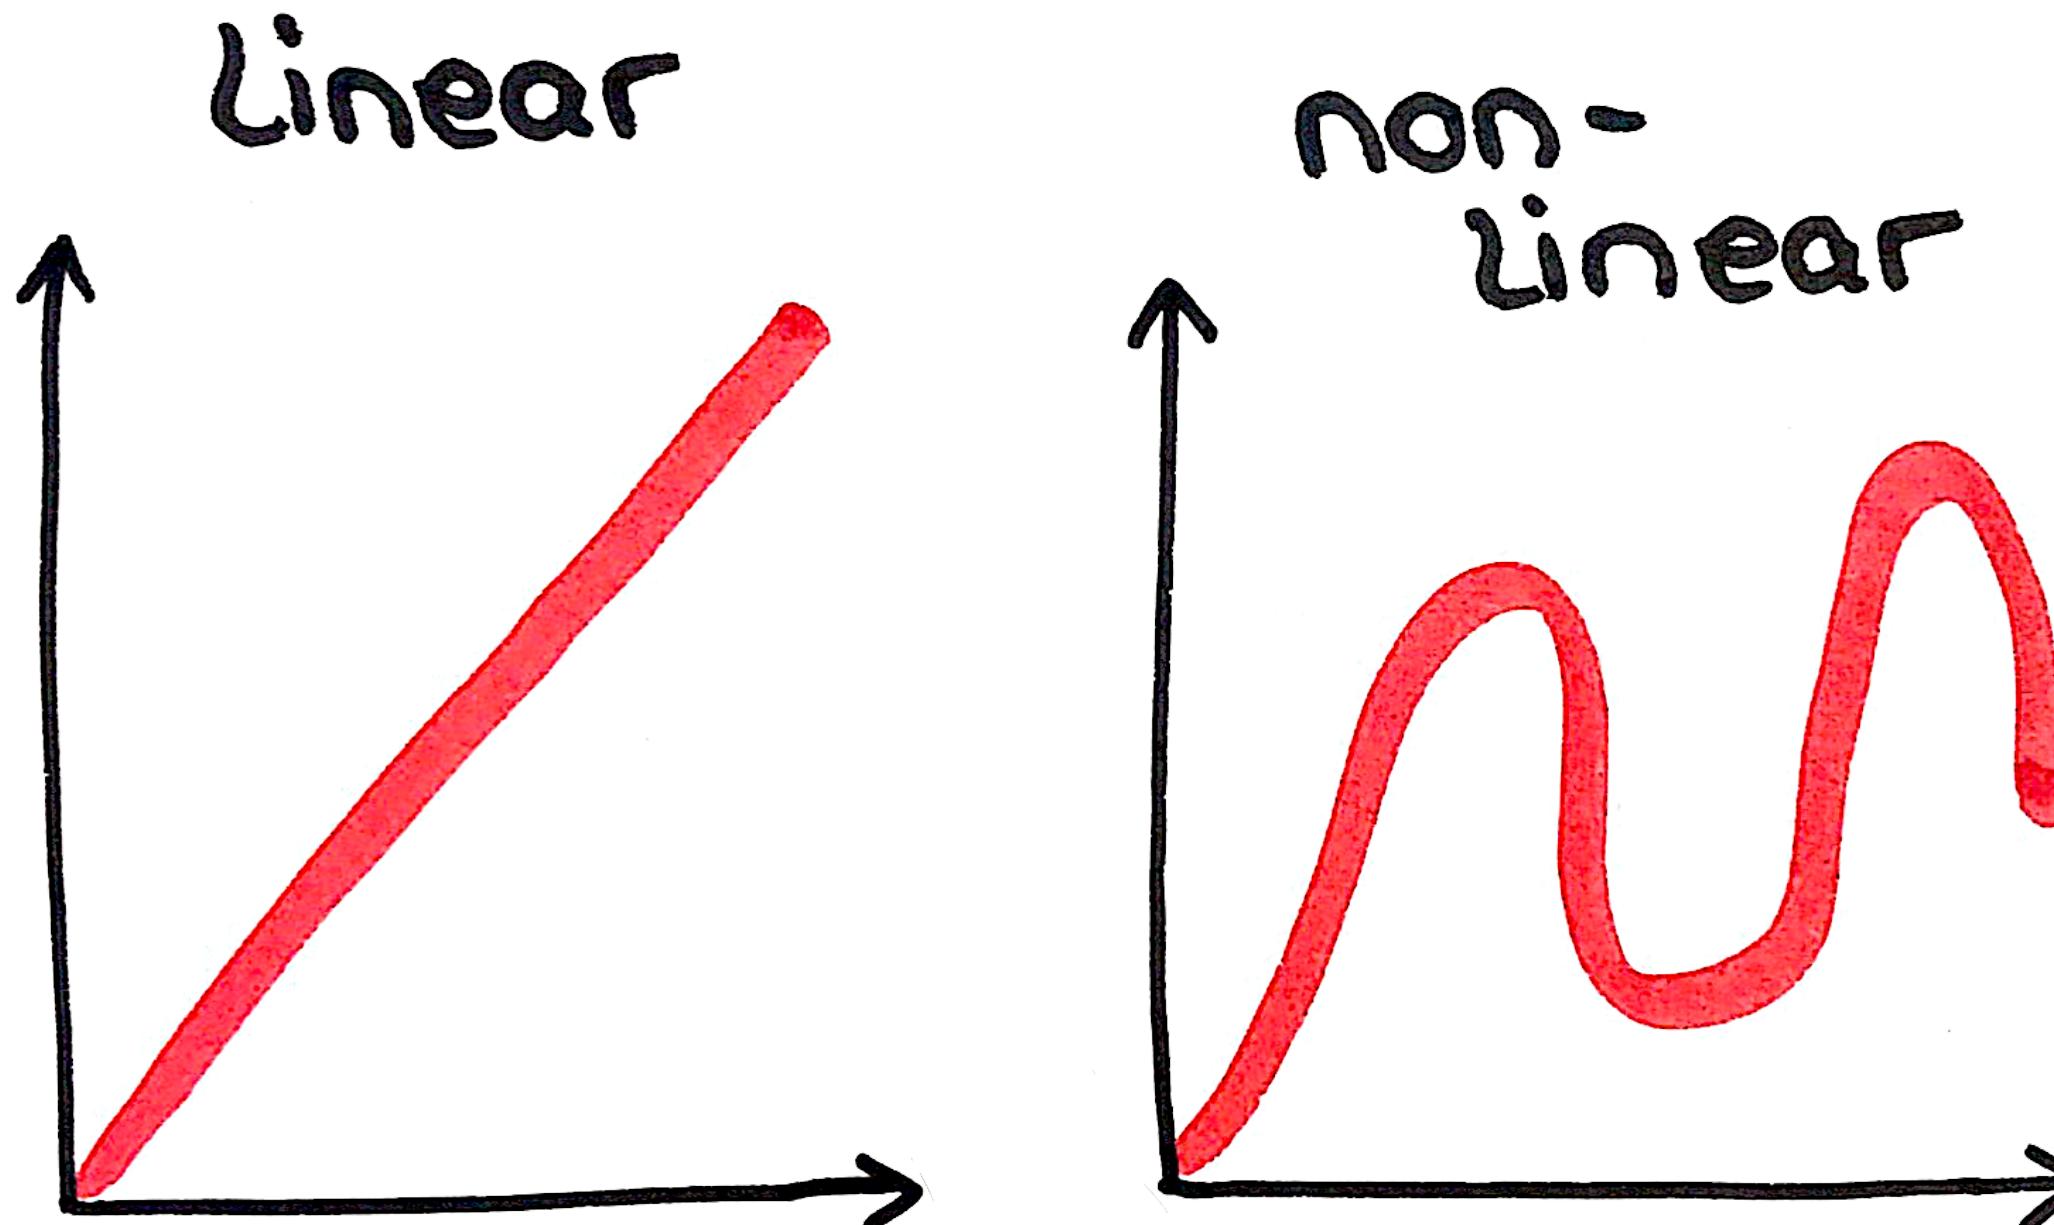
Activation functions normalise input

- every problem can be described with a mathematical function

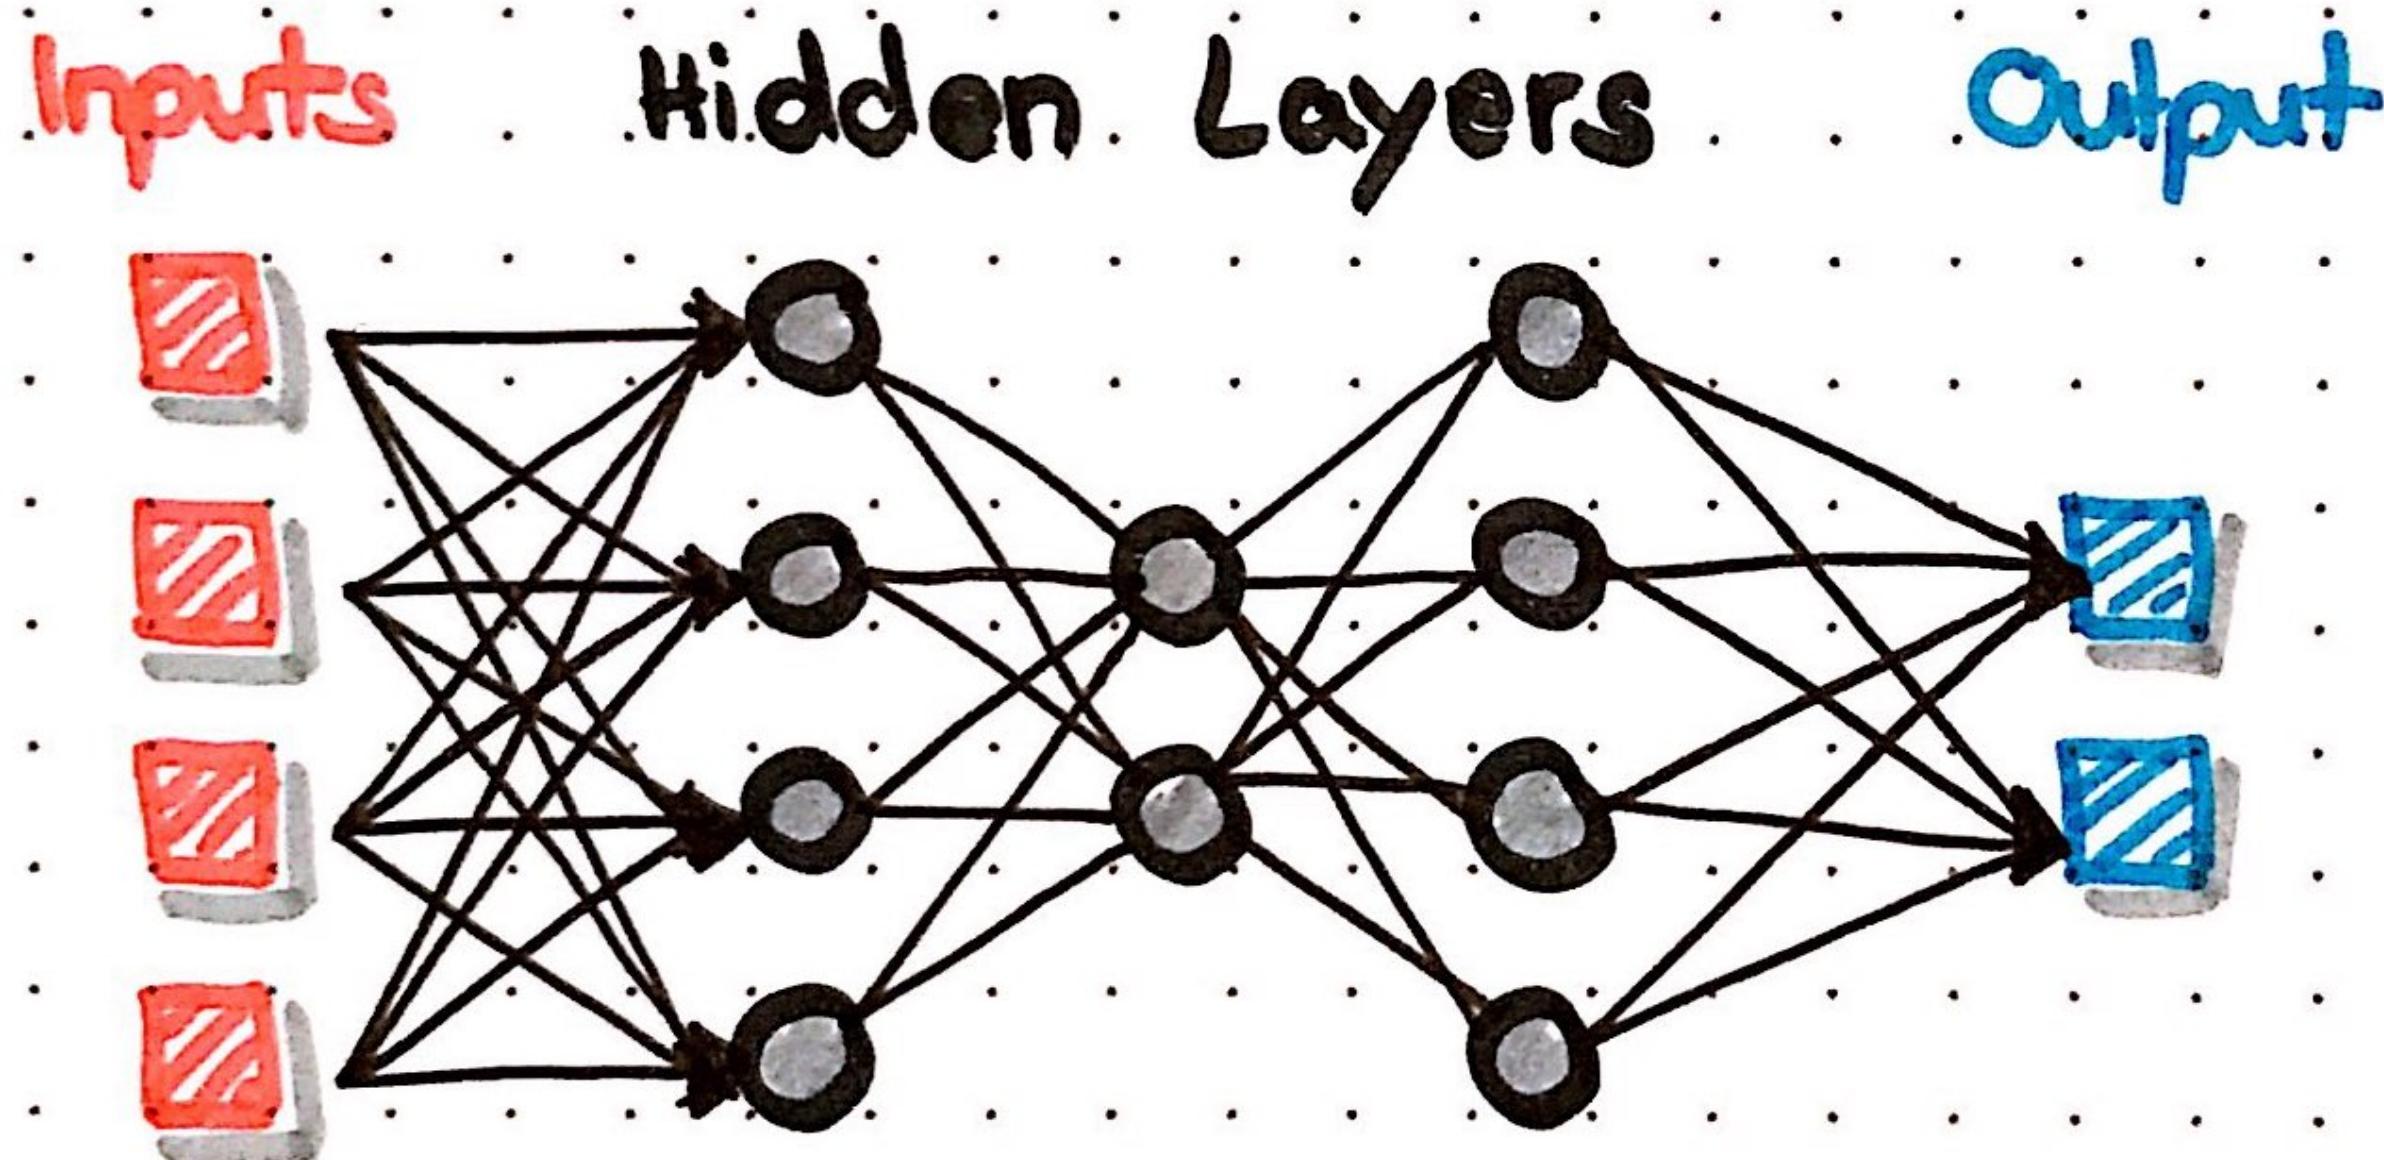


Activation functions make non-linearity possible

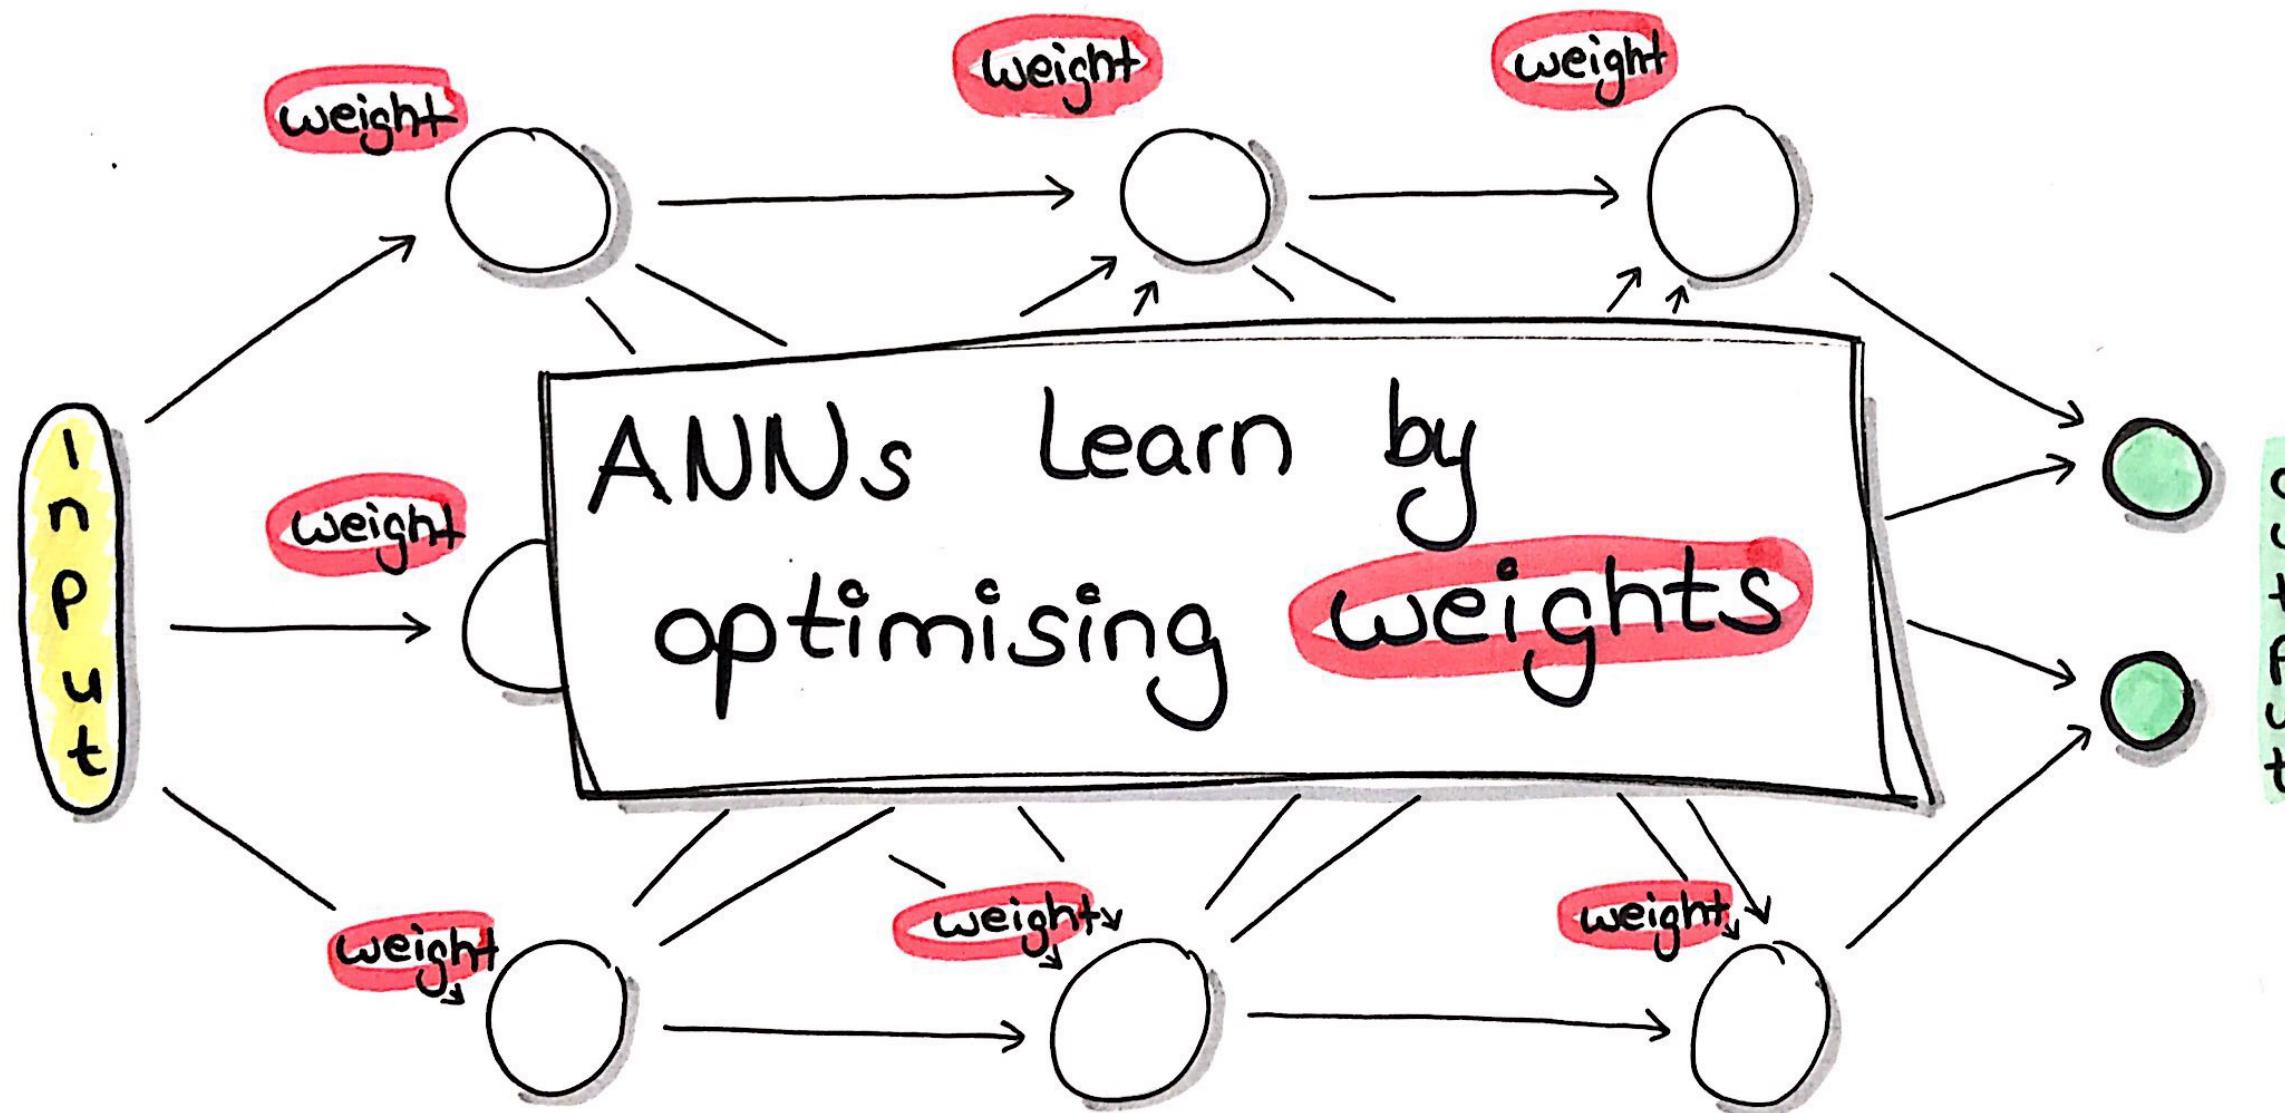
- non-linear activation functions allow us to approximate ANY mathematical formula with neural nets



Multi-Layer Perceptrons



How does a neural net learn?



How does a neural net learn?

The Softmax function

$$x * \omega + b = y$$

input weight bias Output

Model Training ~ finding good weights & biases

	score	probability	
class a	2	→ 0.7	correct
class b	1	→ 0.2	
class c	0.1	→ 0.1	

soft - max

How does a neural net learn?

Cross-entropy



probability distribution

one-hot encoded labels

class a
class b
class c

0.7
0.2
0.1

distance
 $\leftarrow \rightarrow$
cross-entropy

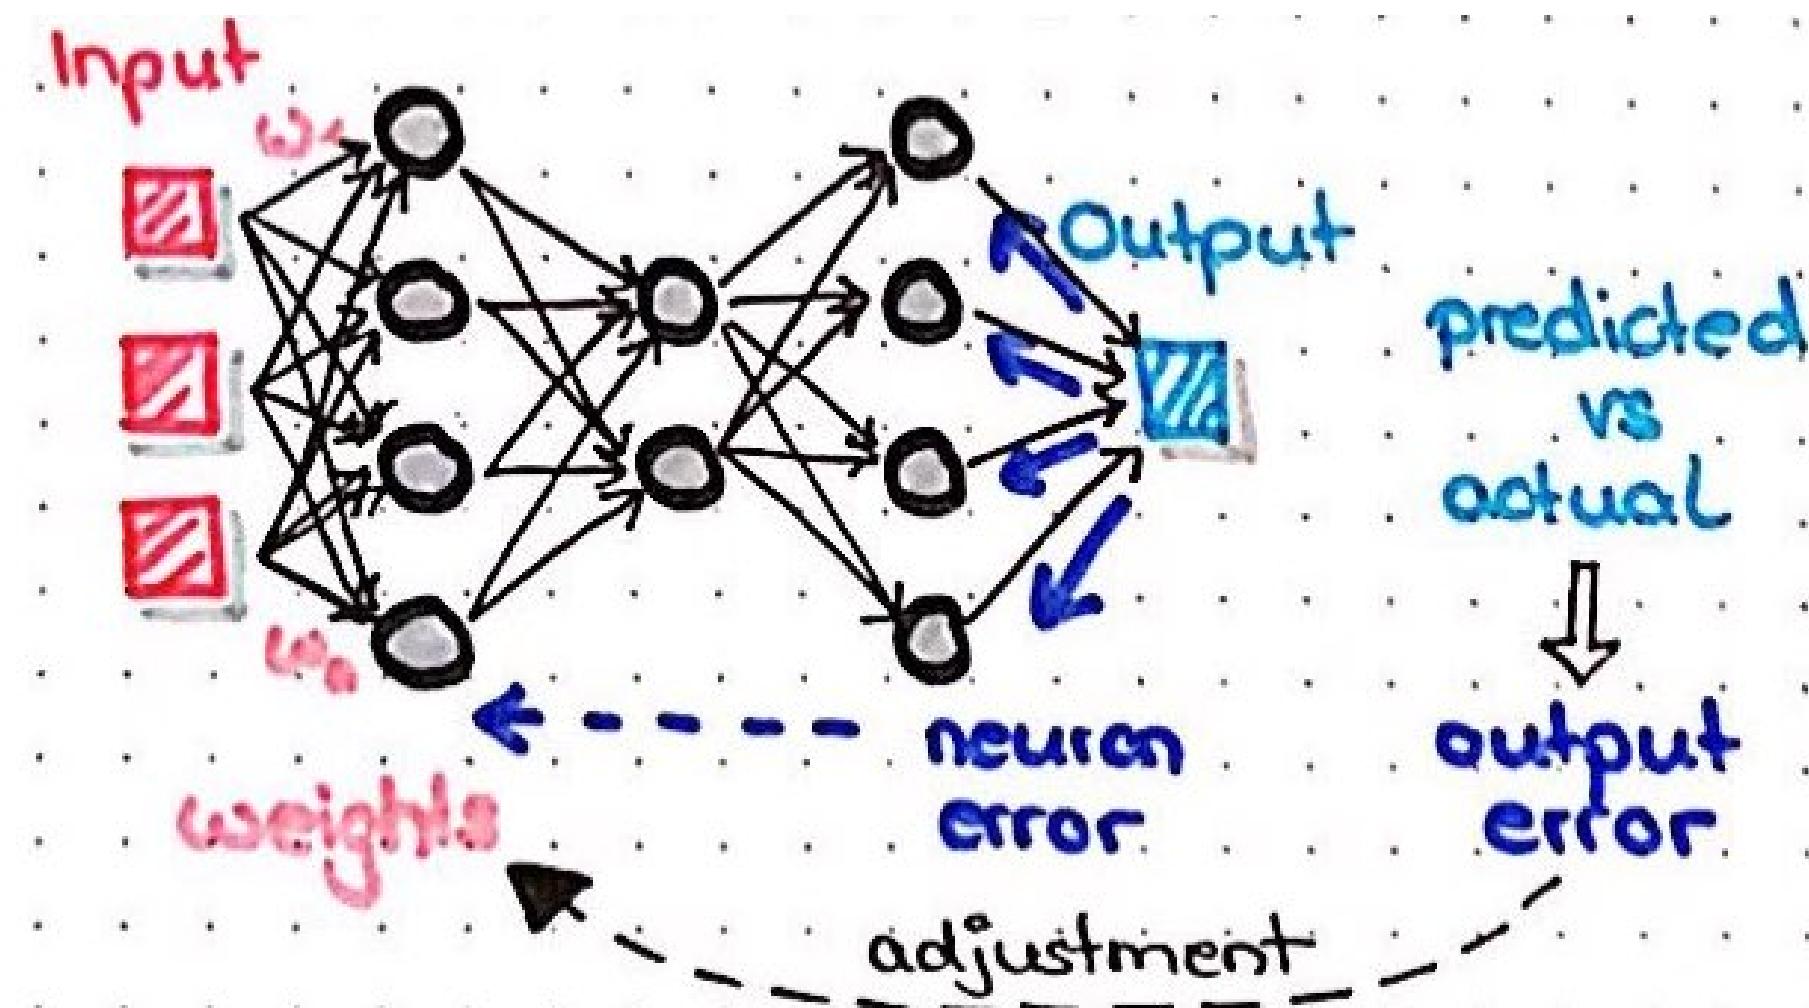
Model Training:

find weights & biases that minimize cross-entropy

Loss: average cross-entropy over training set.

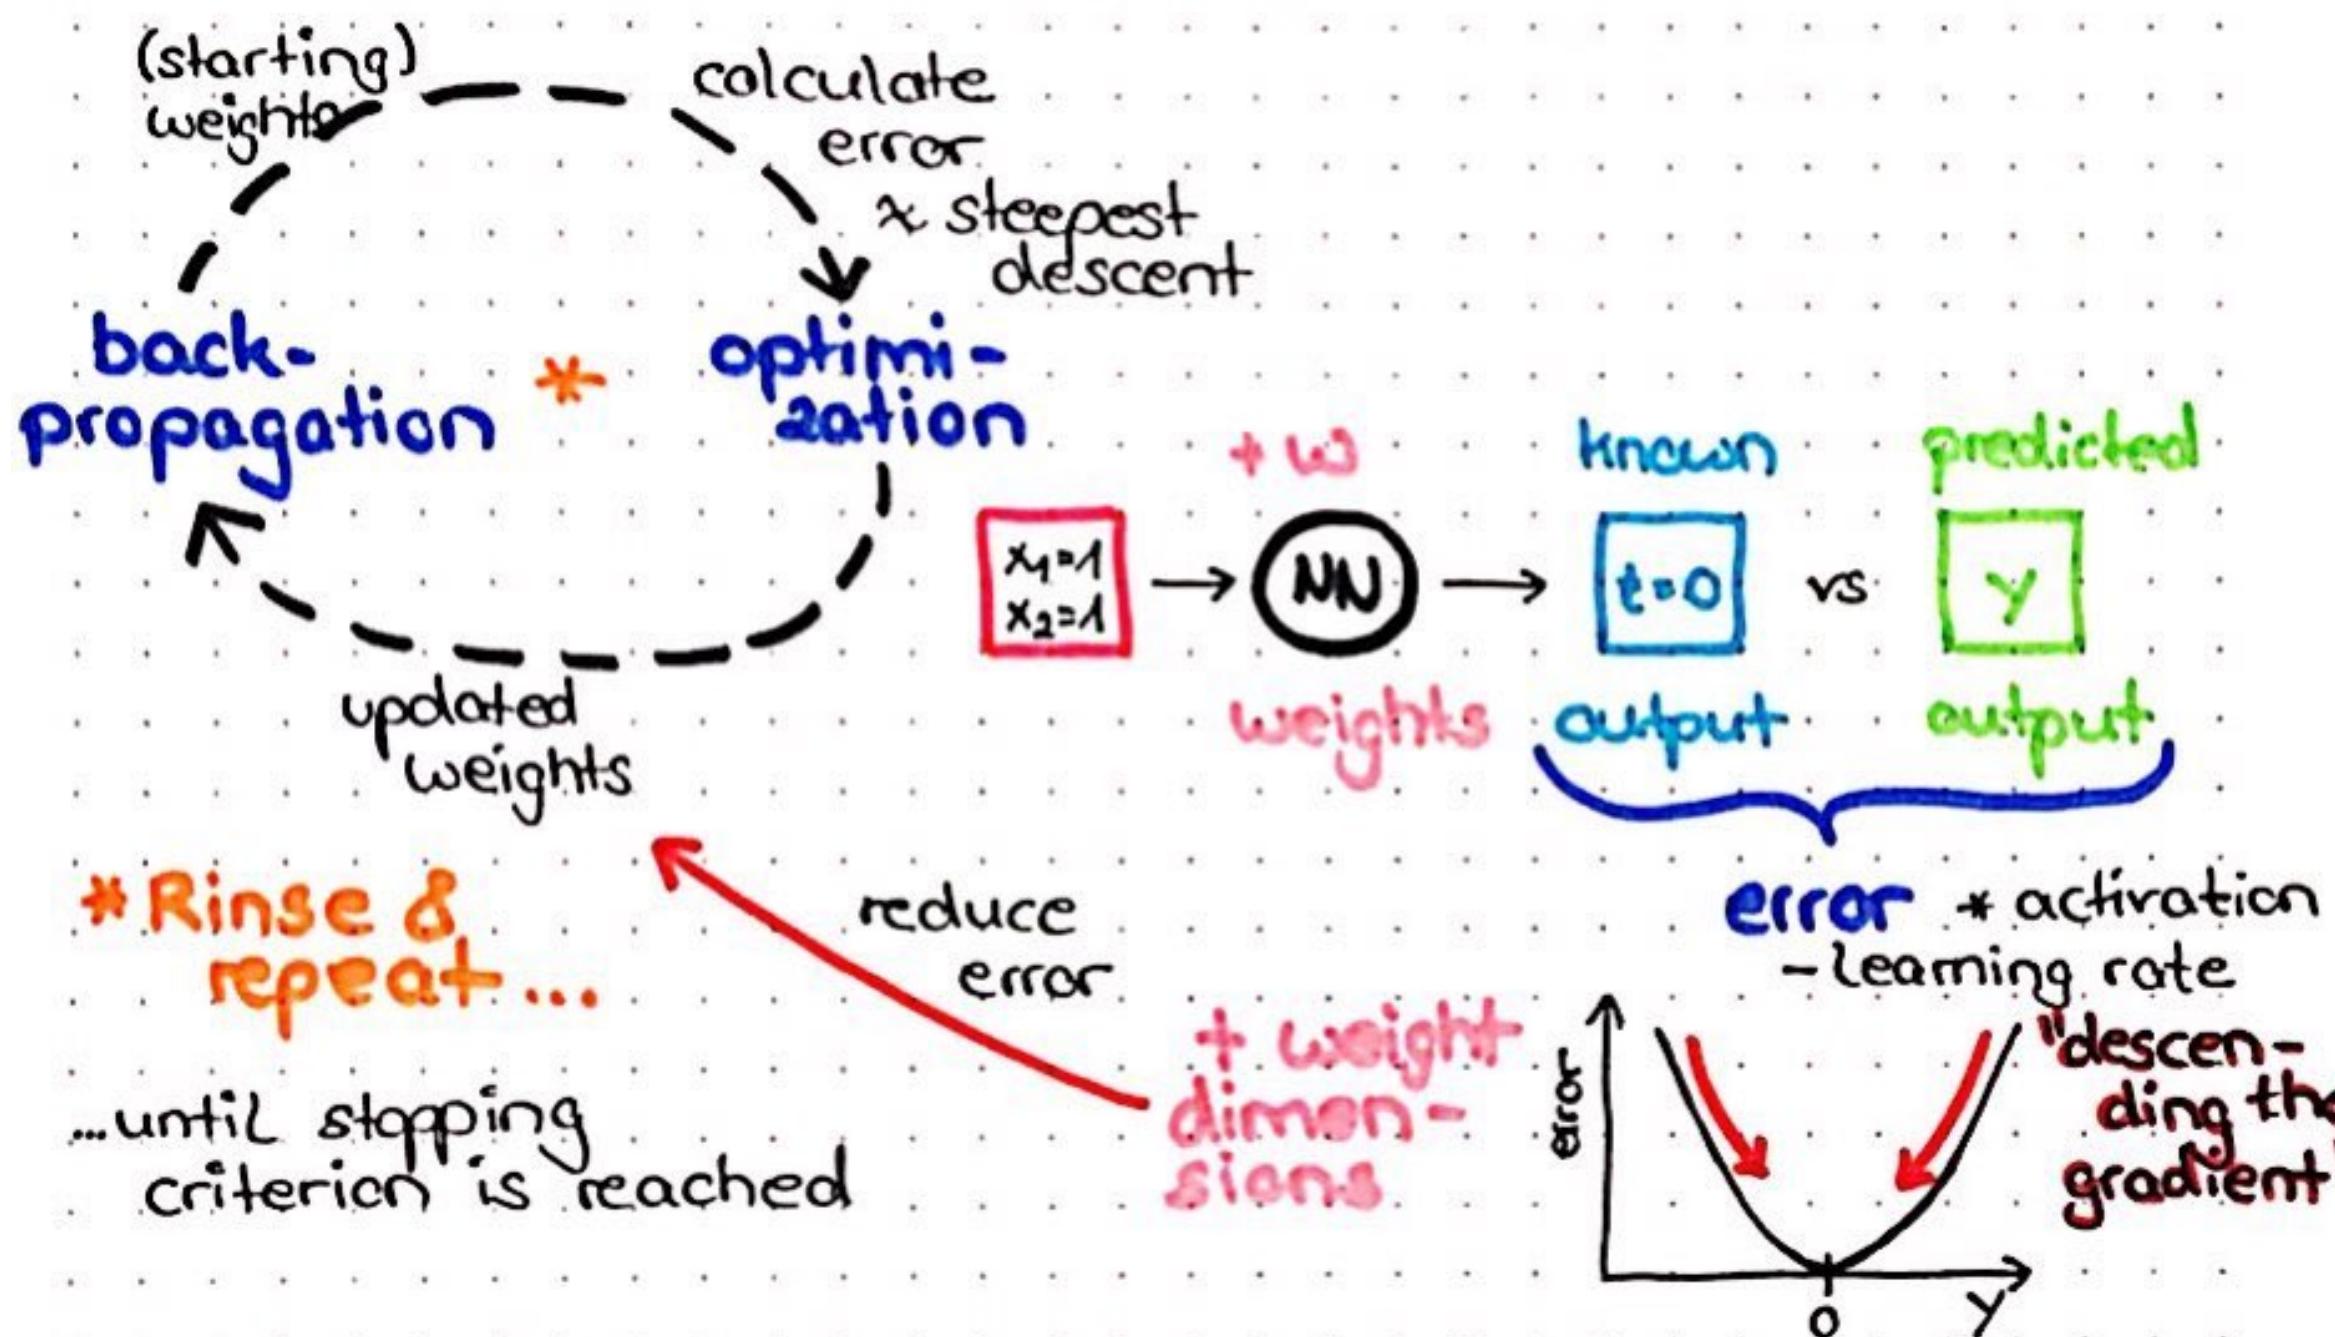
How does a neural net learn?

Backpropagation

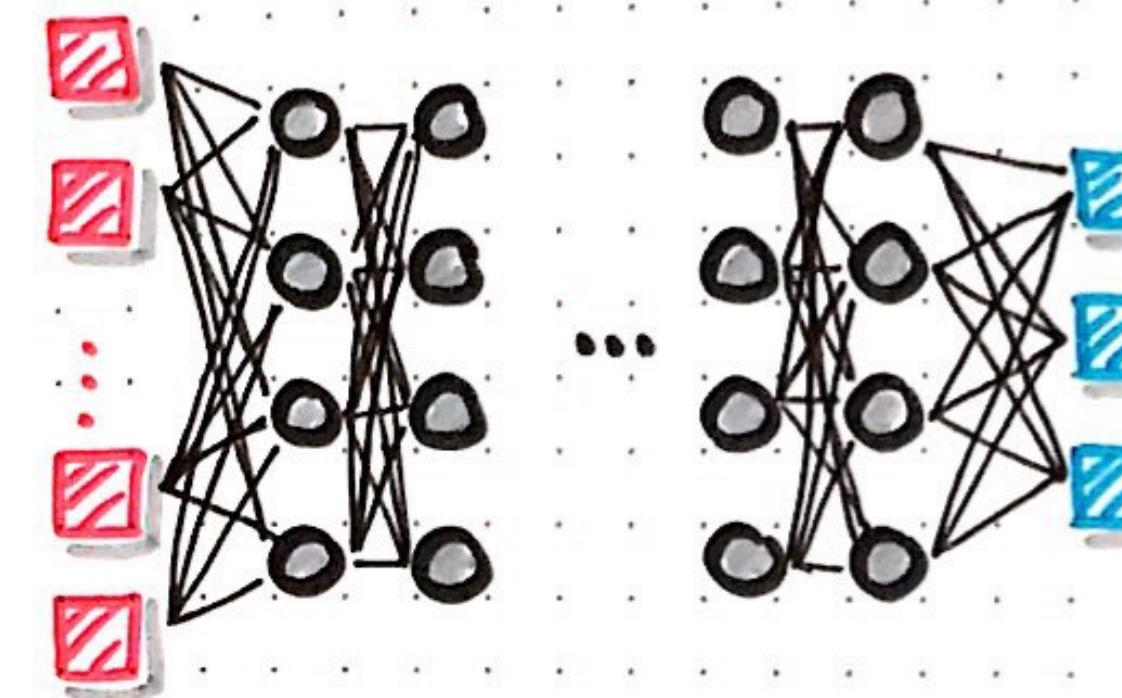


How does a neural net learn?

Gradient descent optimization

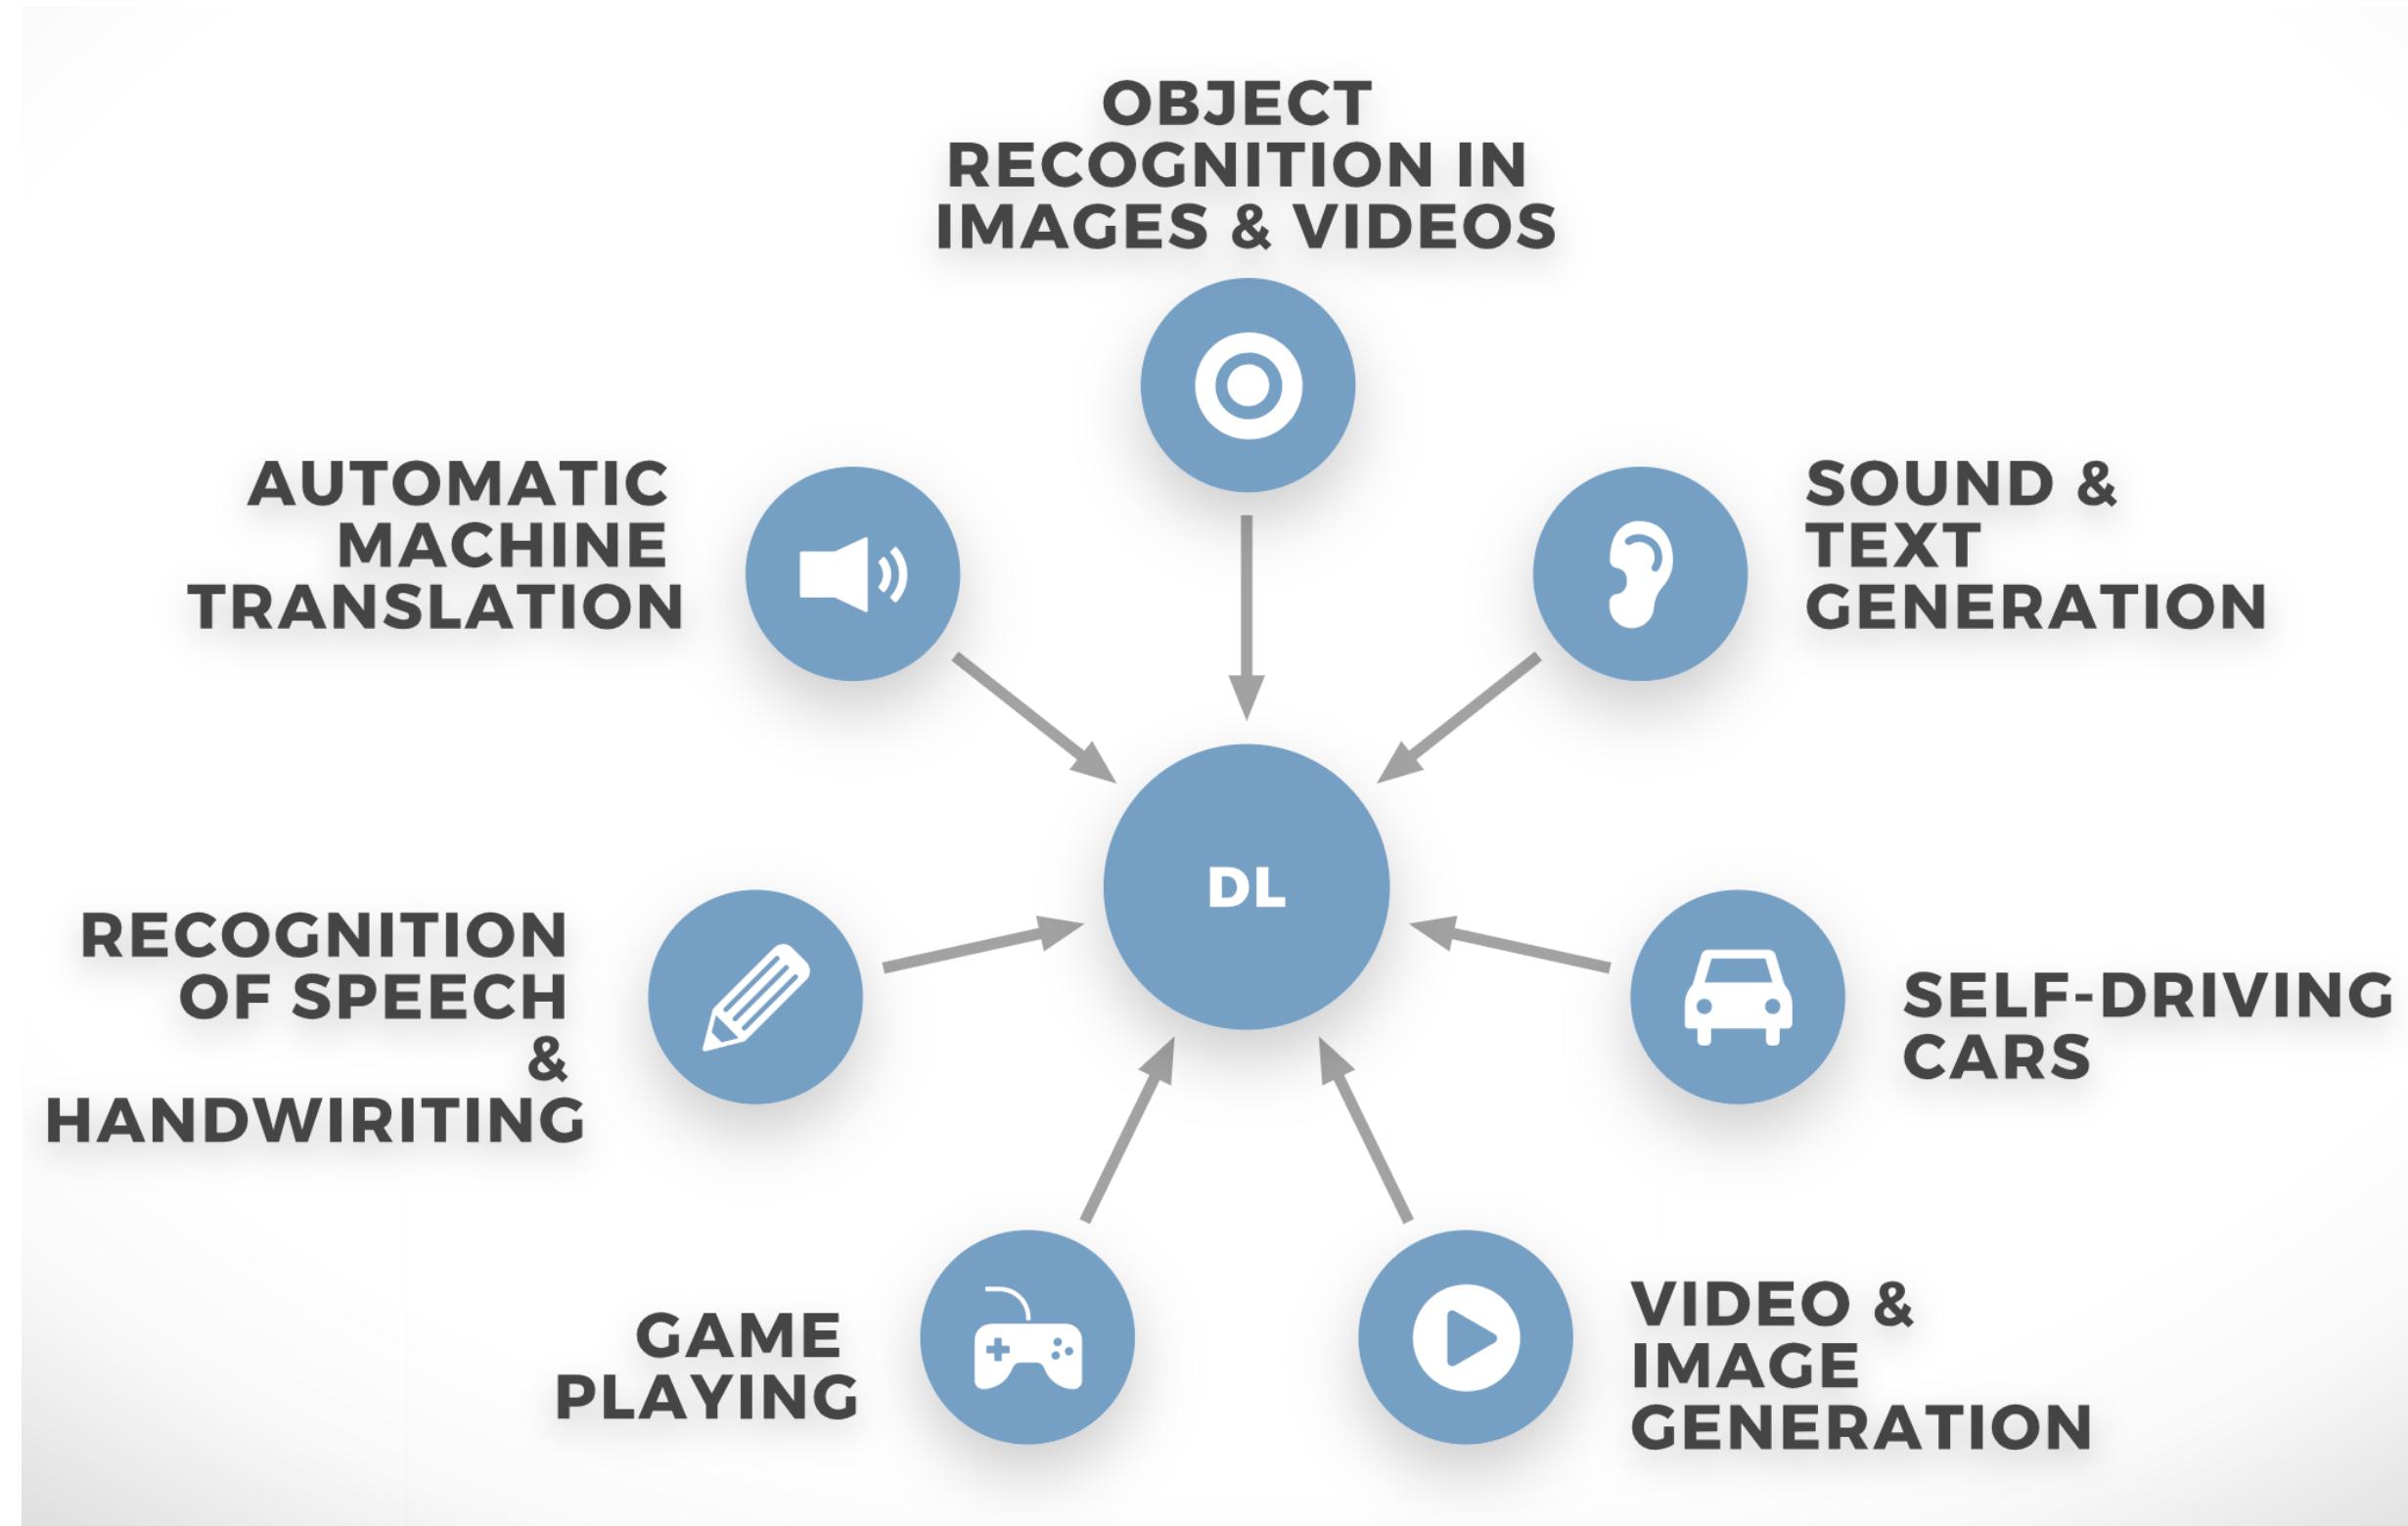


Deep Learning



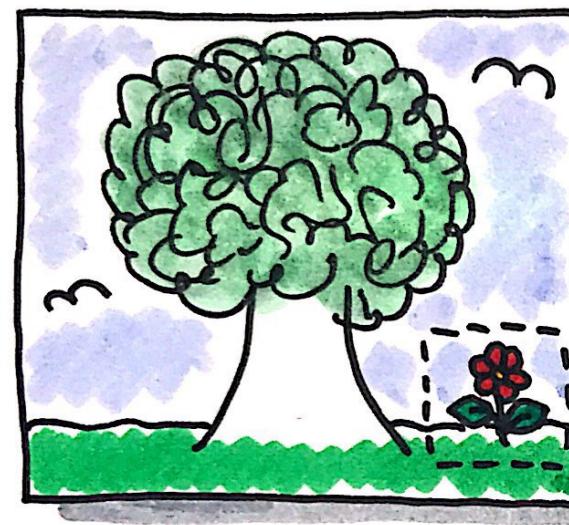
- supervised, semi-supervised or unsupervised
- NLP, speech recognition
- image recognition
- object classification
- recommender systems
- etc.

Deep Learning in the wild



How does a computer learn to "see"?

Convolutional Neural Nets



- image classification
class = tree
- object detection
flower

Computer
Vision

How does a computer learn to "see"?

Jupyter Lab and Python

- i(nteractive) Python
- browser based notebook
- code + markdown + output



Jupyter tips'n'tricks

- ?: search documentation

```
In [5]: len?
```

- ??: look at source code

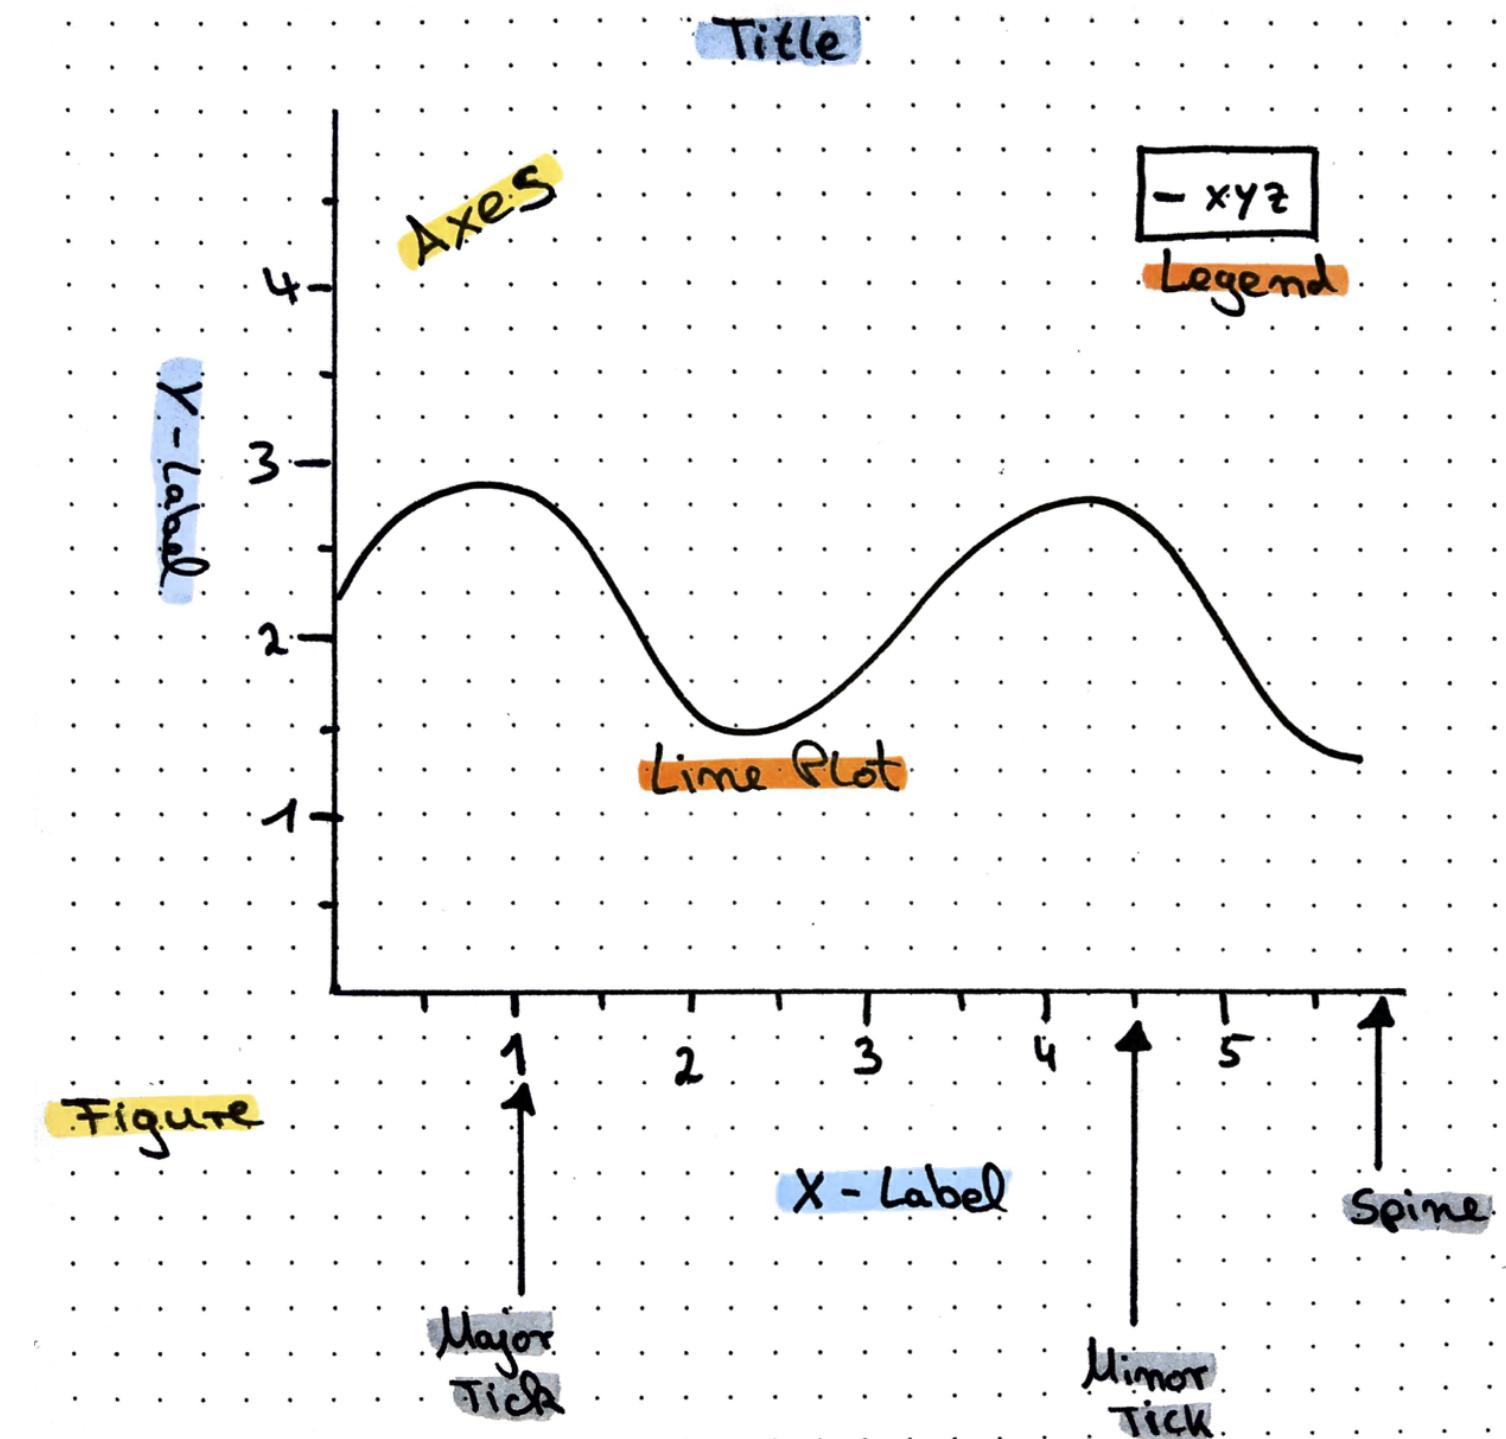
```
In [6]: len??
```

- tab: autocomplete
- shift + tab: look at function call

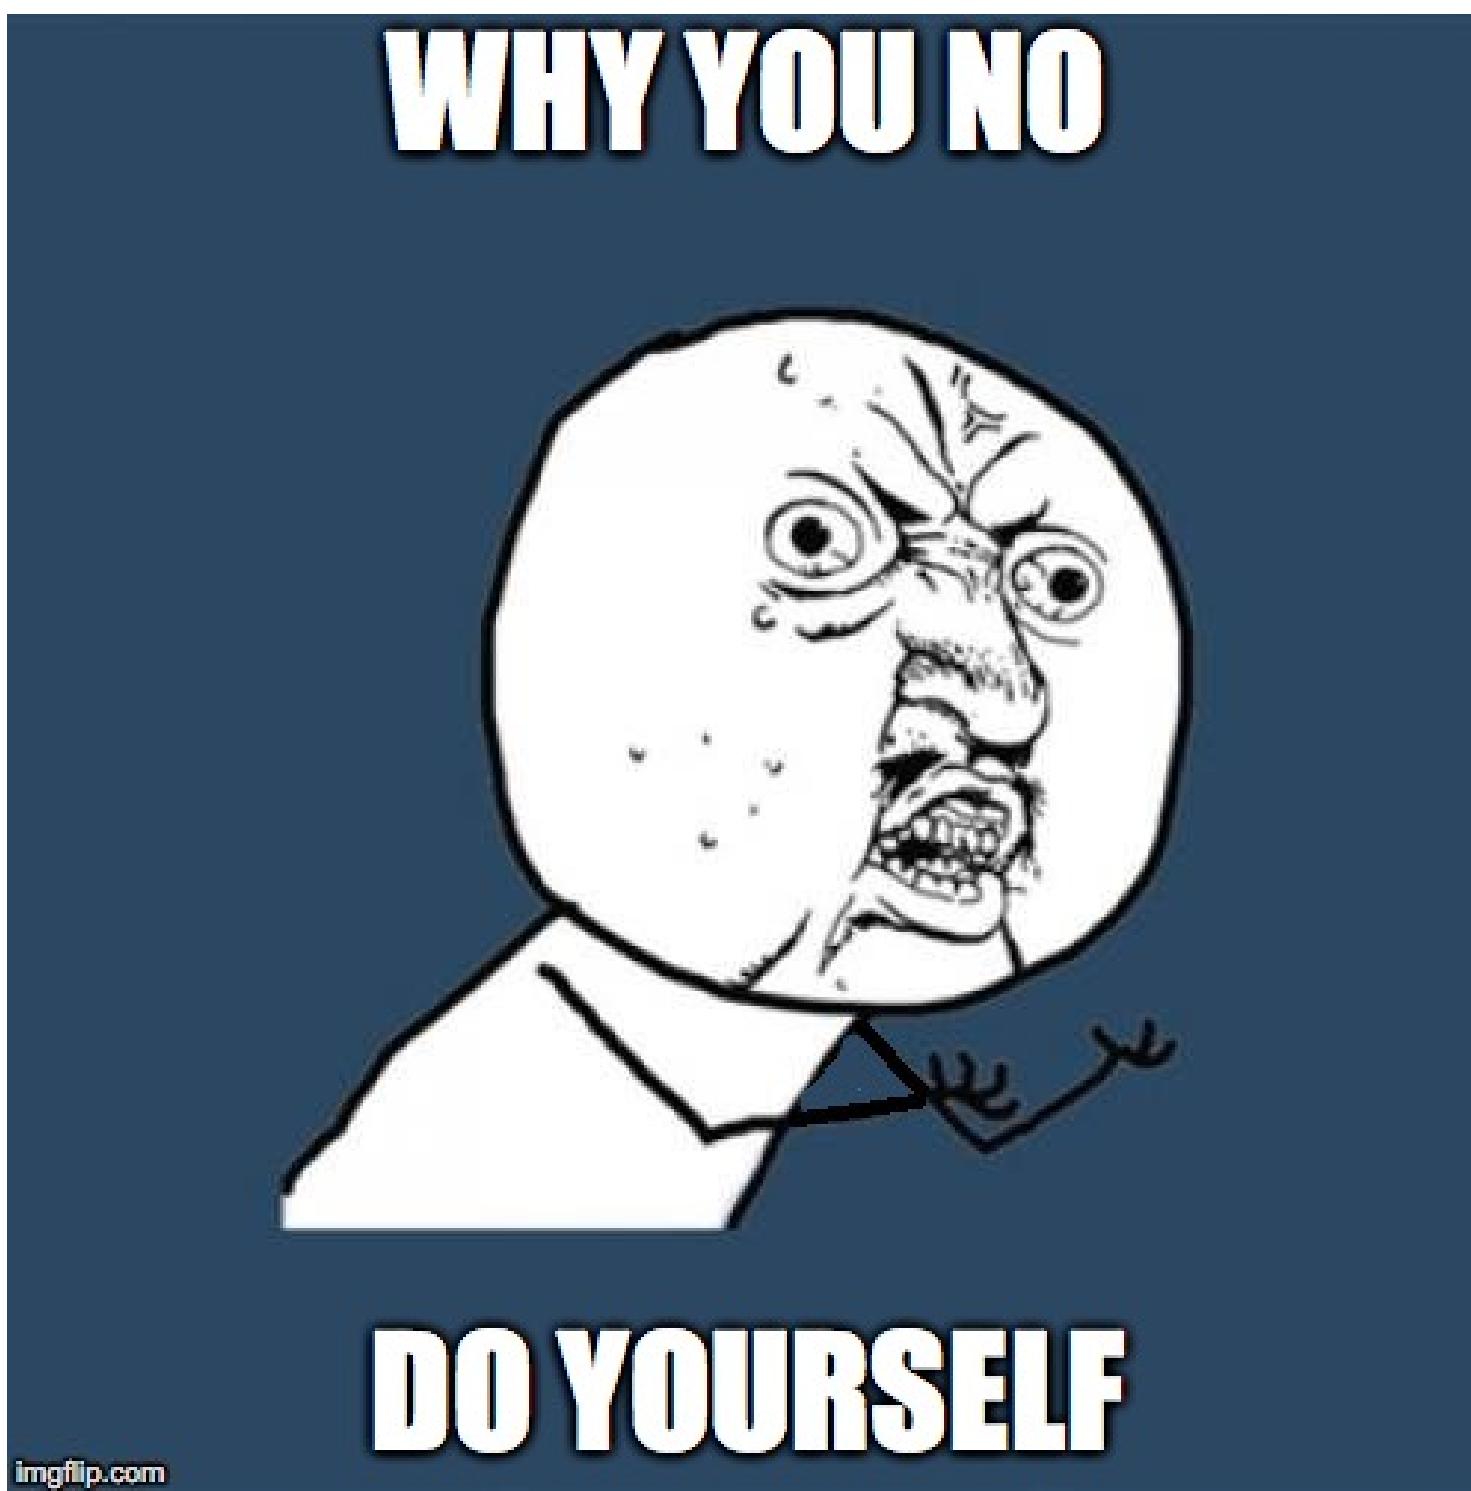


Visualization with Matplotlib

- Matplotlib is a Python 2D plotting library
- Can be used in Python scripts, Python/IPython shells, Jupyter notebooks



Matplotlib

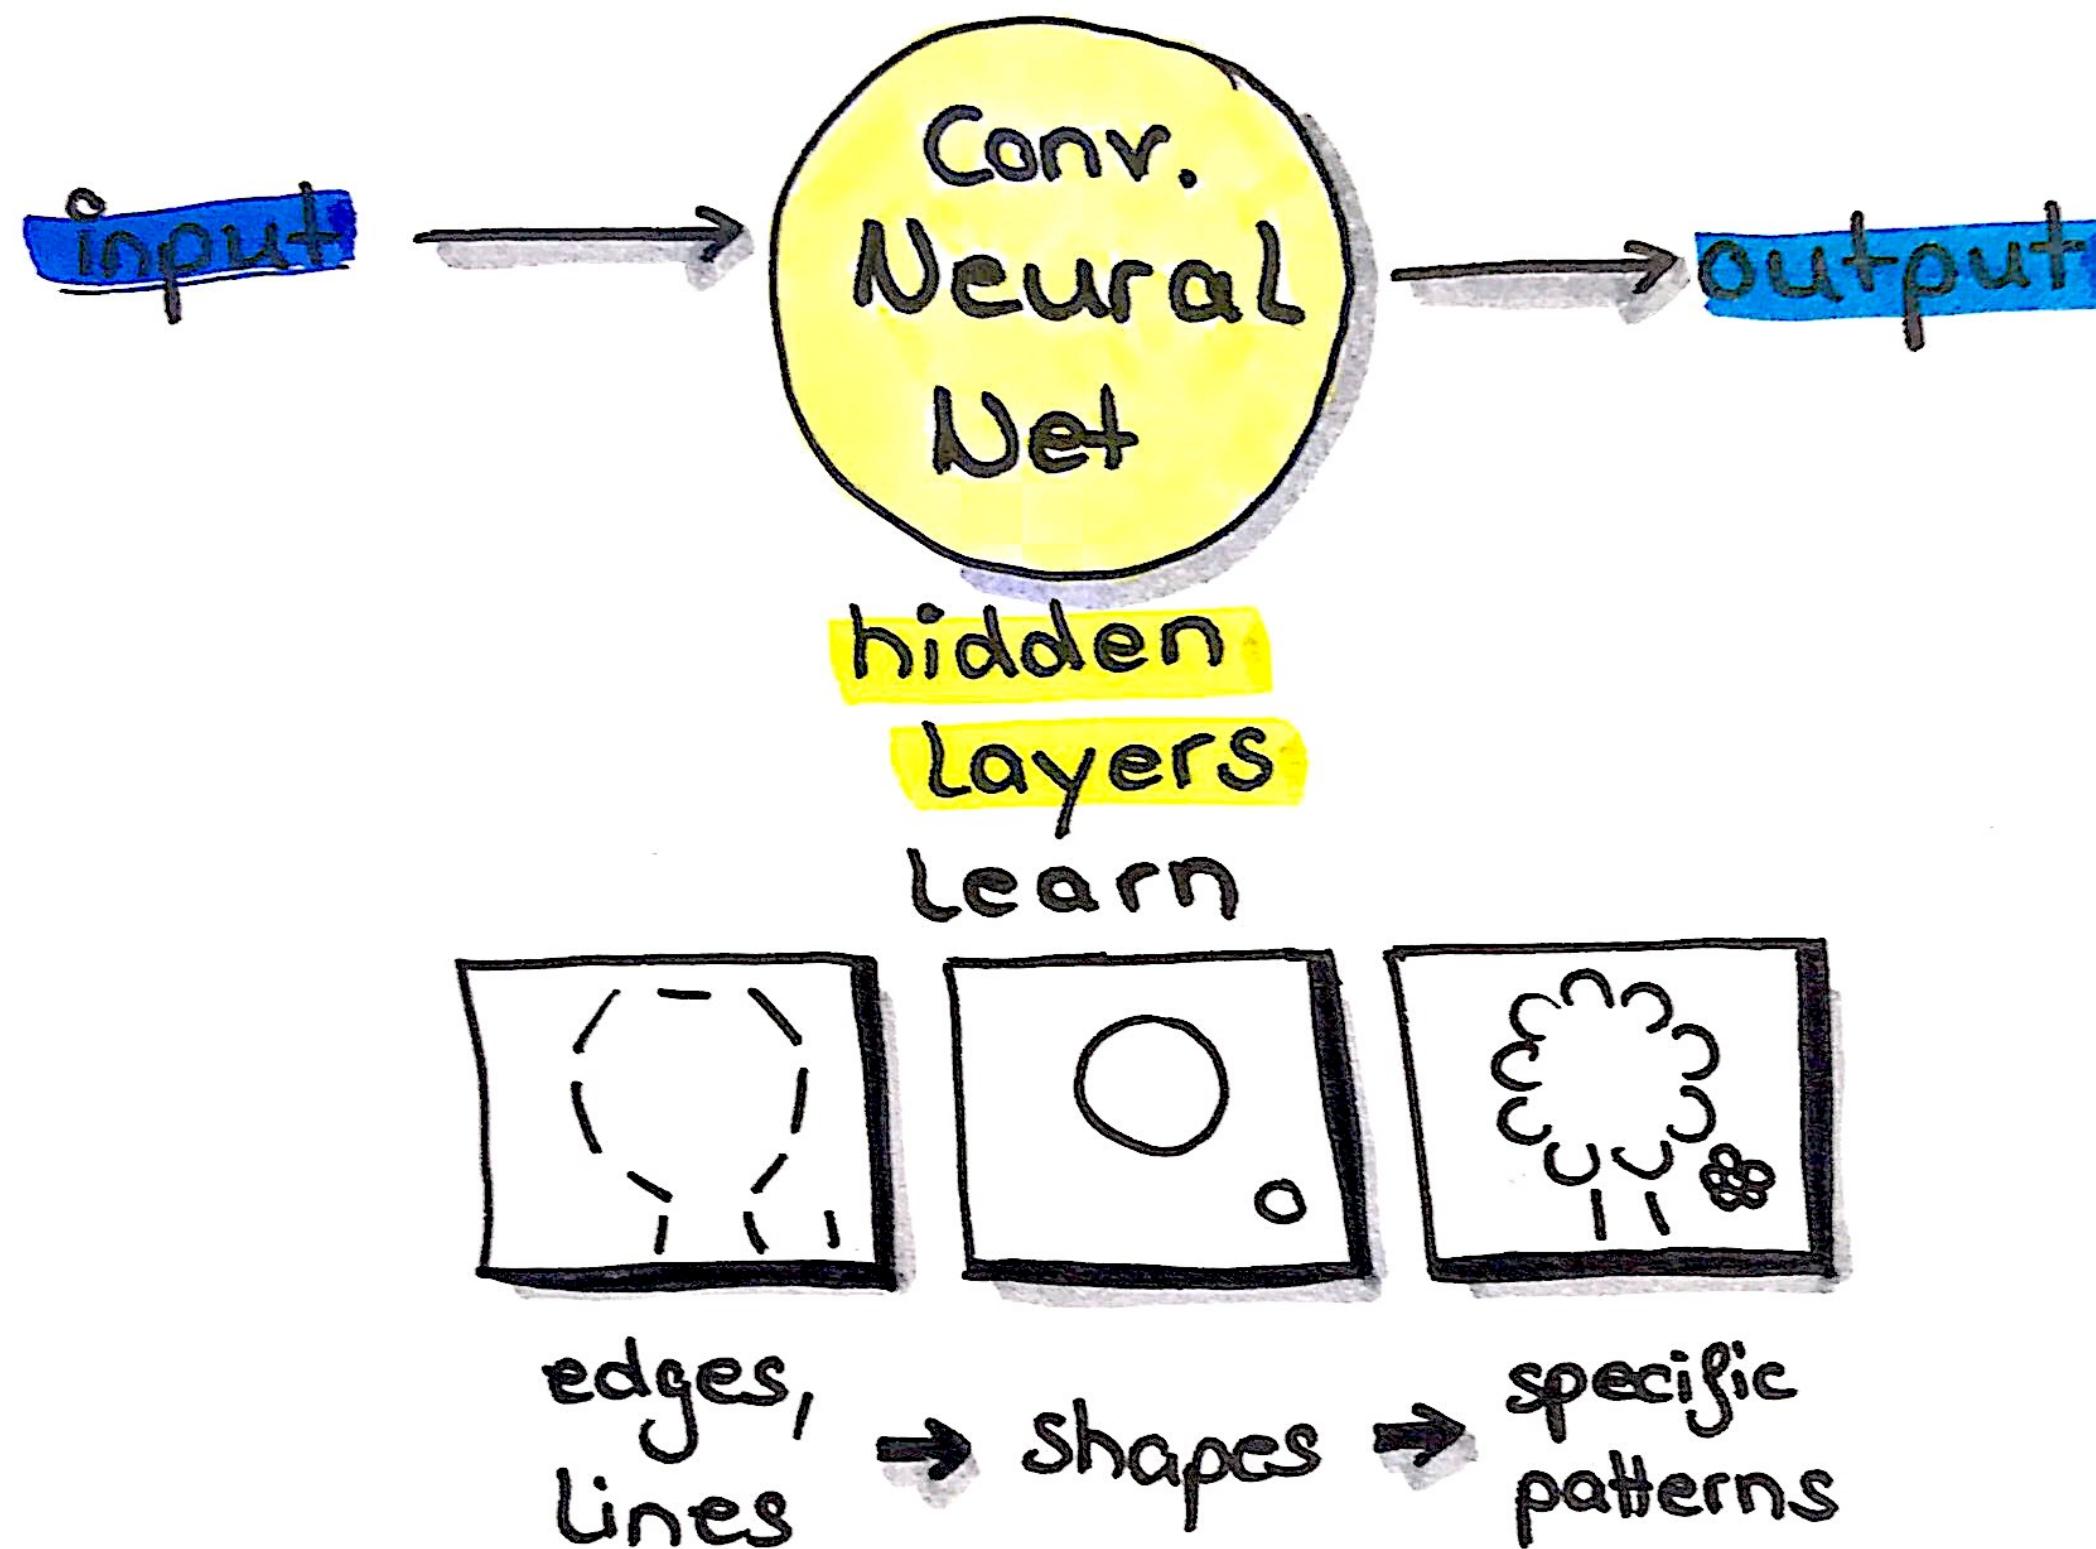


Convolutional Neural Nets

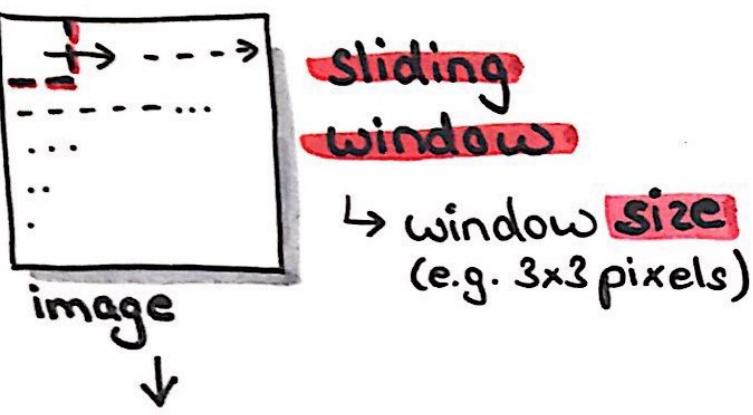
MLPs vs CNNs

- pixels are considered independent
- computationally faster
- Learned : weights
- pixels are considered as groups of connected information (context)
- analysed as chunks (= windows)
- Learned : filters

ConvNets



ConvNets



1	0	-1
4	0	-1
1	0	-1

vertical Lines

1	1	1
0	0	0
-1	-1	-1

horizontal Lines

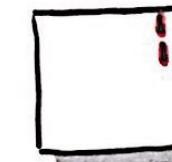
- **Filters**

detect shapes & patterns
in window chunks

- multiple filters are combined

- filters are learned

padding



- fake pixels are created at the edge of images to incorporate border pixels

Convolutional Layers

apply filters



output



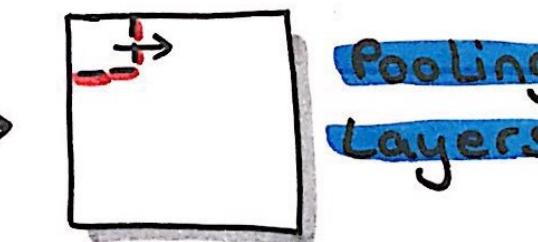
feature maps

→ stacks of

feature maps

stride

- how much overlap to have in sliding window

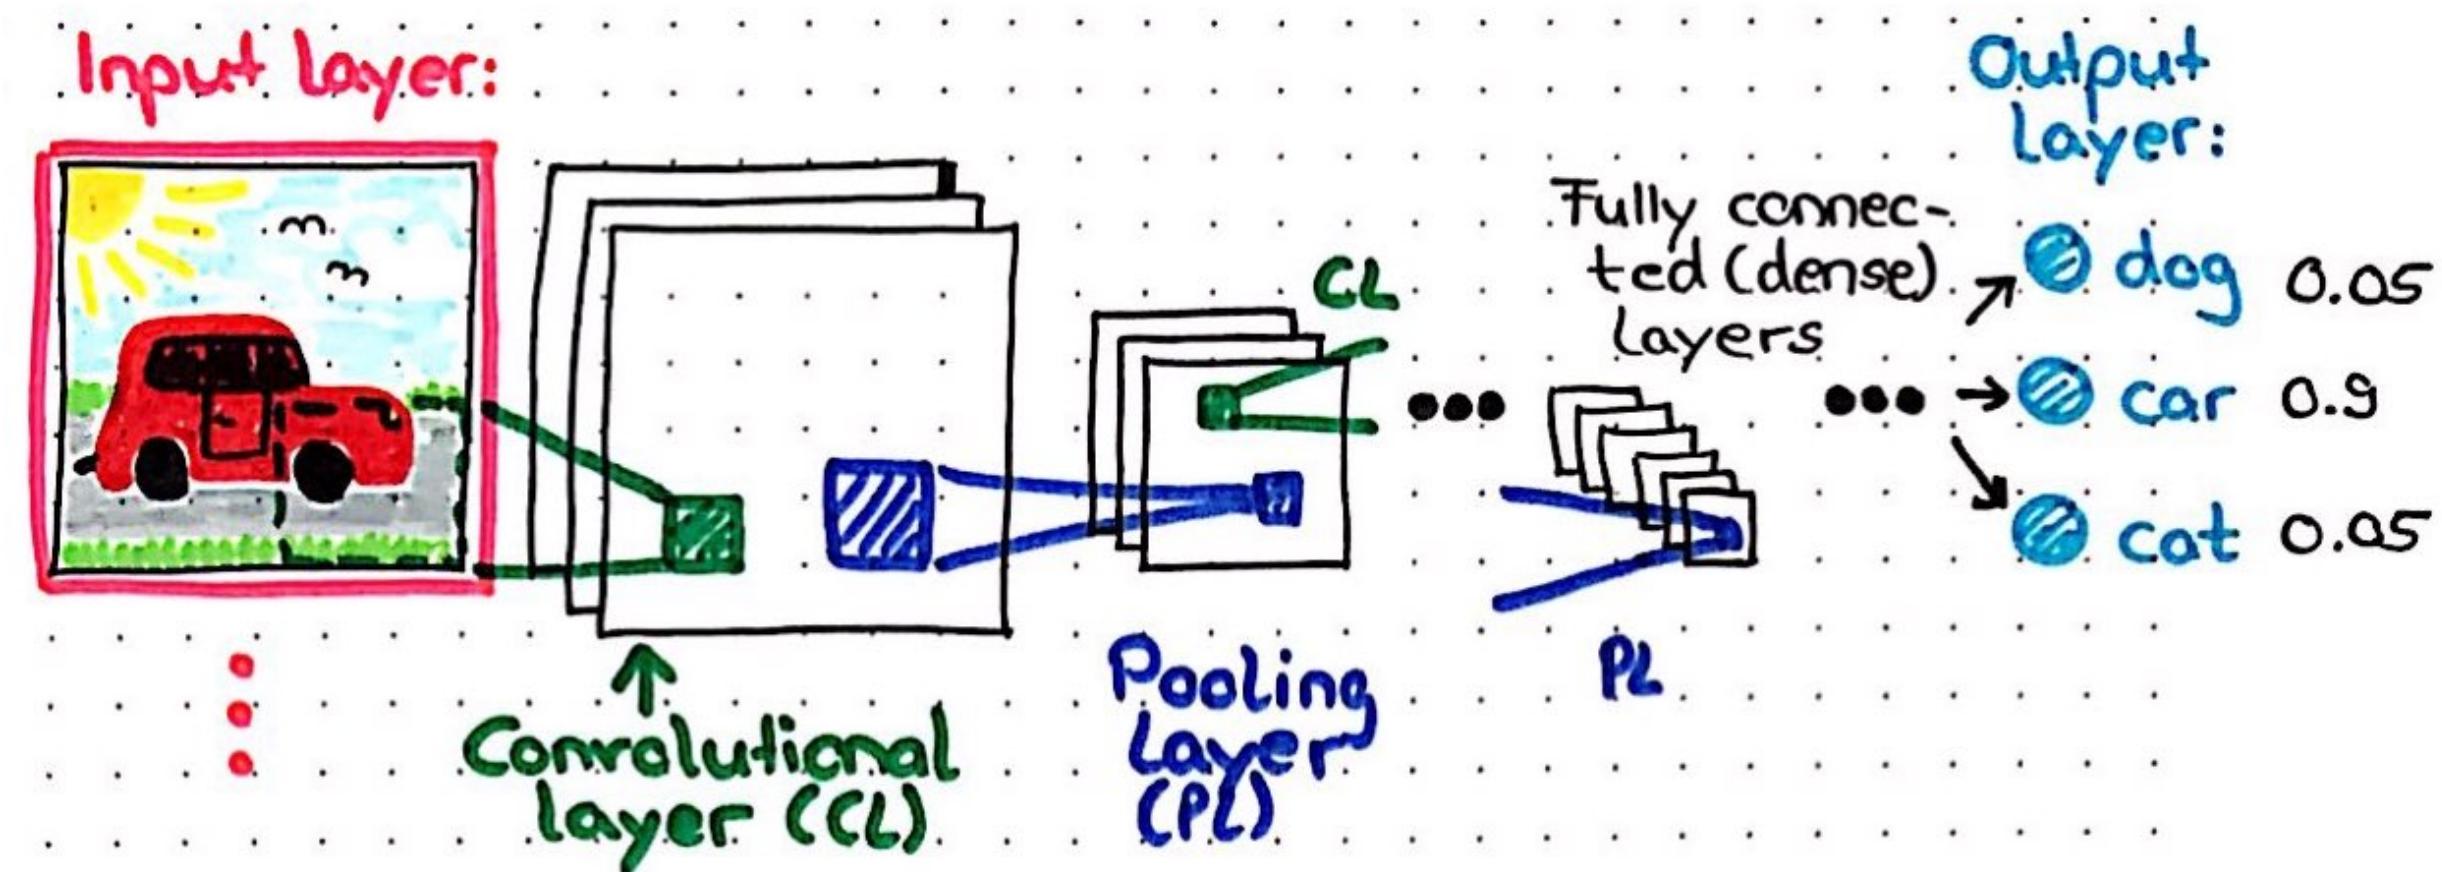


e.g. max pooling

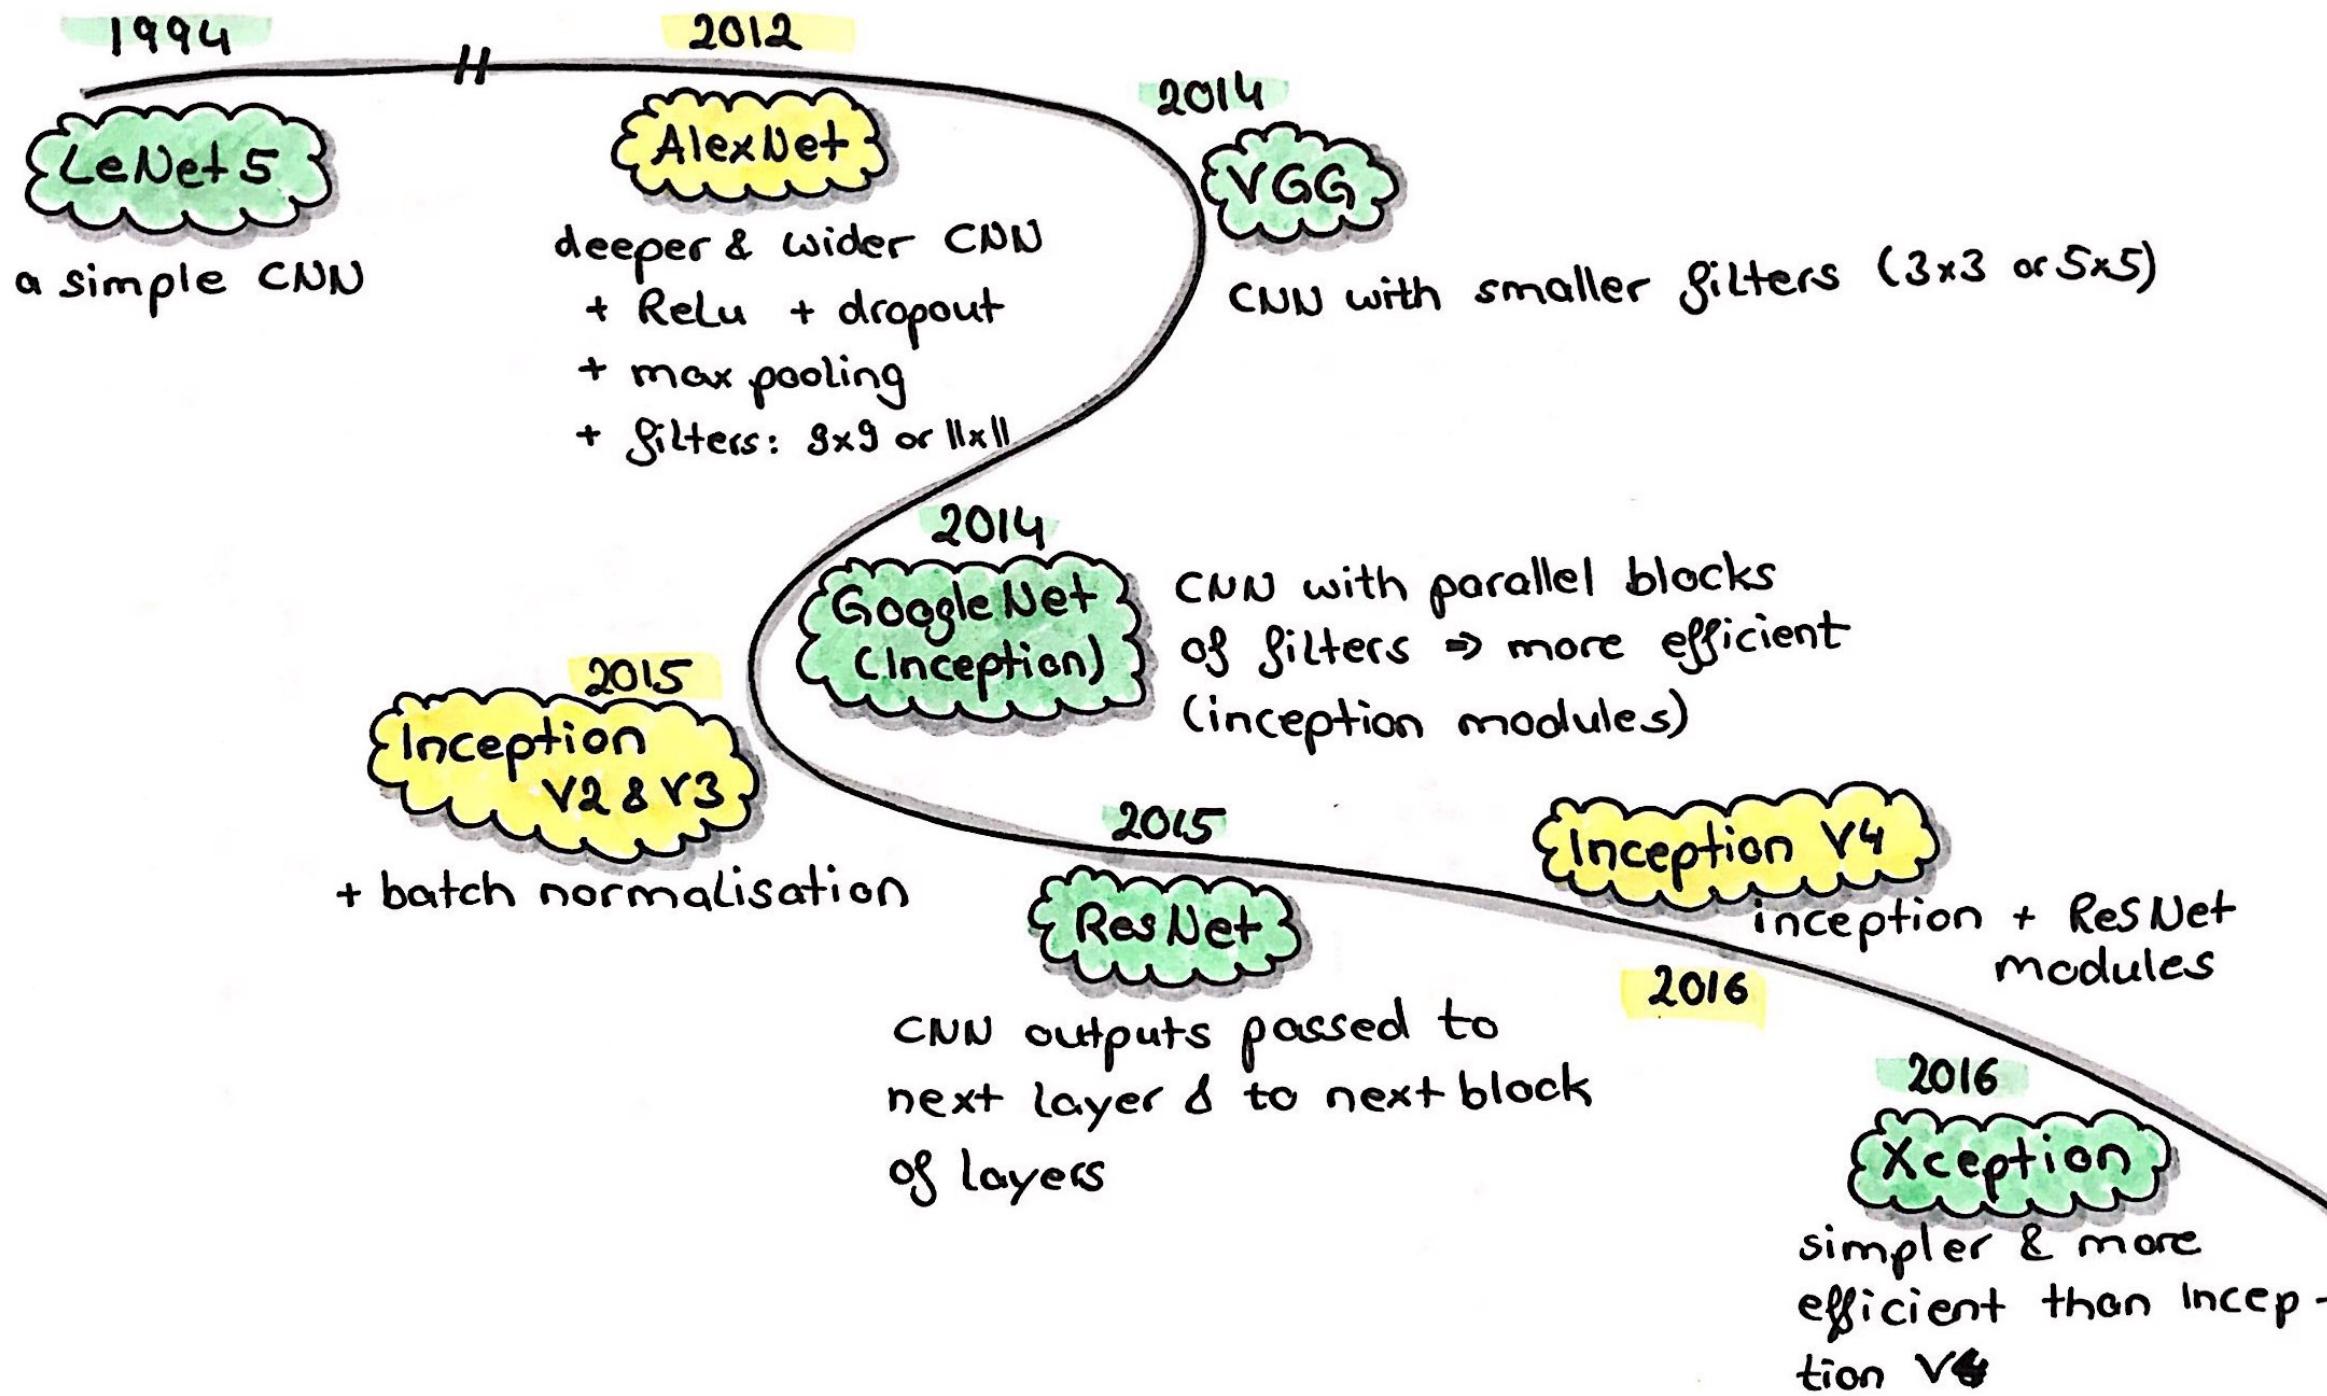
1	0	0
4	8	0
6	1	1

- reduces compute time
- boils information down

ConvNets

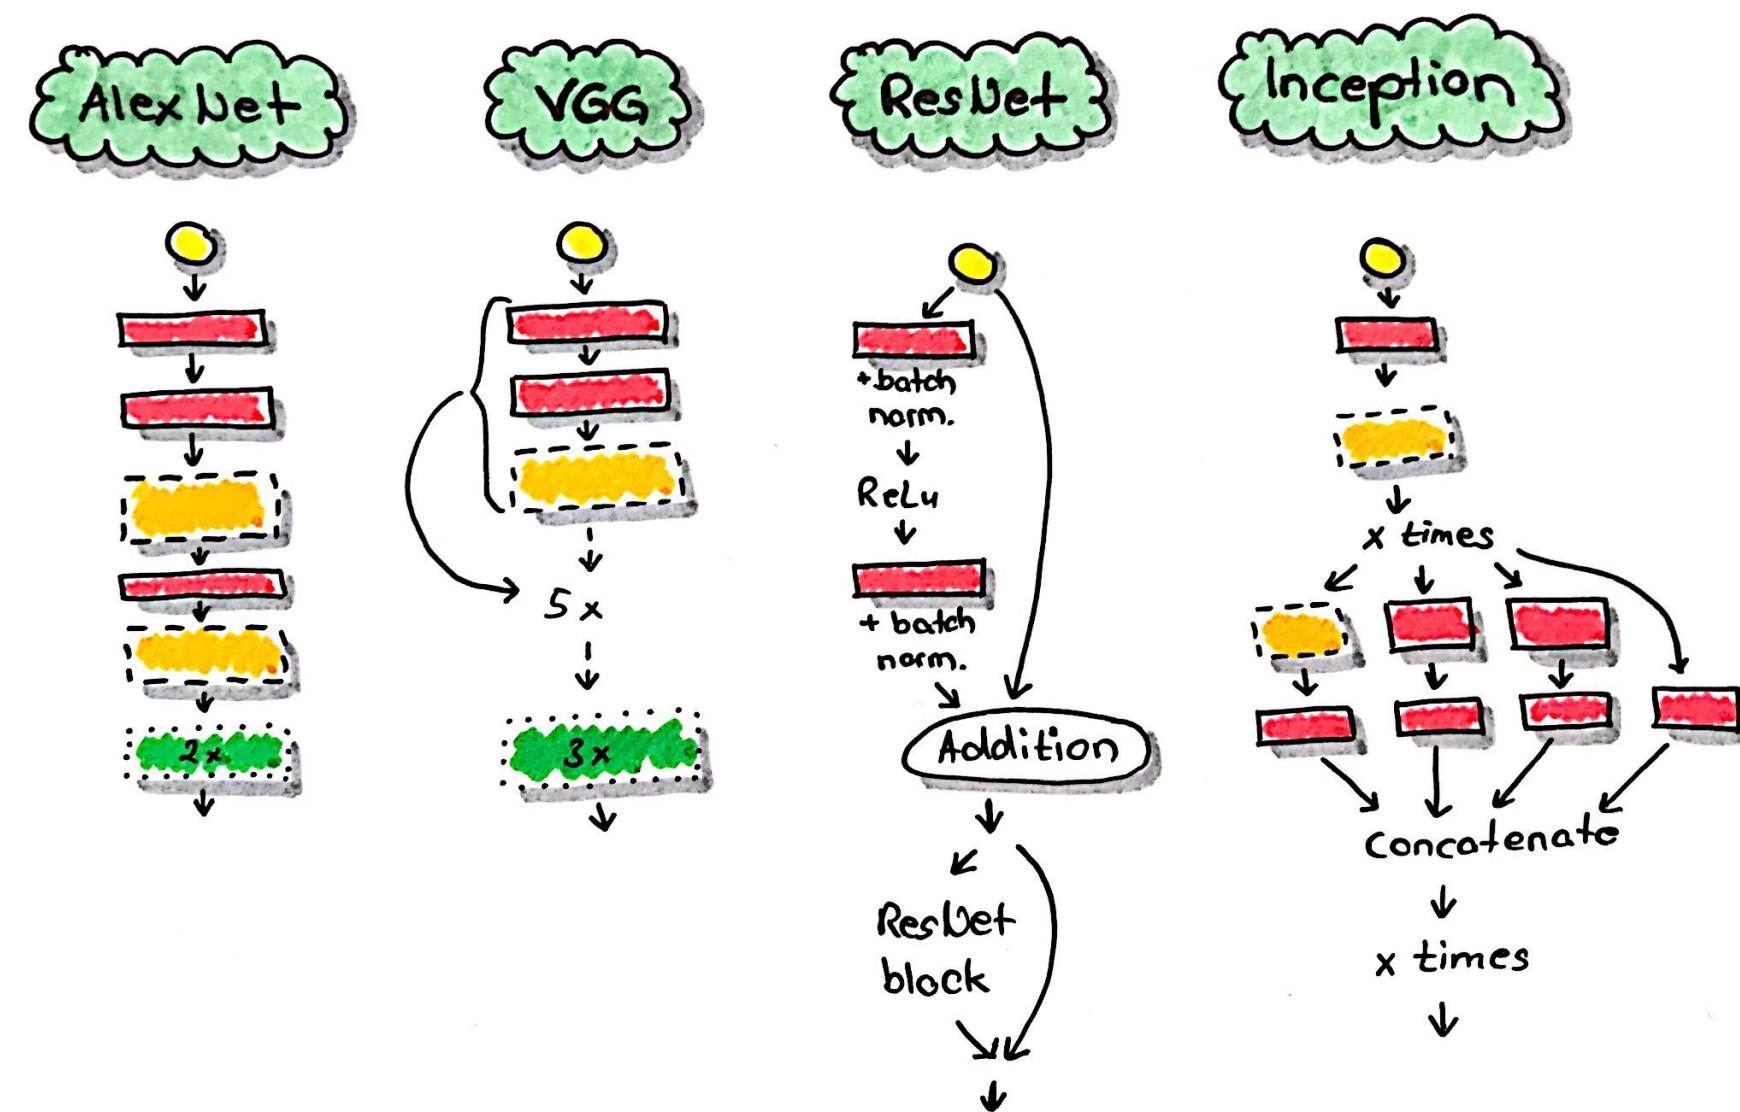


Evolution of neural nets for image recognition

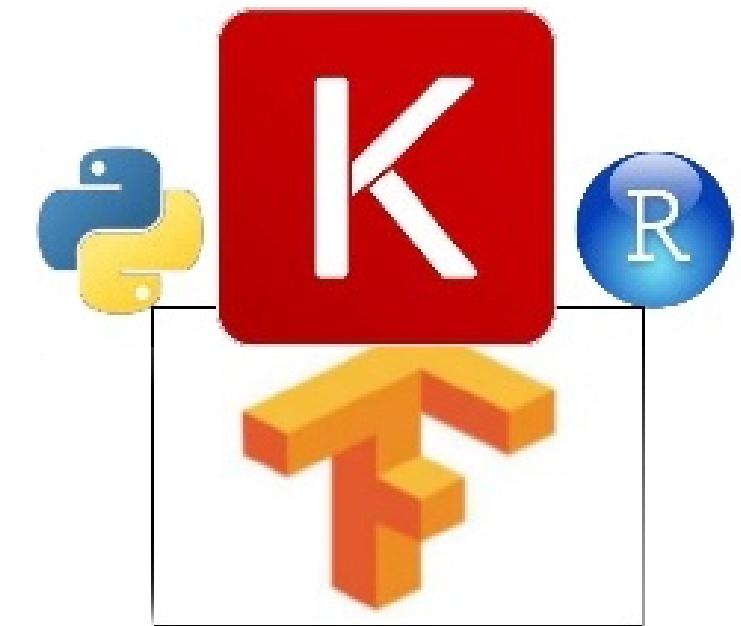


CNN architectures

- input image
- convolutional layer
- pooling layer
- dense layer



Introduction to TensorFlow



What are tensors?



Tensors = multidimensional arrays

Dimension

1

1	2	3
---	---	---

vector

2

1	2	3
4	5	6
7	8	9

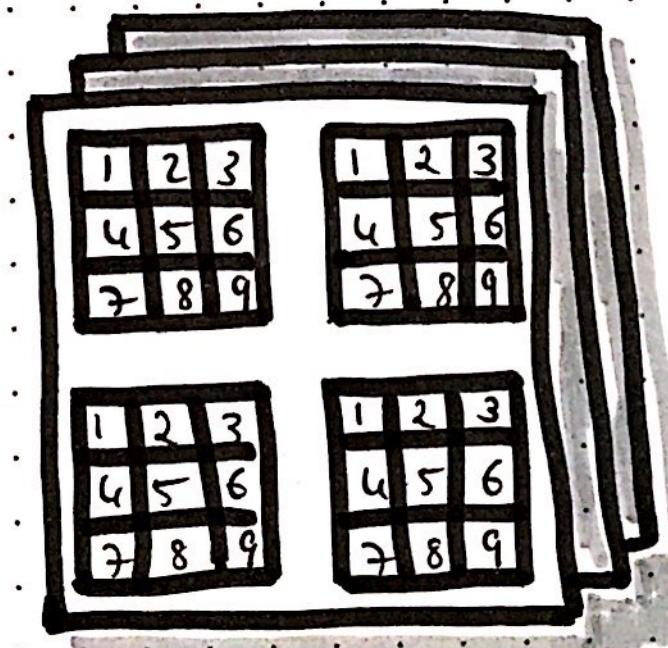
matrix

3

1	2	3
4	5	6
7	8	9

3-dim.
array

n



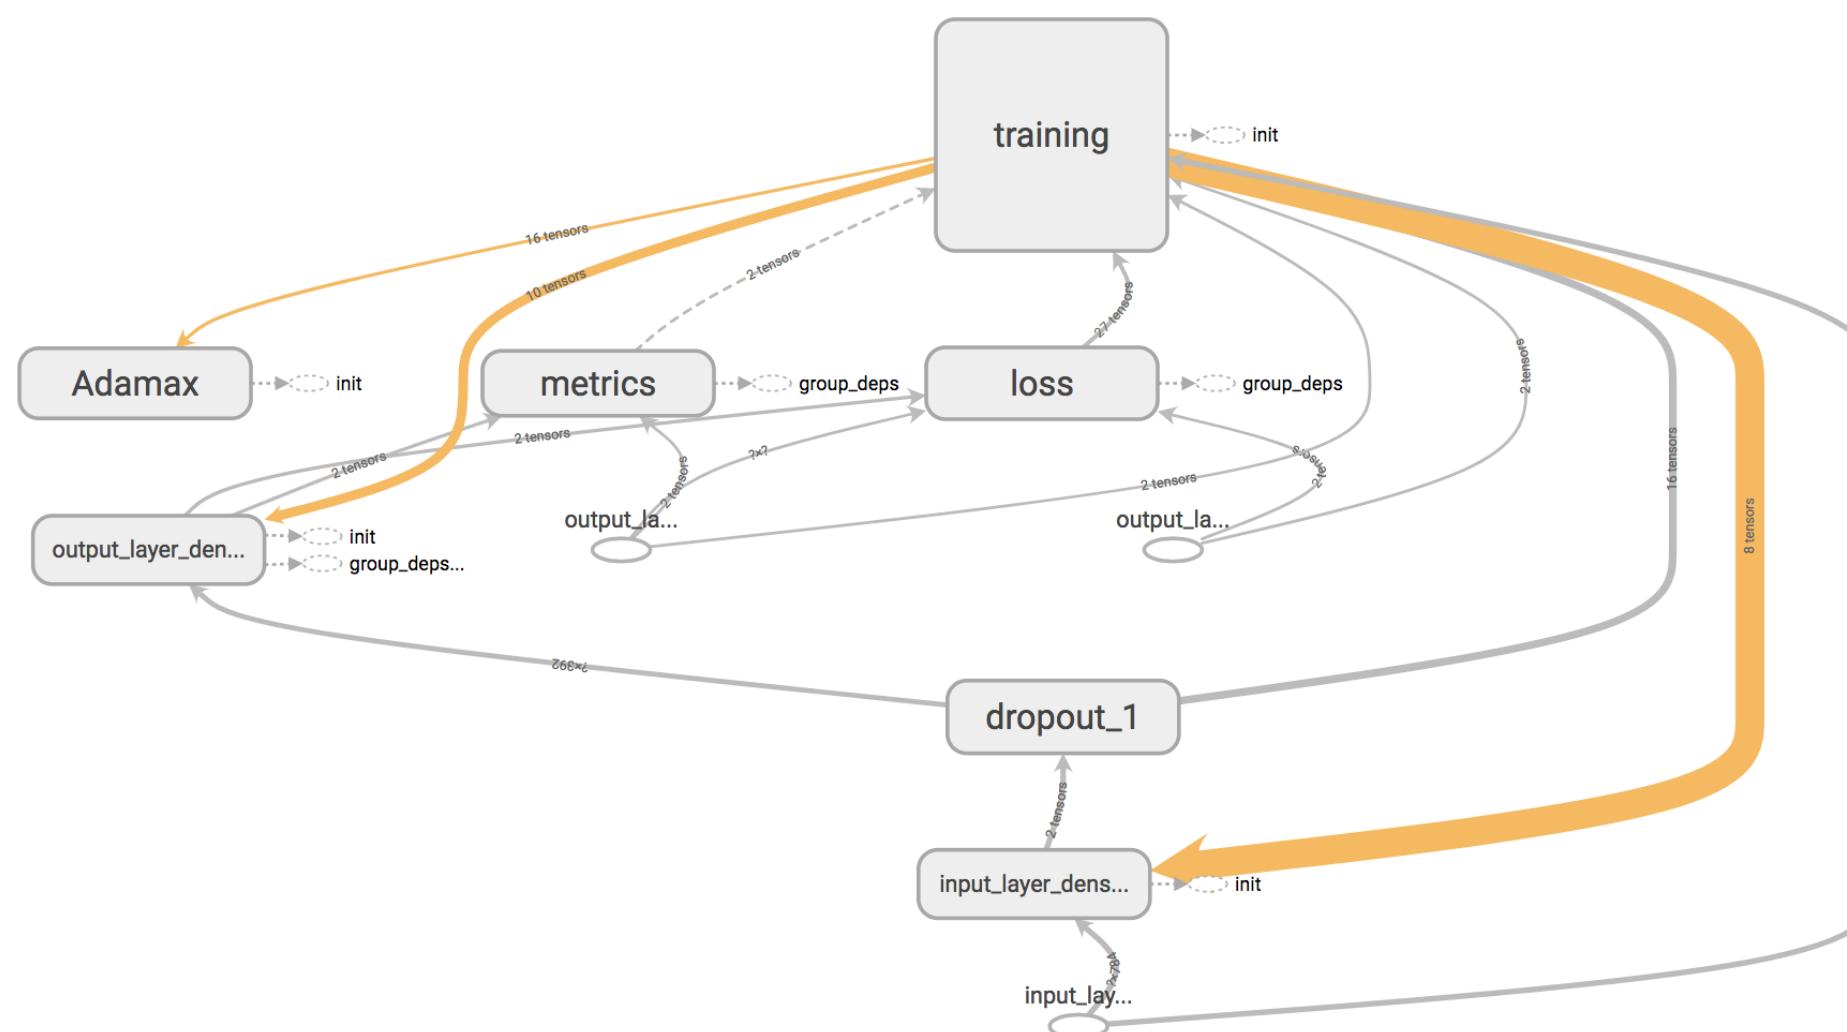
n-dim.
array

Tensor "Flow"



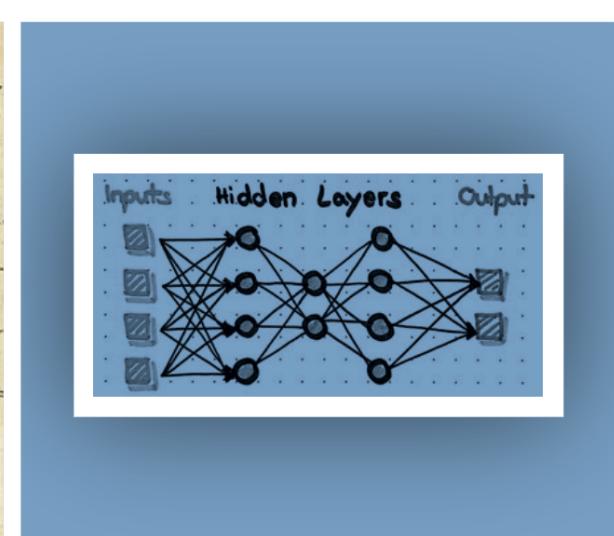
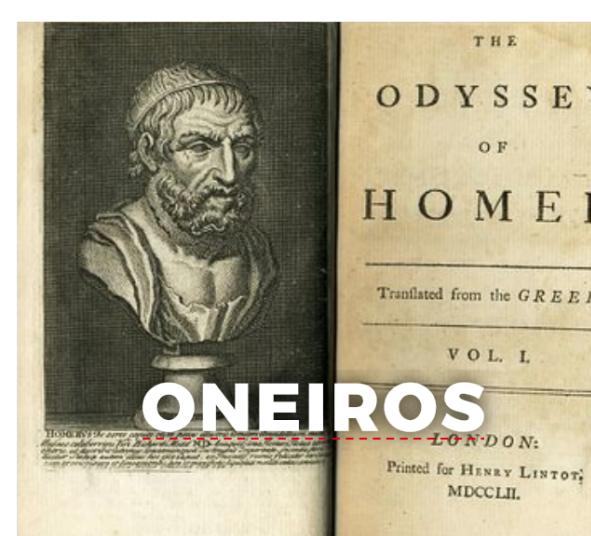
Graphs in TensorBoard

Main Graph



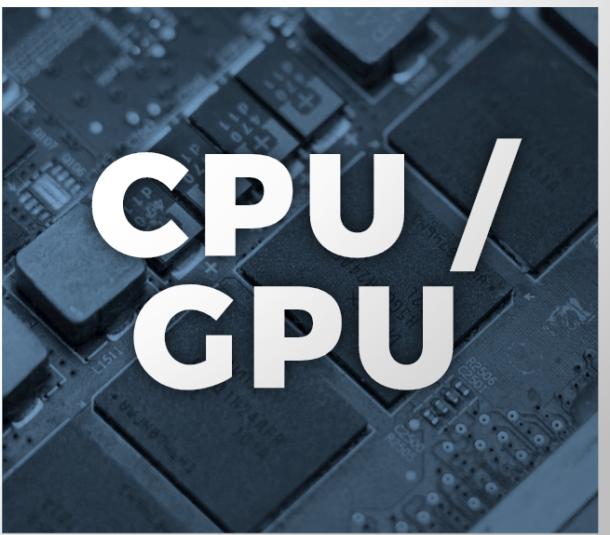
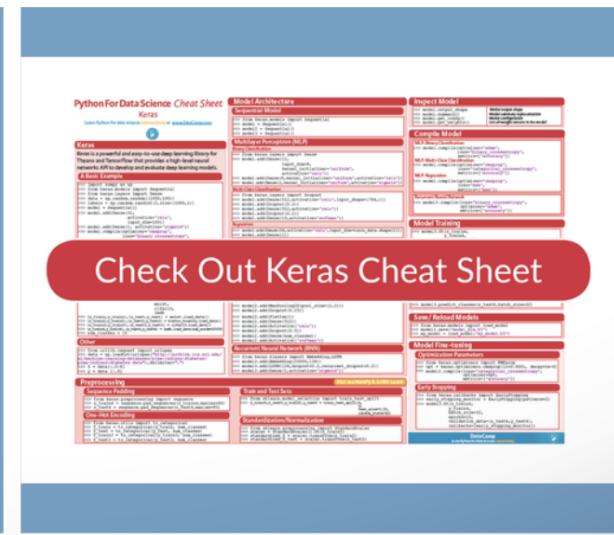
Auxiliary Nodes

Keras High-Level API for TensorFlow

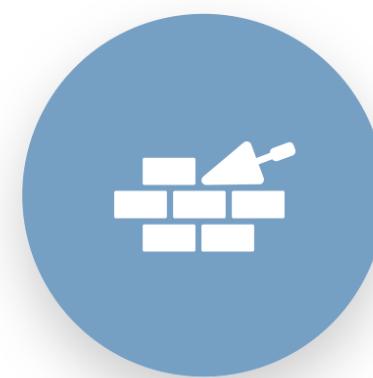


"Oneiroi are beyond our unravelling -
who can be sure what tale they tell?
Not all that men look for comes to pass.
Two gates there are that give passage to fleeting Oneiroi;
one is made of horn, one of ivory.
The Oneiroi that pass through sawn ivory are deceitful,
bearing a message that will not be fulfilled;
those that come out through polished horn have truth behind them,
to be accomplished for men who see them."

Homer, Odyssey 19. 562 ff (Shewring translation).



Keras APIs



SEQUENTIAL MODELS

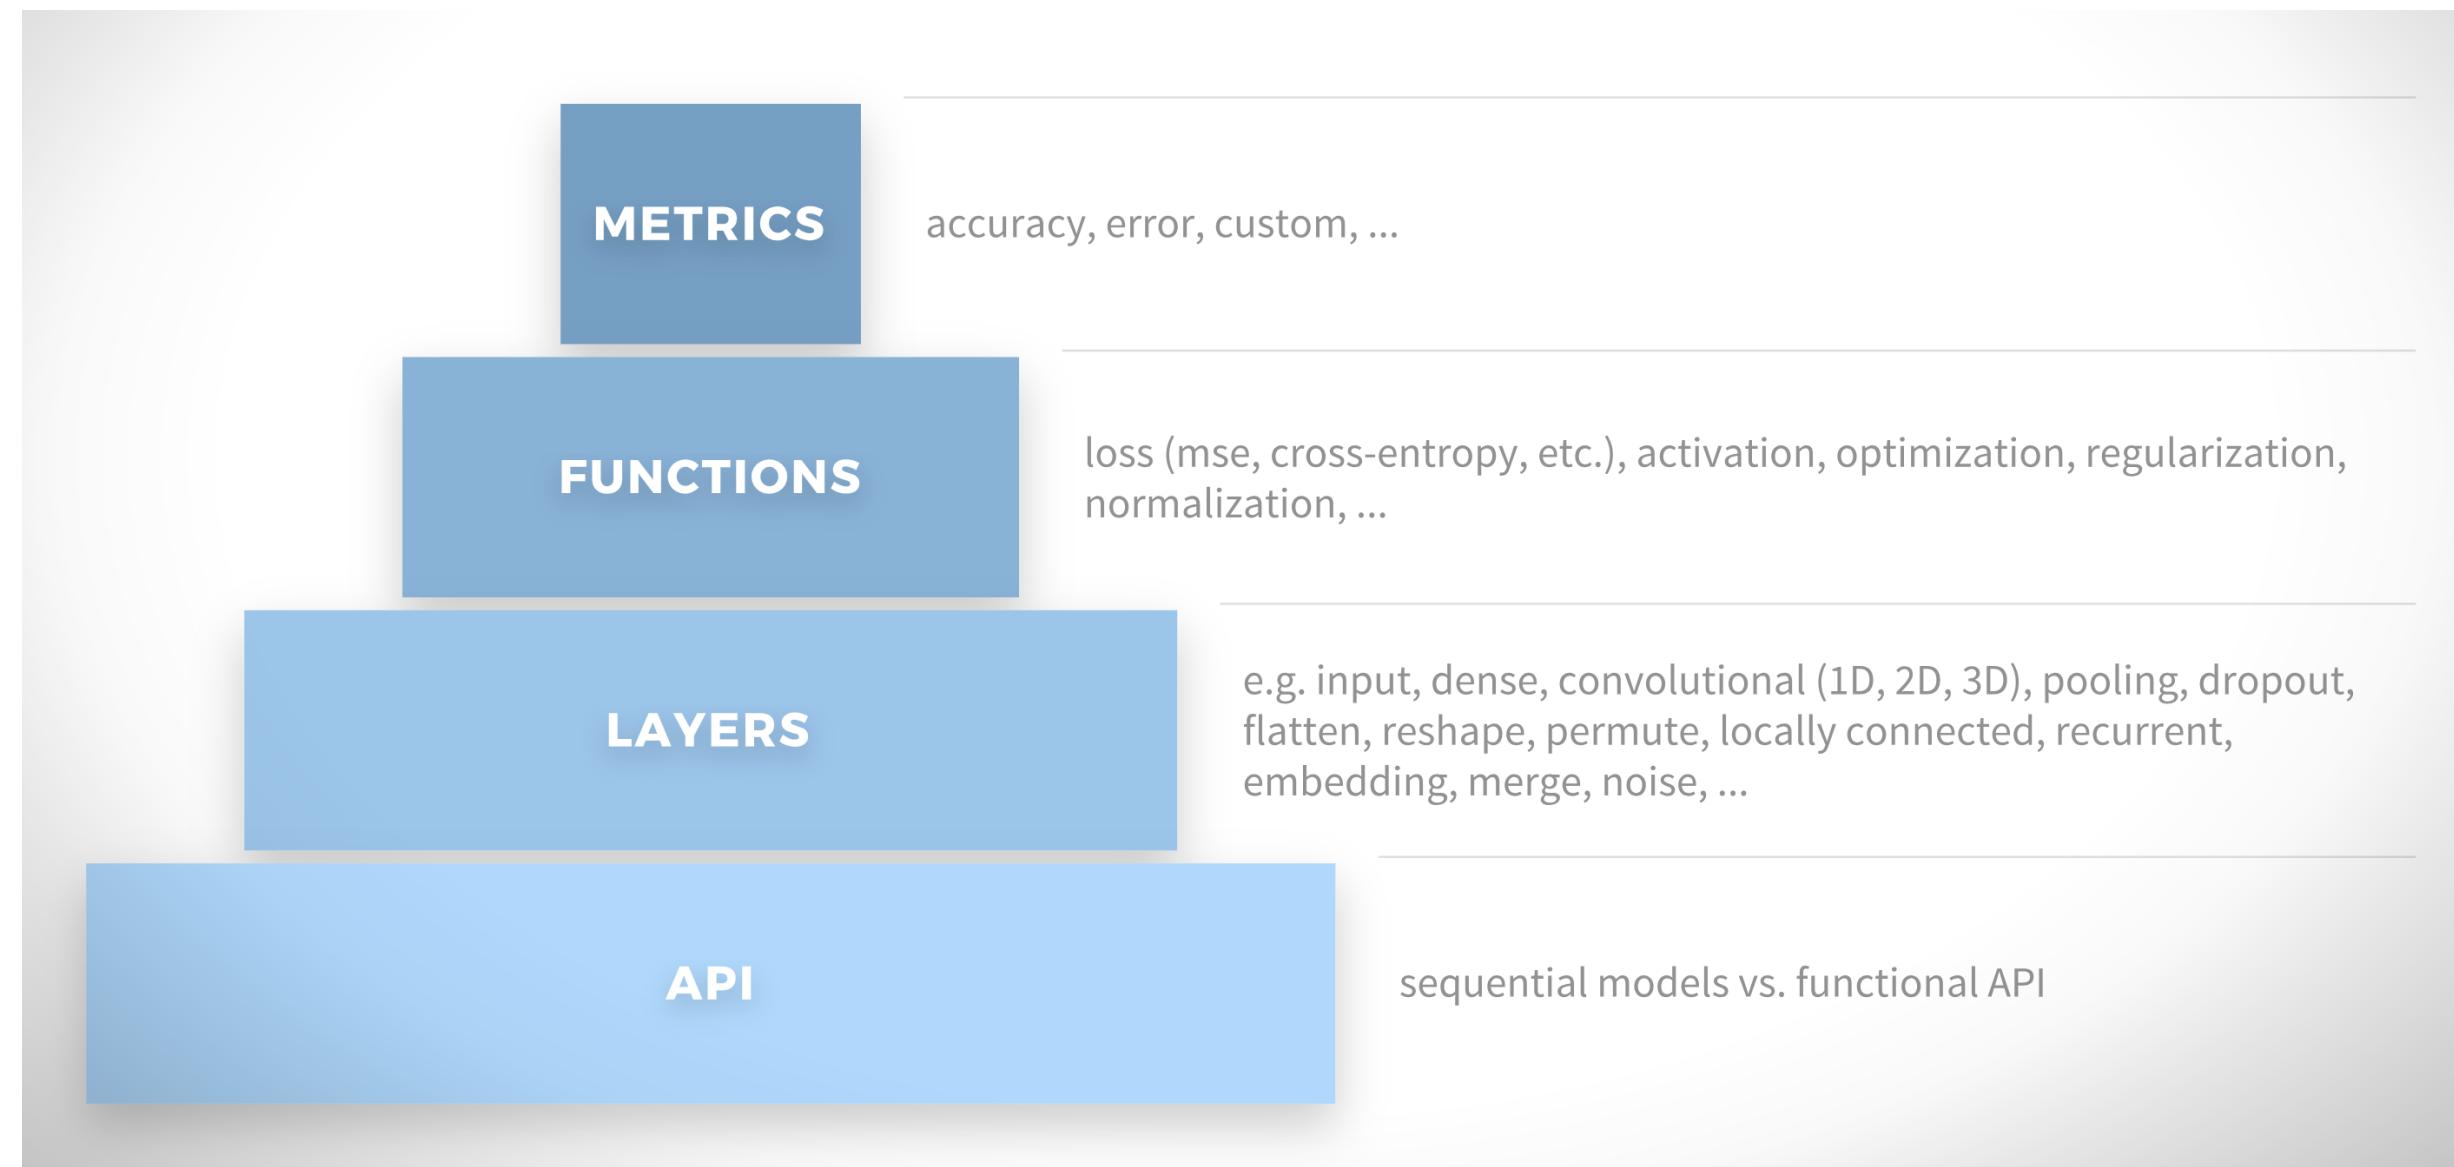
simple
suitable for most cases
linear order of layers
only one direction from input to output



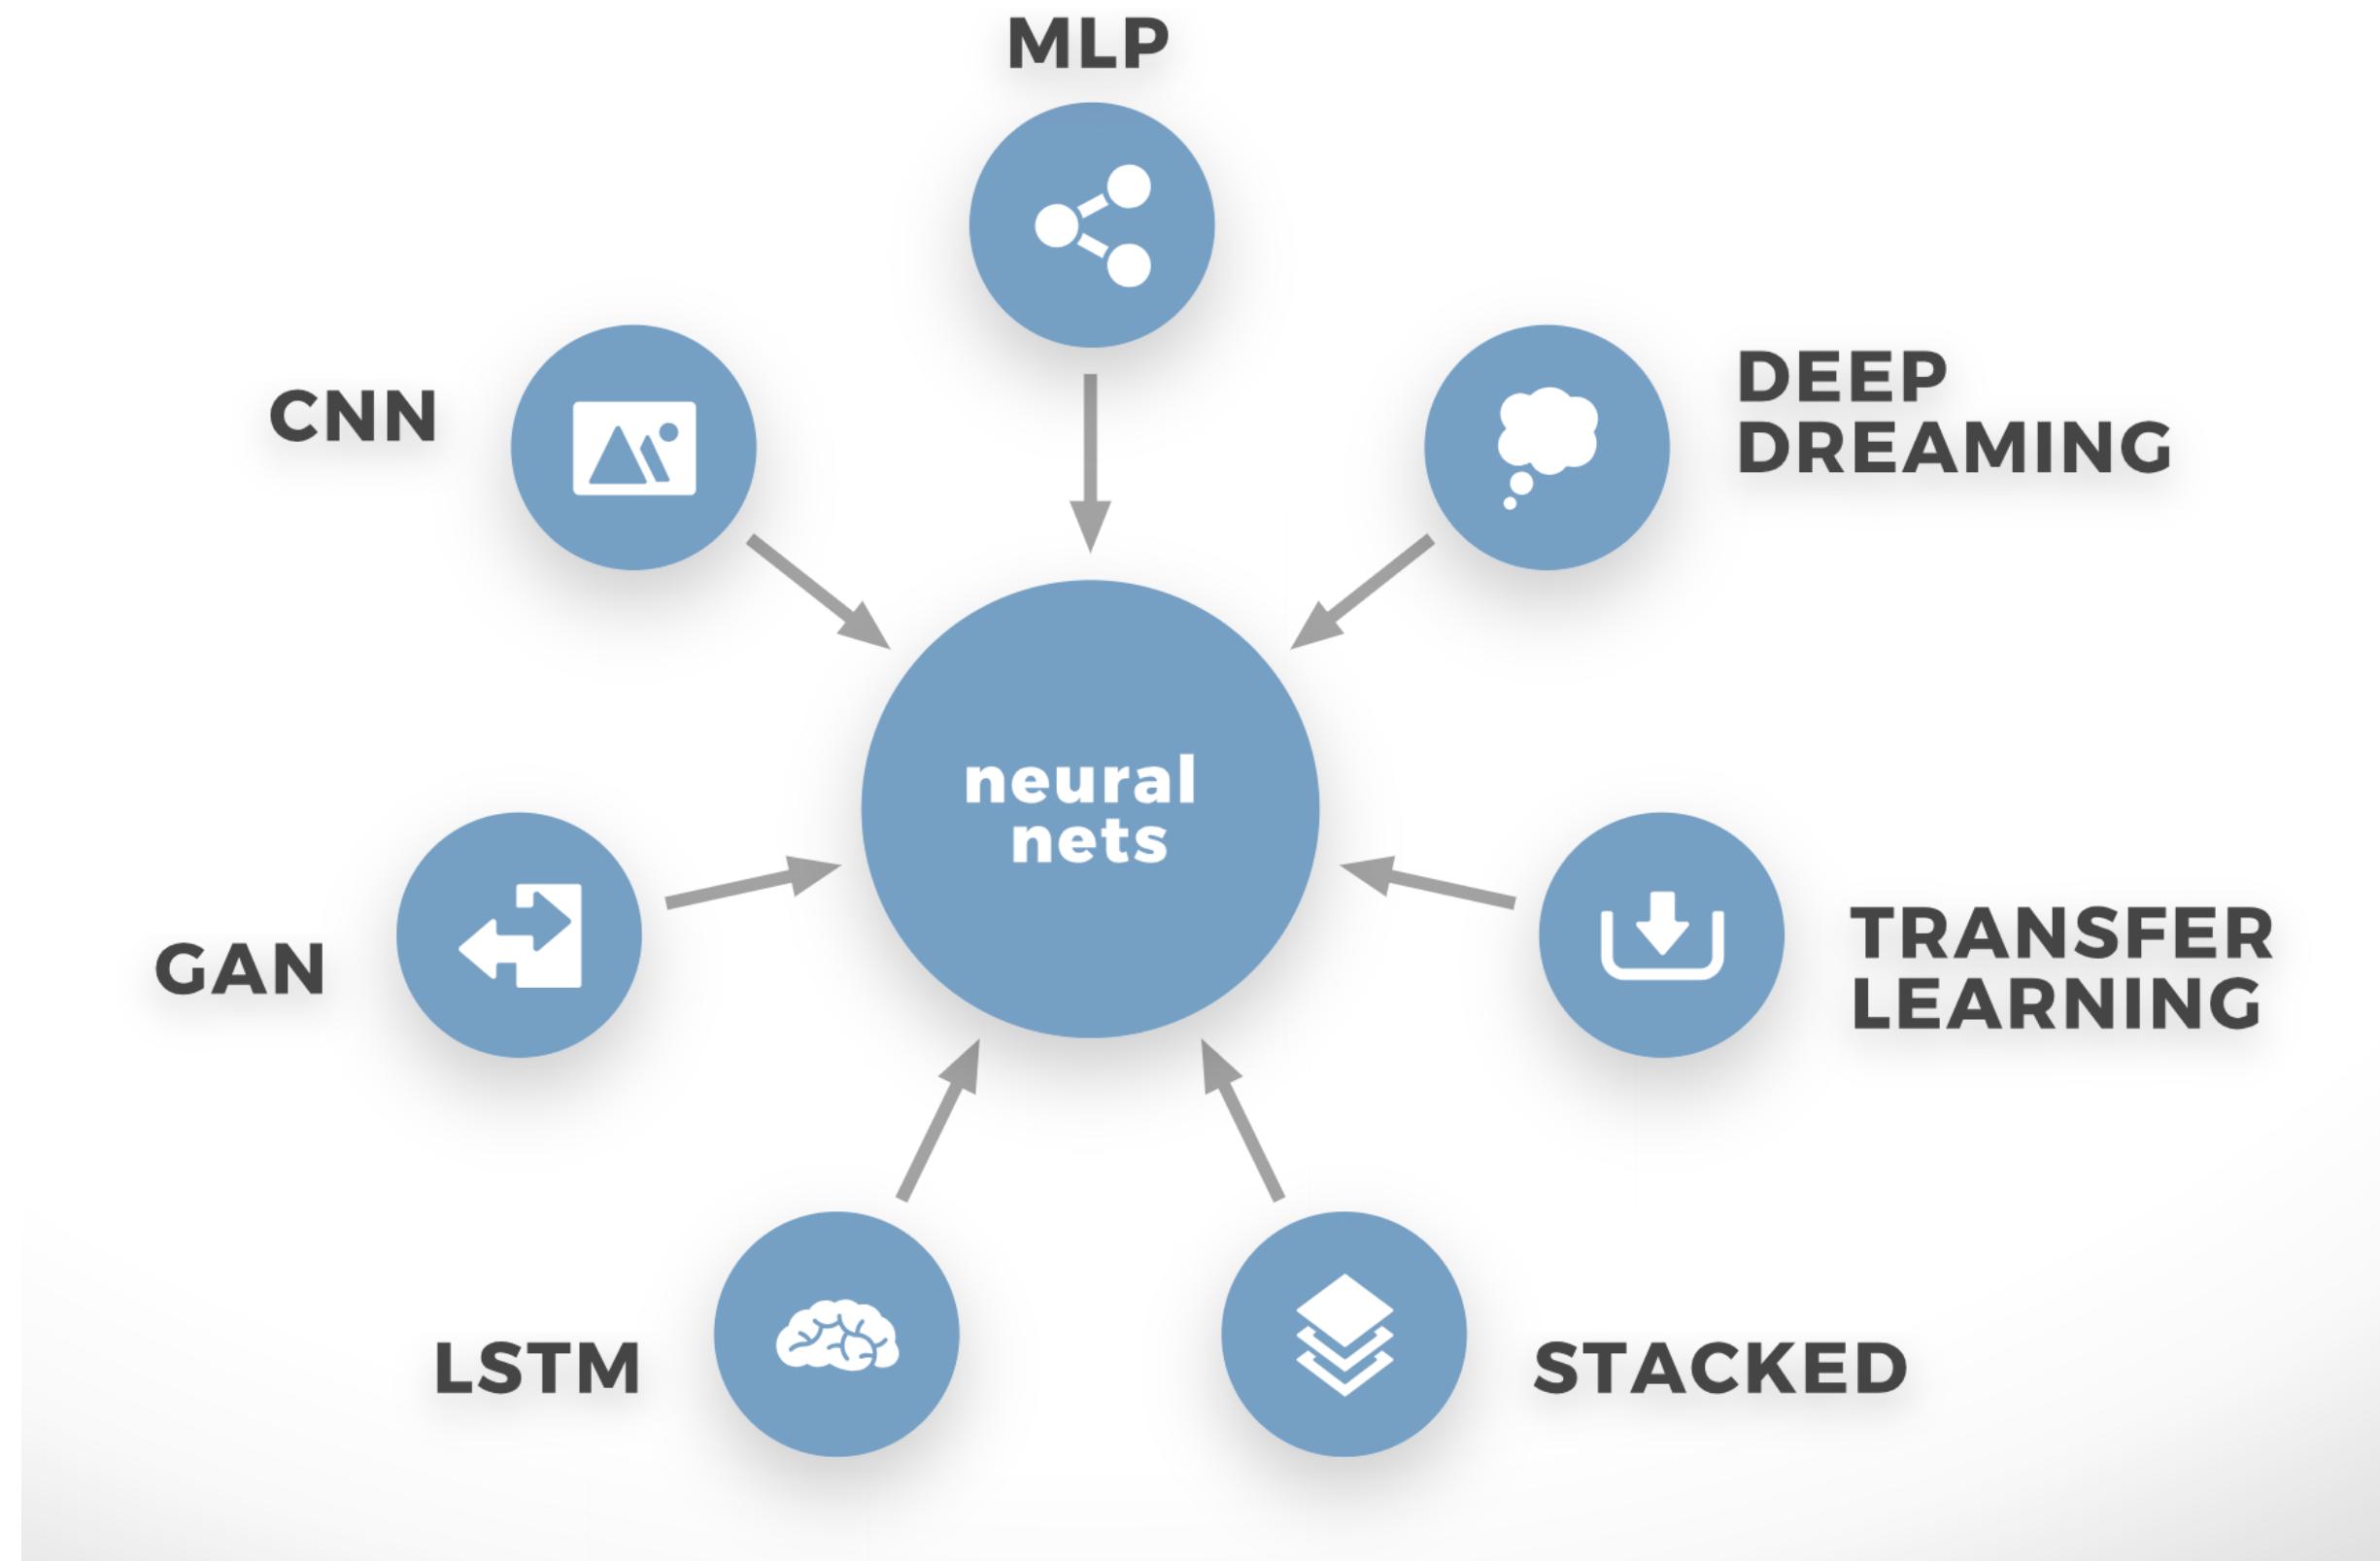
FUNCTIONAL API

more complex
suitable for complex models
can have multiple in- or outputs
layers can be non-sequential, e.g. LSTM

Keras layers



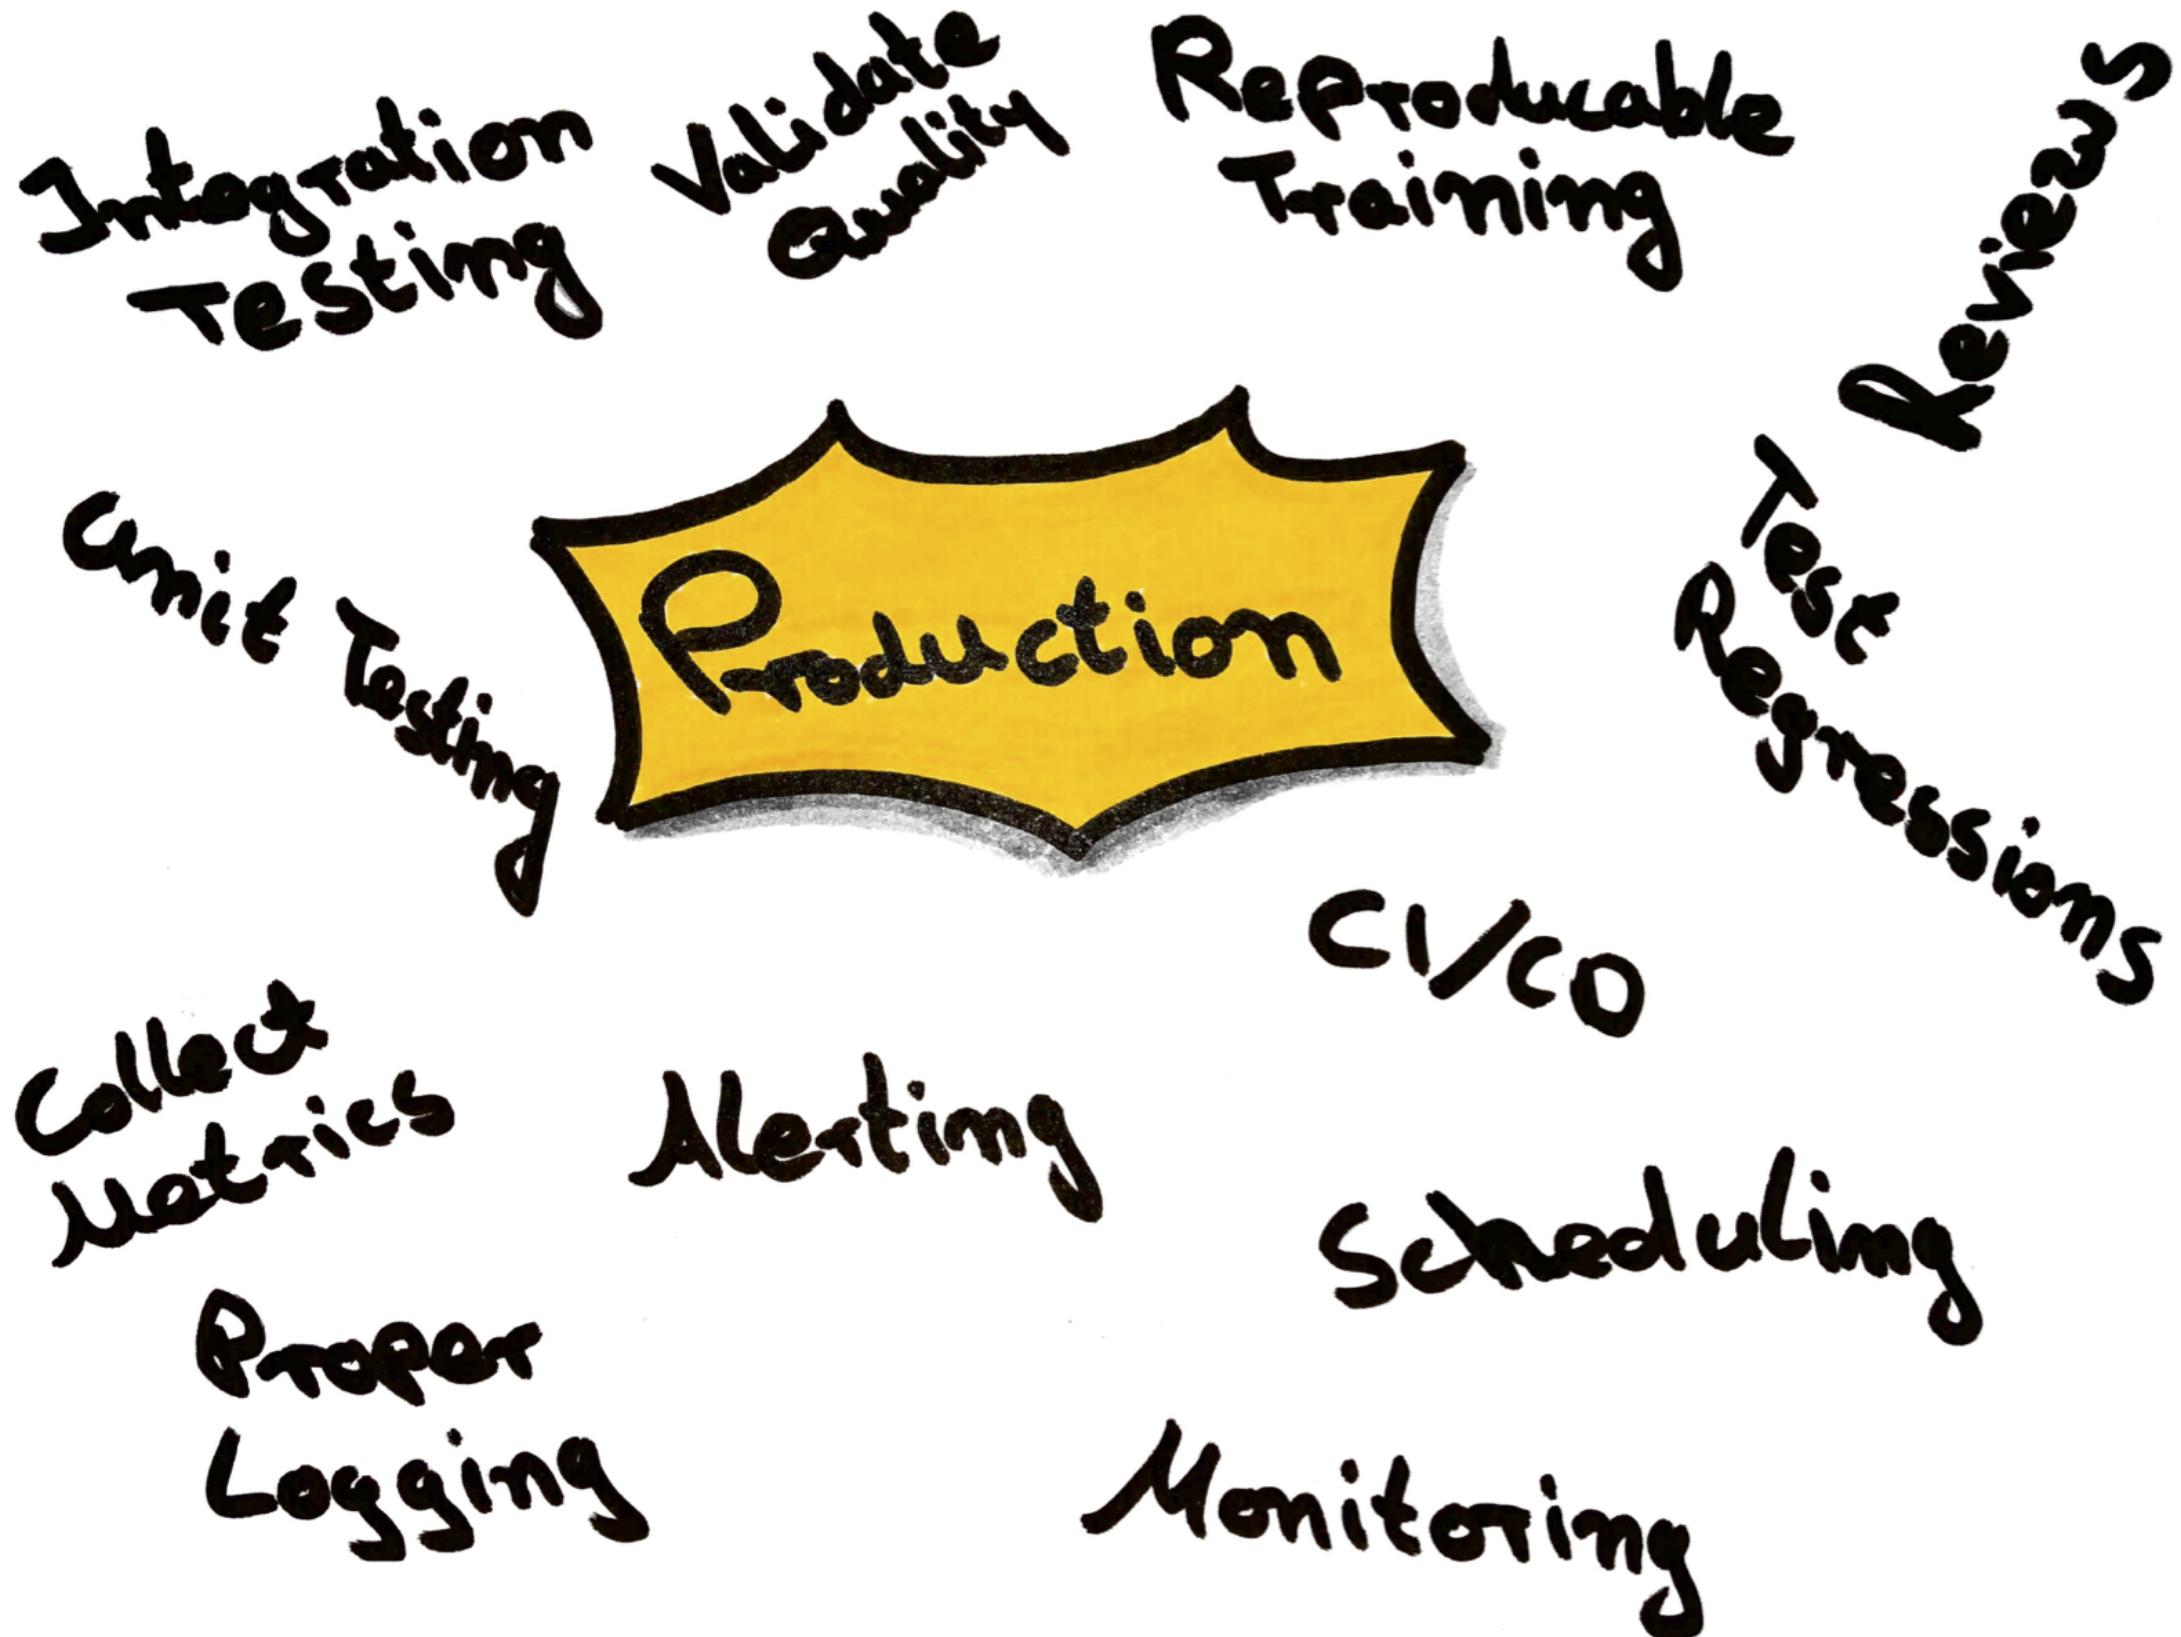
Endless possibilities



It works, now DEPLOY it!



Production ready



A bit about Luigi

Luigi helps to stitch long running tasks together into pipelines

It contains a wide toolbox of task templates (e.g. Hive, Pig, Spark, Python)

How to compose workflows?

A workflow consists of Targets, Tasks and Parameters

Targets correspond to a file or a database entry or some other kind of checkpoint

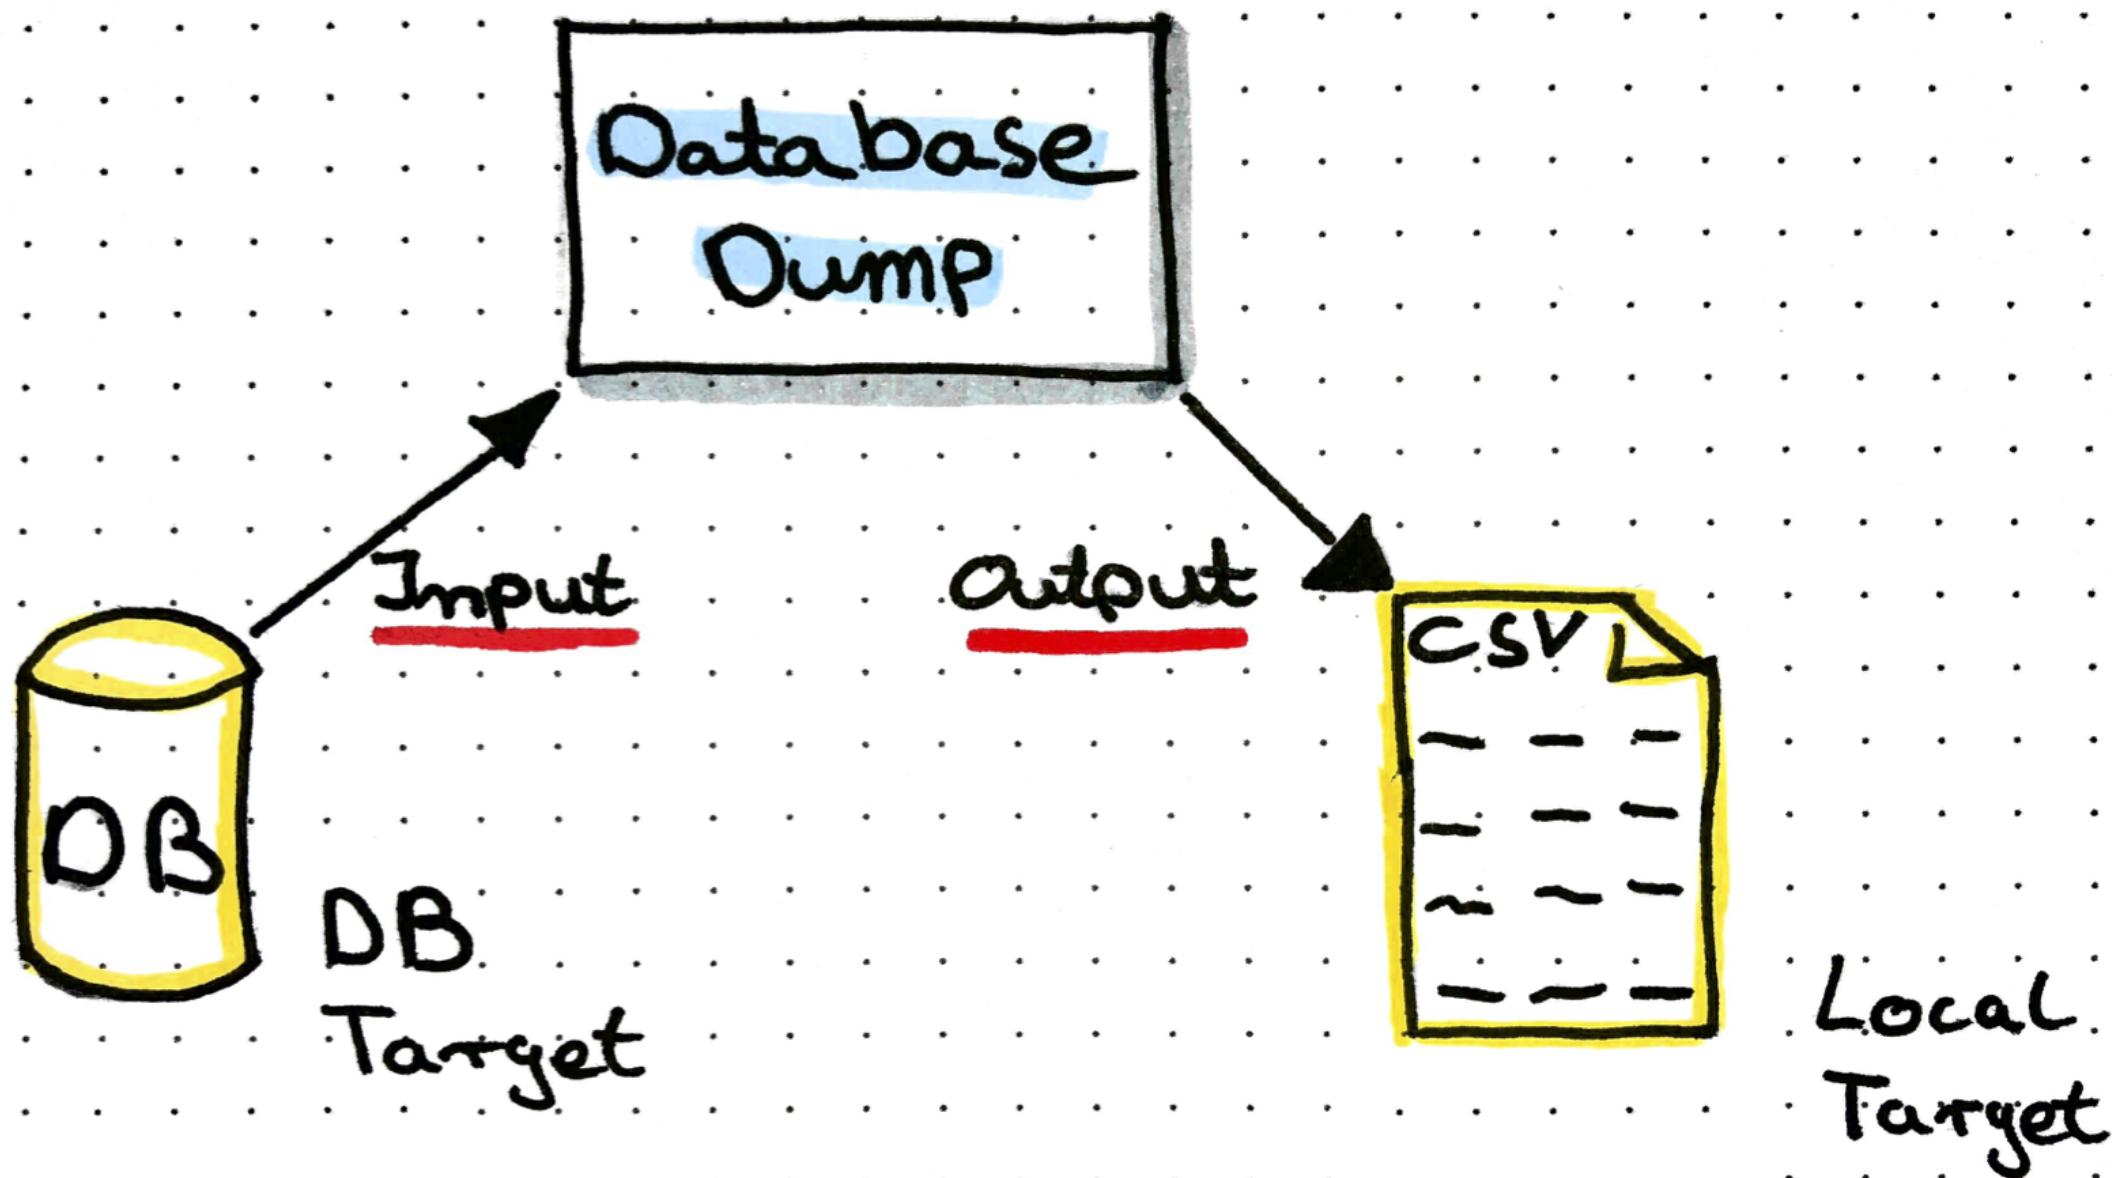
Tasks consume Targets of other tasks, run a computation, then output a target

Parameters take care of task parameterization

Targets

- Files on disk or database entries
- Checkpoints that prevent tasks from multiple executions
- A lot of implementations already exist in the Luigi framework
- LocalTarget (File), RemoteTarget (SSH), HDFSTarget, MySQLTarget, ...

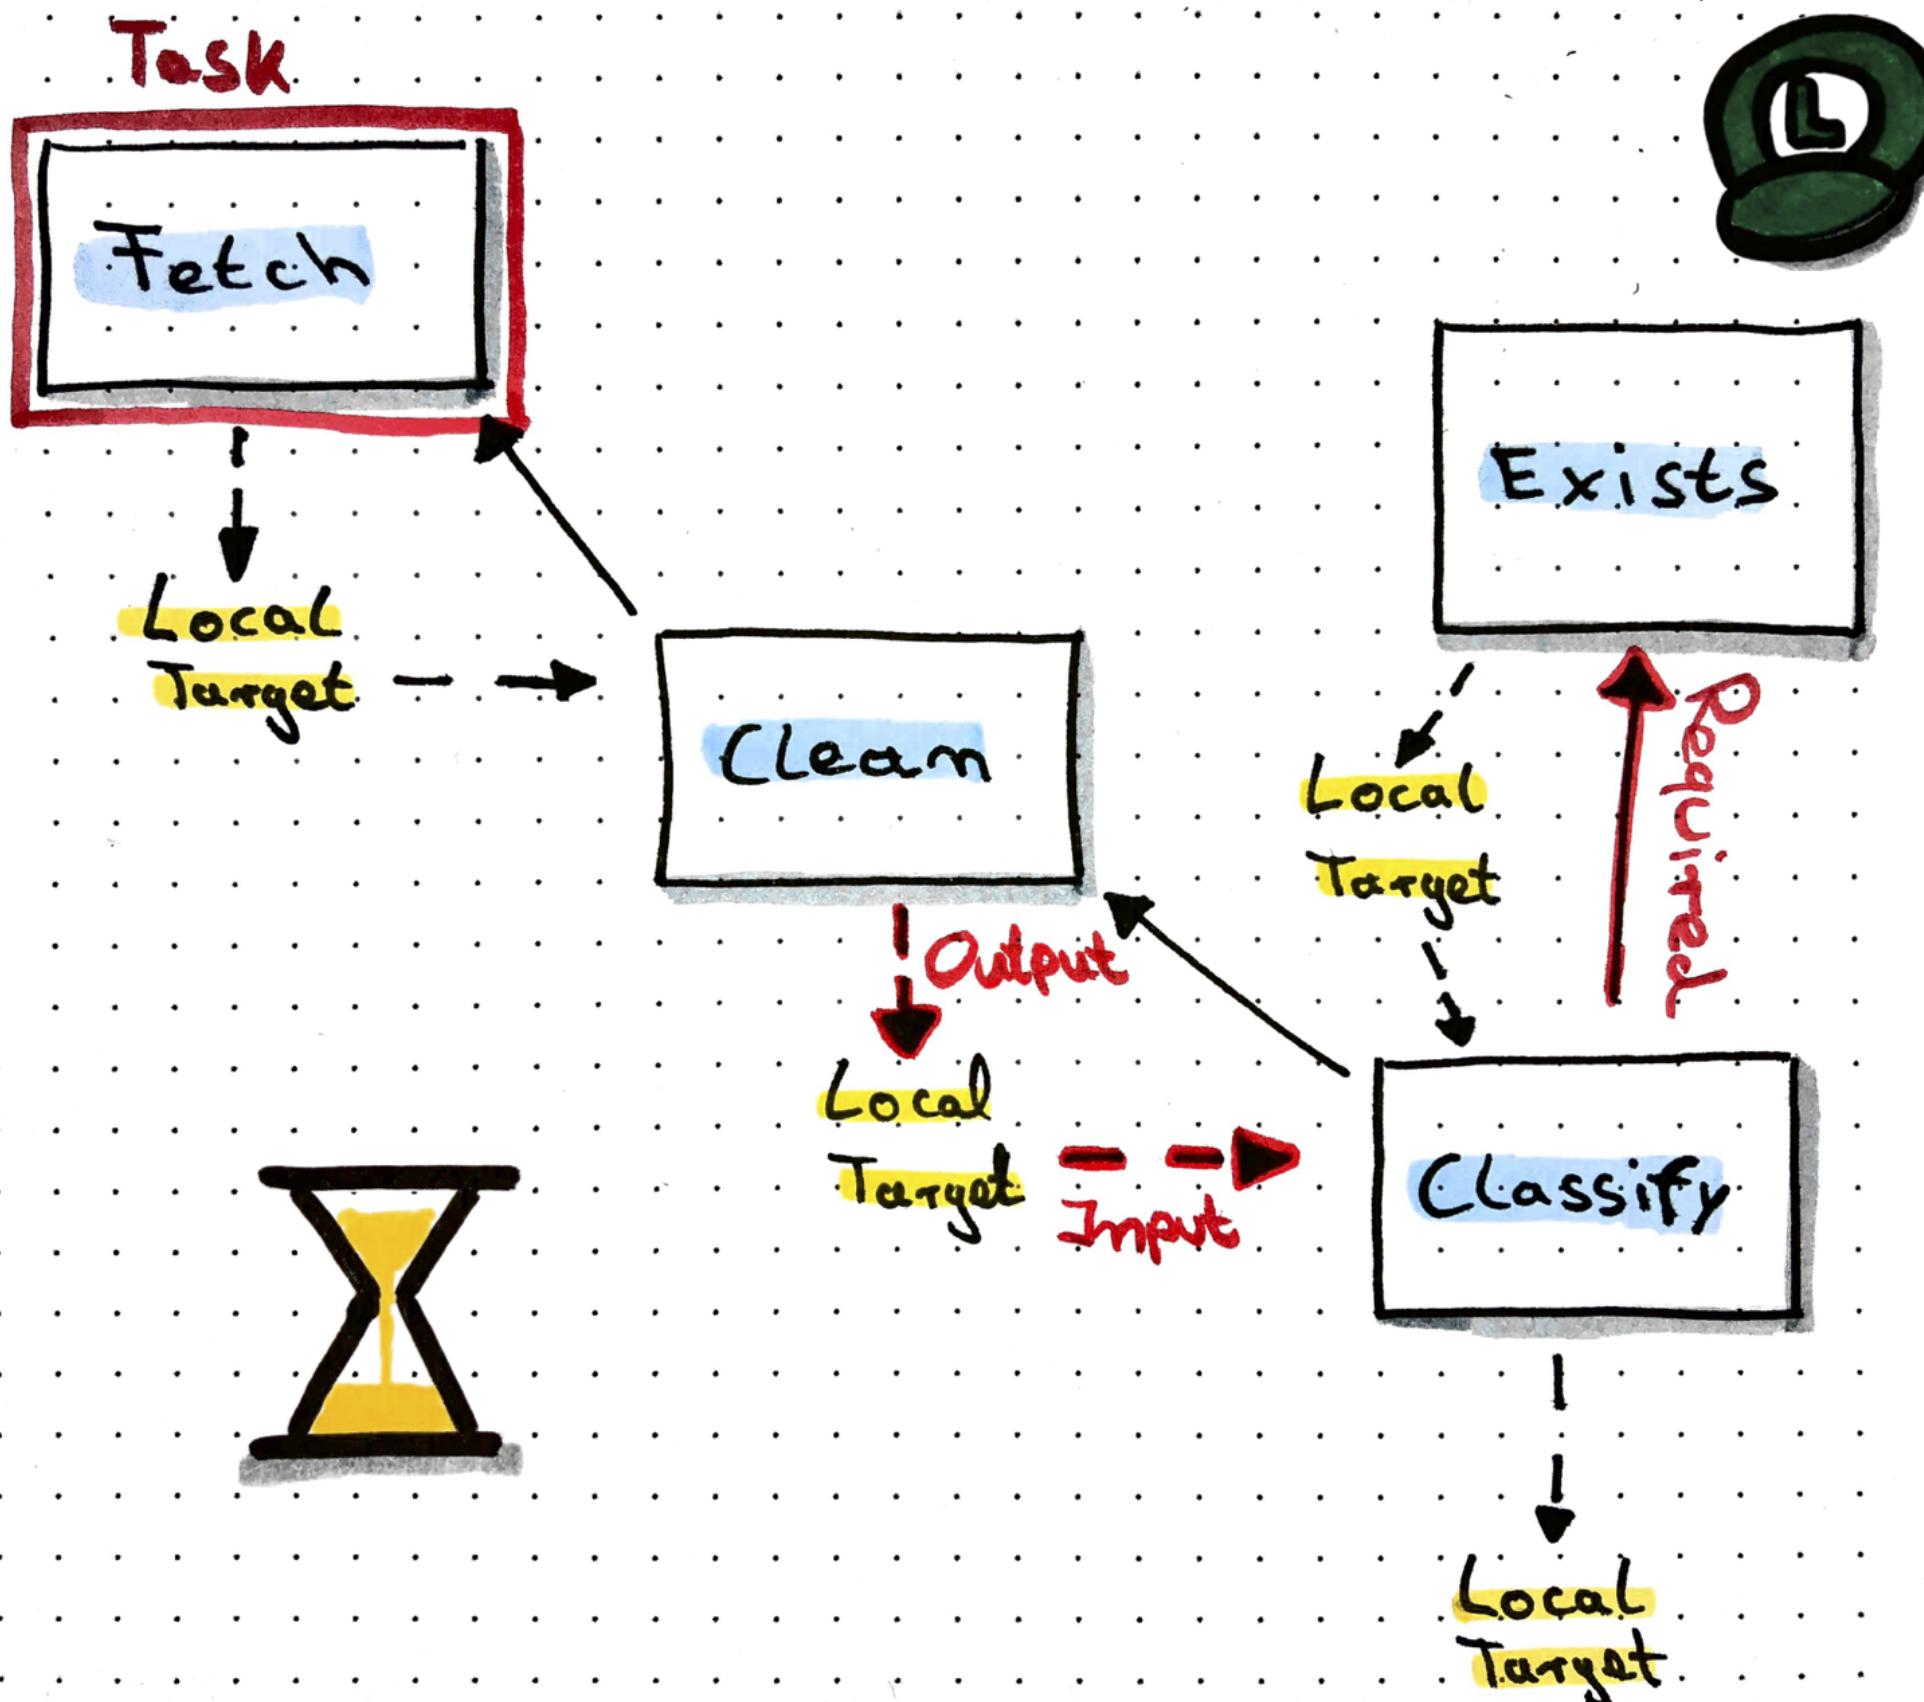
Targets



Tasks

- Implement the actual processing
- Consume targets, process data and save results in new target
- Respect dependencies to other tasks
- Implemented via Python-Classes

Tasks

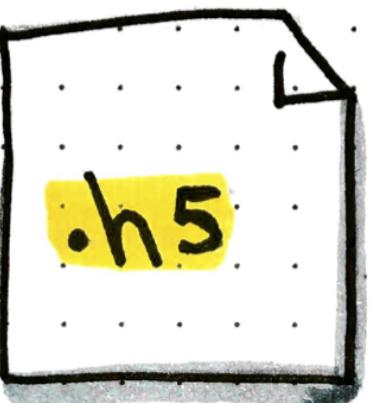


Parameters

- like constructur parameters
- Luigi takes care of parameter validation
- Again, a lot of implementations already exist in the Luigi framework
- IntParameter, BoolParameter, DateParameter, etc...

Parameters

".../{name}/{version}/model.h5"

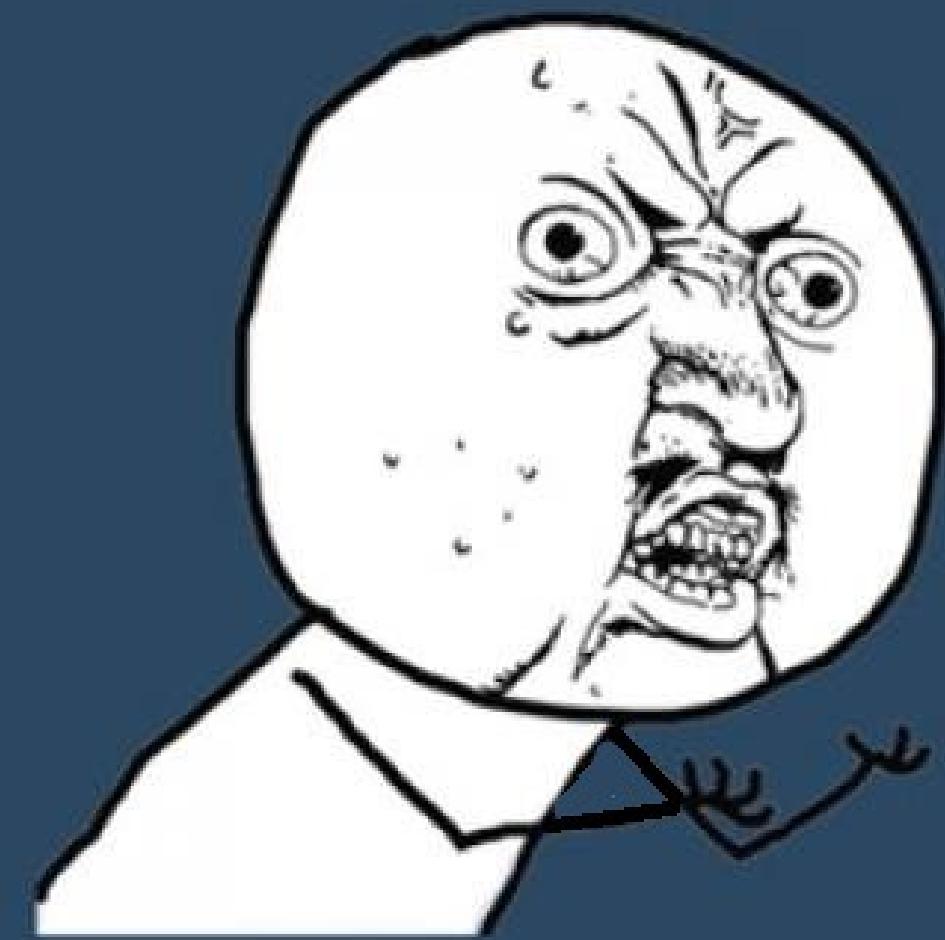


filepath



name = "cnn"
version = 1

WHY YOU NO

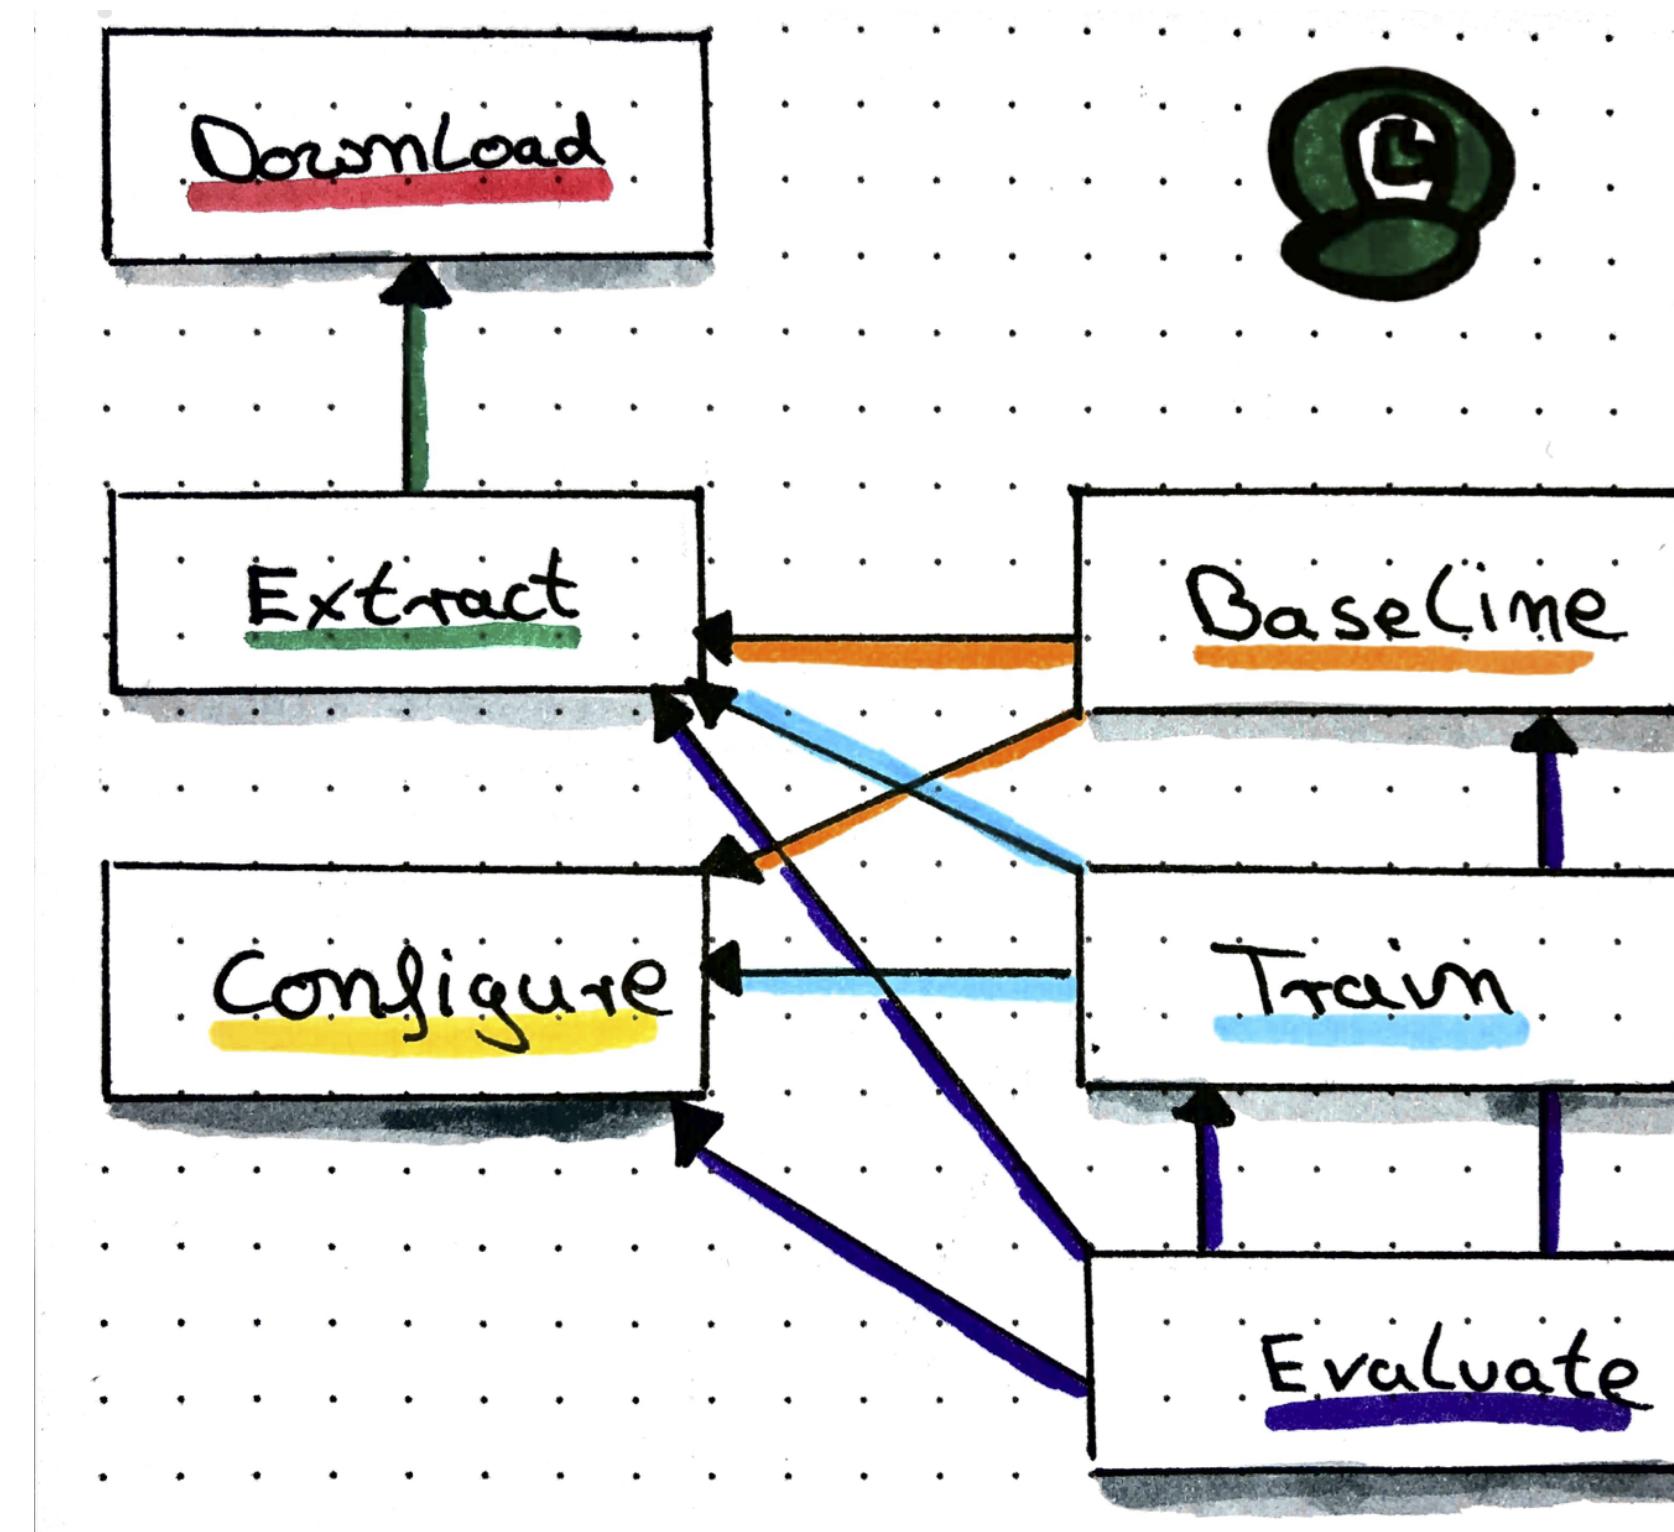


DO YOURSELF

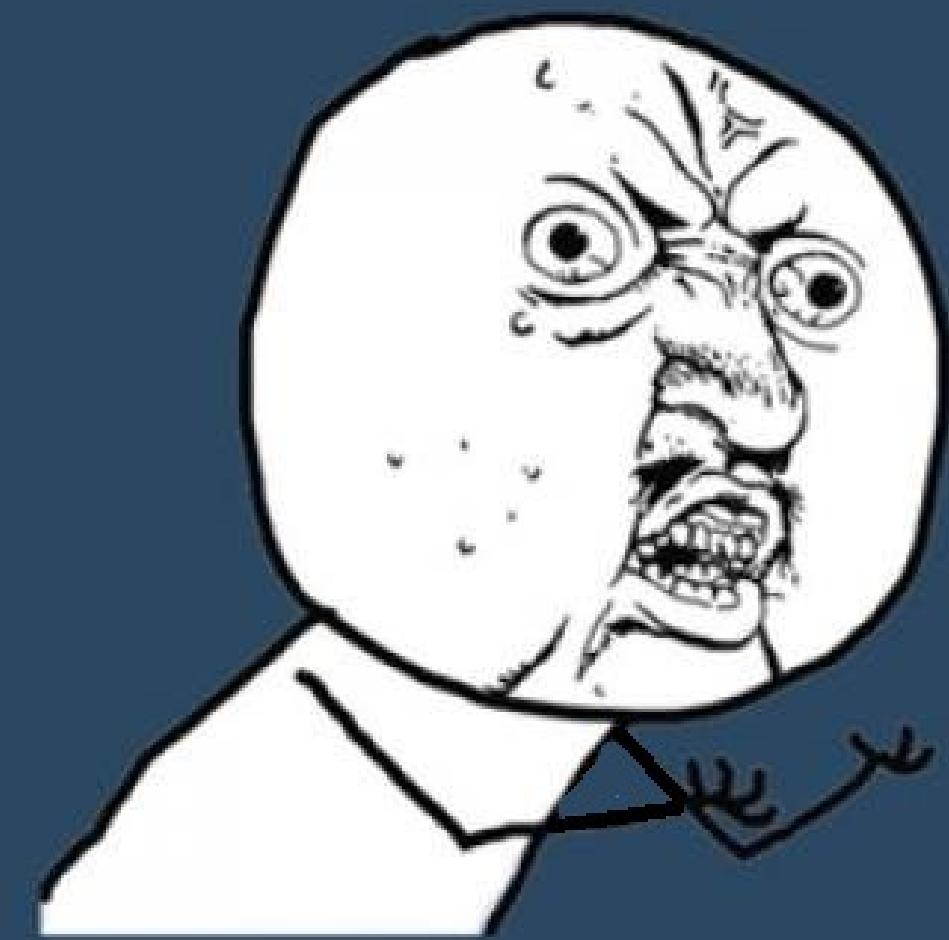
How would a production-ready Workflow look like?

- Download the dataset
- Extract the data
- Create a preprocessing configuration
- Run the baseline validation
- Train the model
- Evaluate the model

Workflow

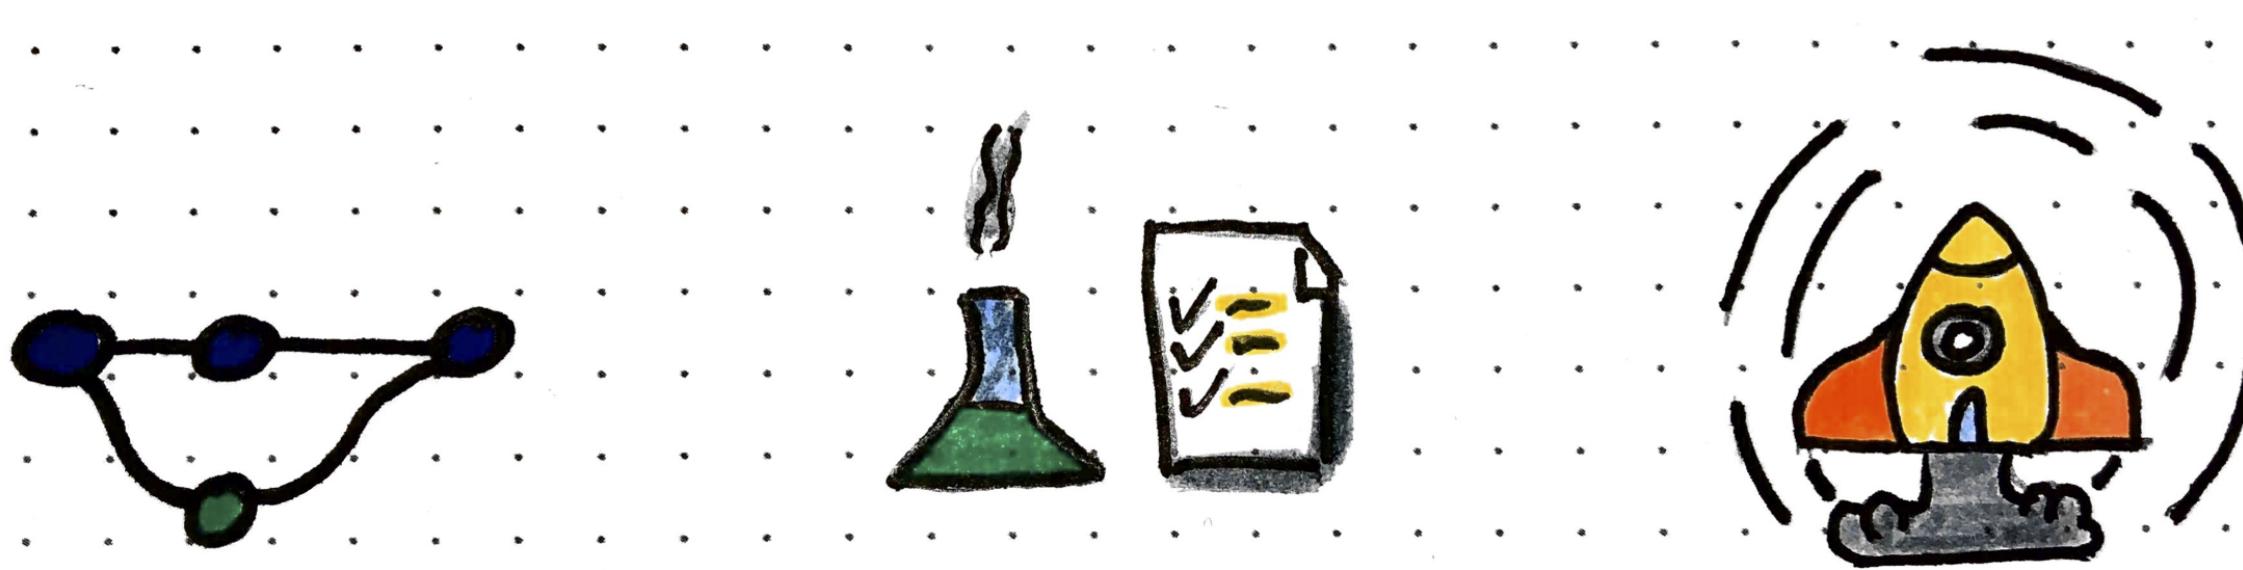


WHY YOU NO

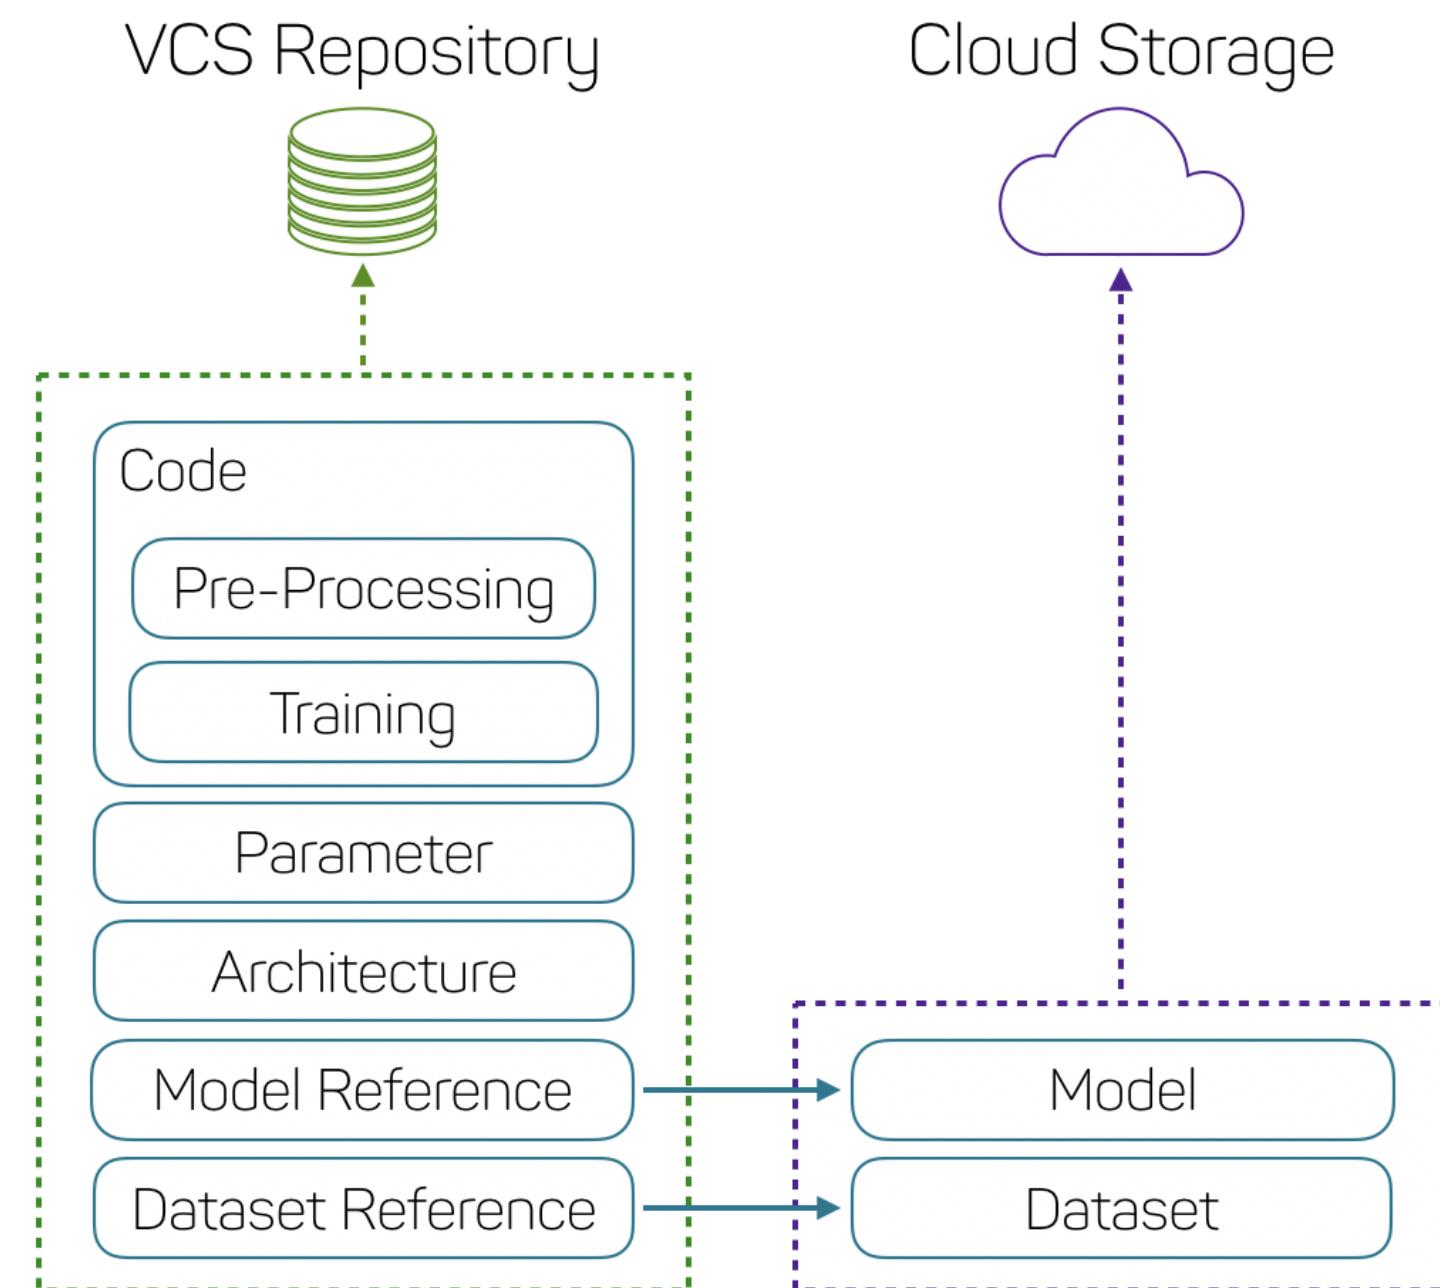


DO YOURSELF

DVC - Data Version Control



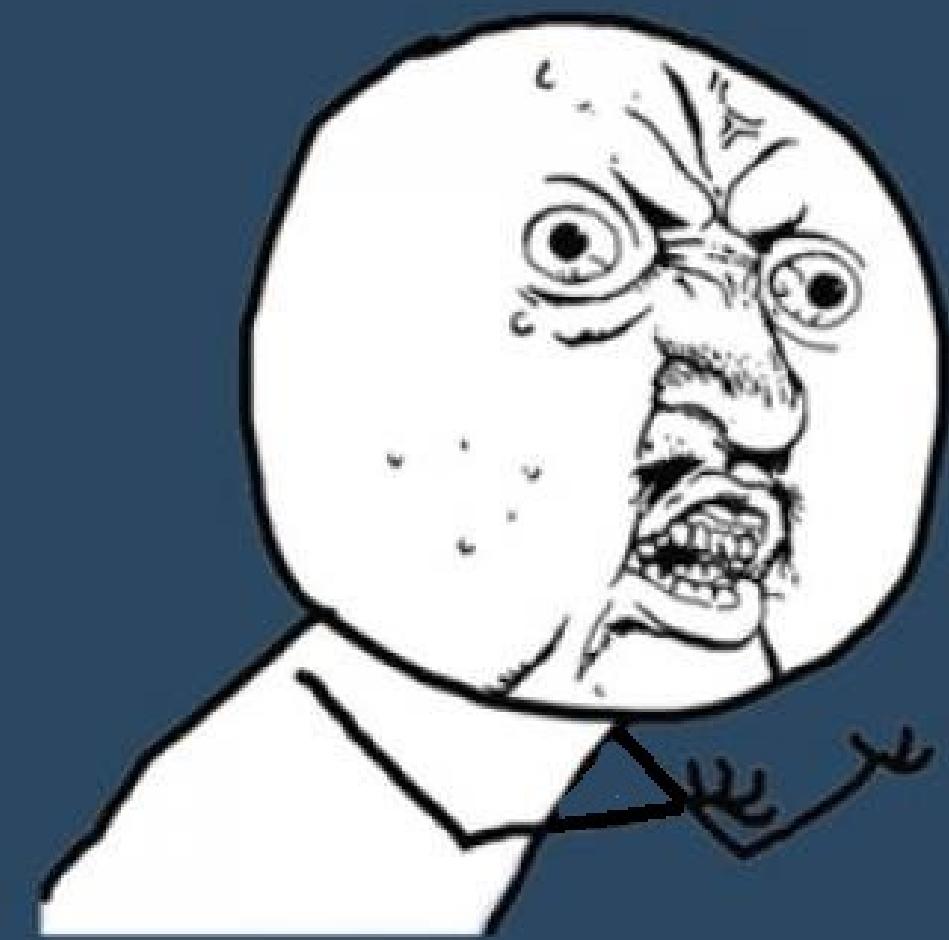
Artifacts and Storage



Using DVC

- DVC feels like Git
- Basic commands like: dvc init, dvc checkout, dvc add, dvc remove ...
- Lightweight pipelines: dvc run, dependencies wiht -d, outputs with -o, -O
- Management commands: dvc metrics, dvc pipeline

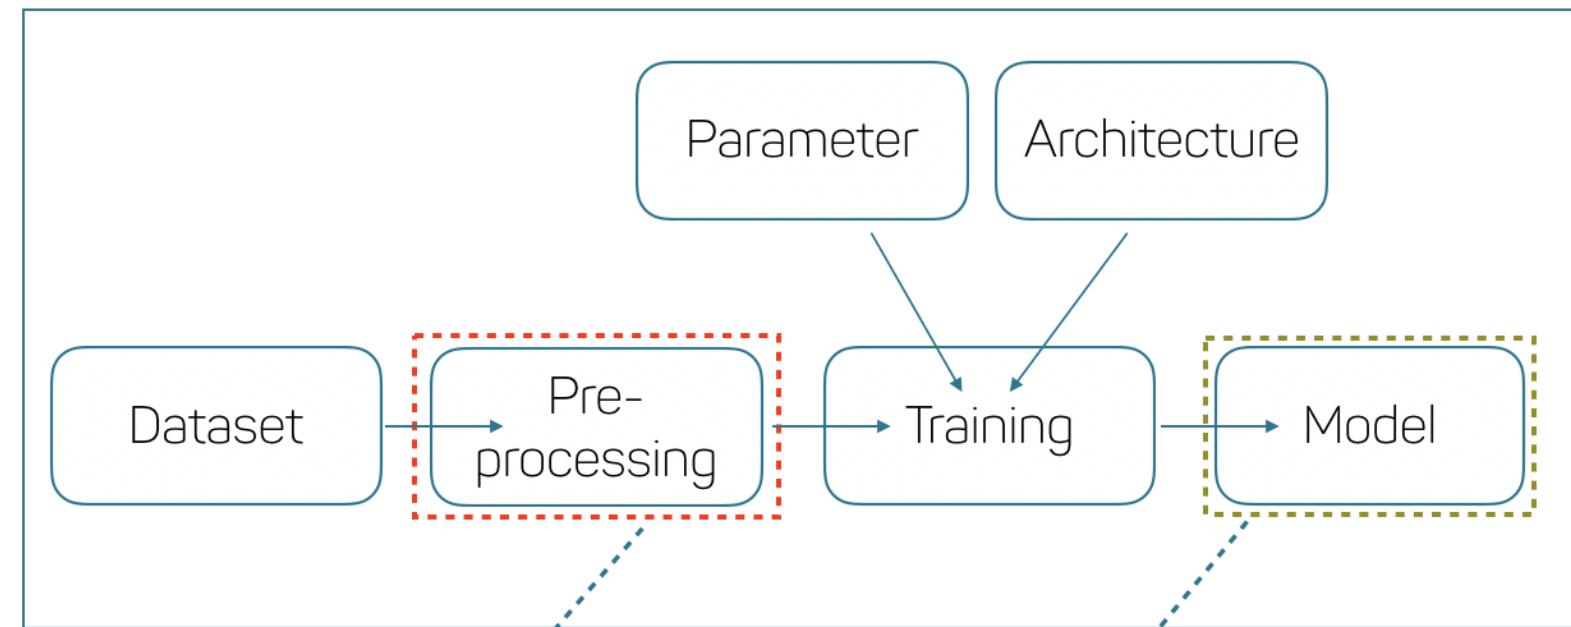
WHY YOU NO



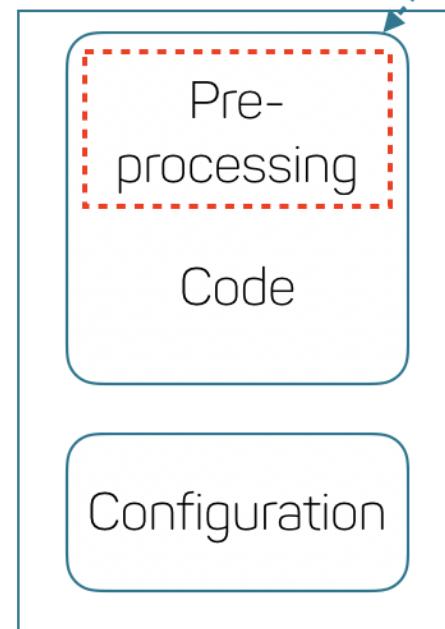
DO YOURSELF

Data Engineering - Development vs Experiments

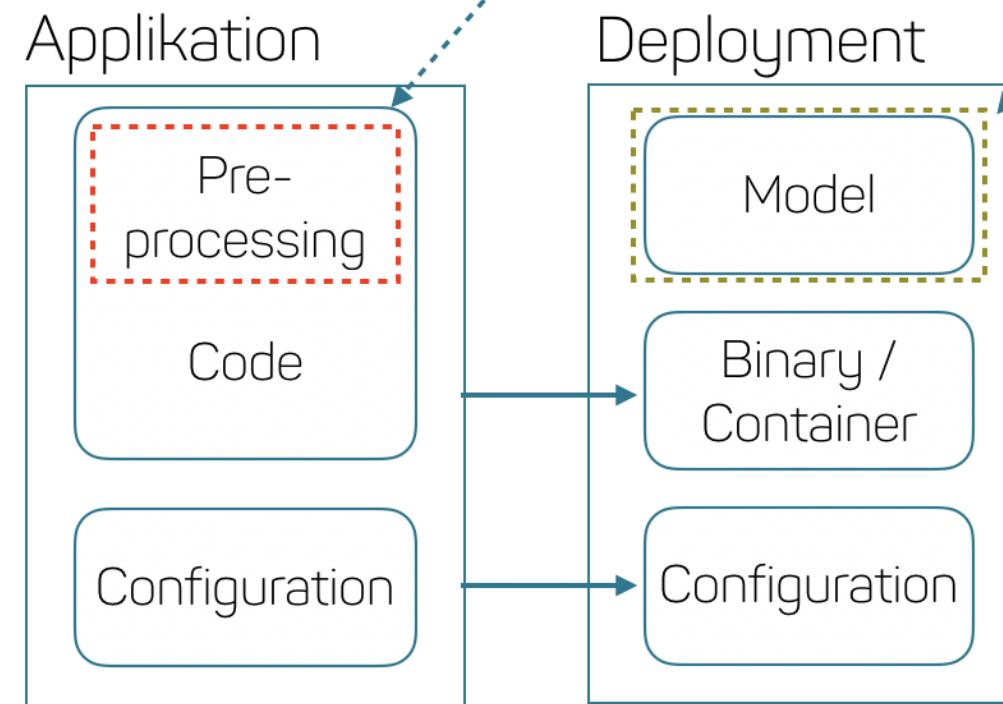
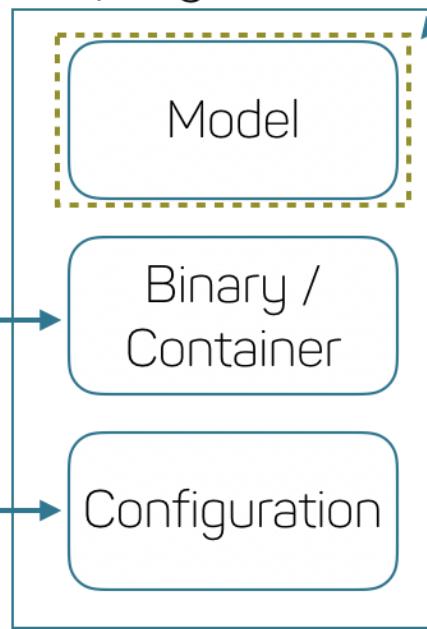
Data-Science Experiment



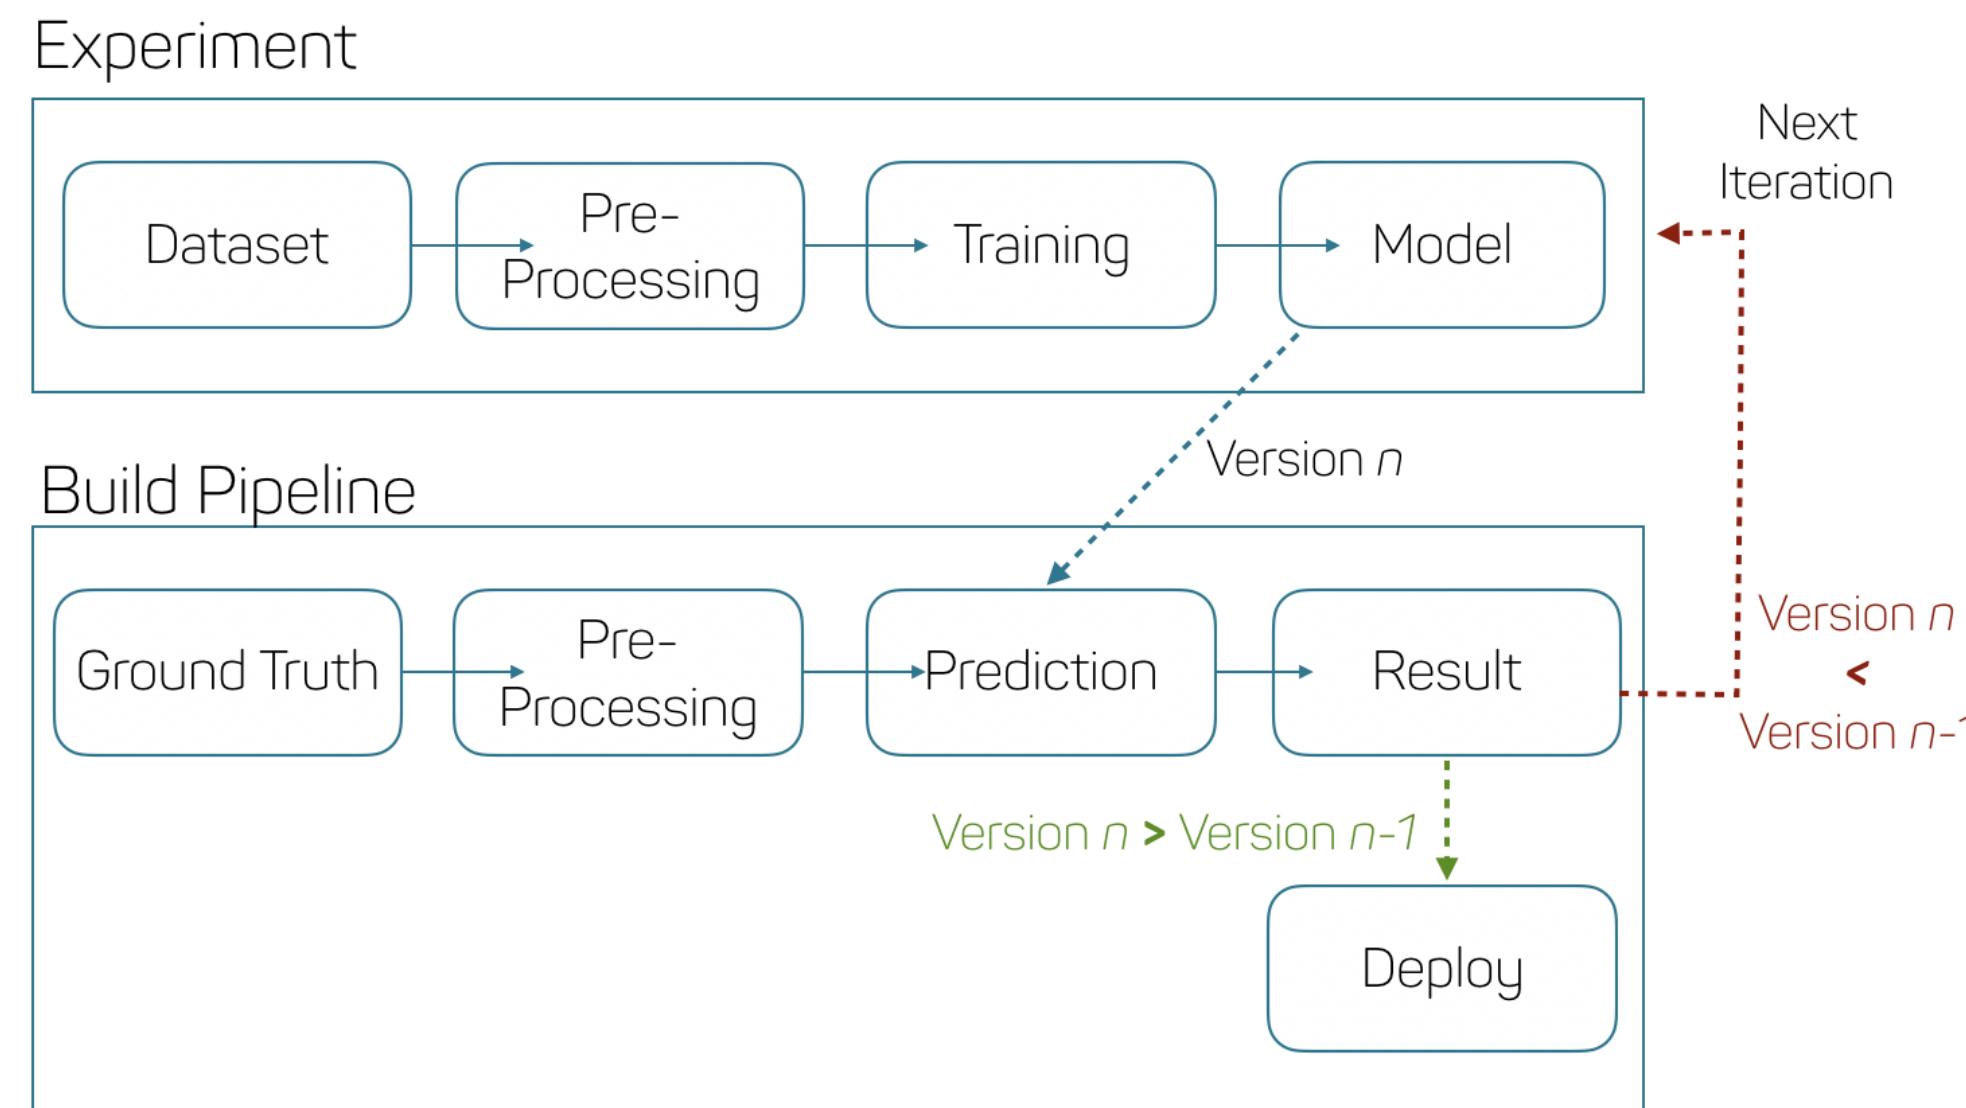
Applikation



Deployment



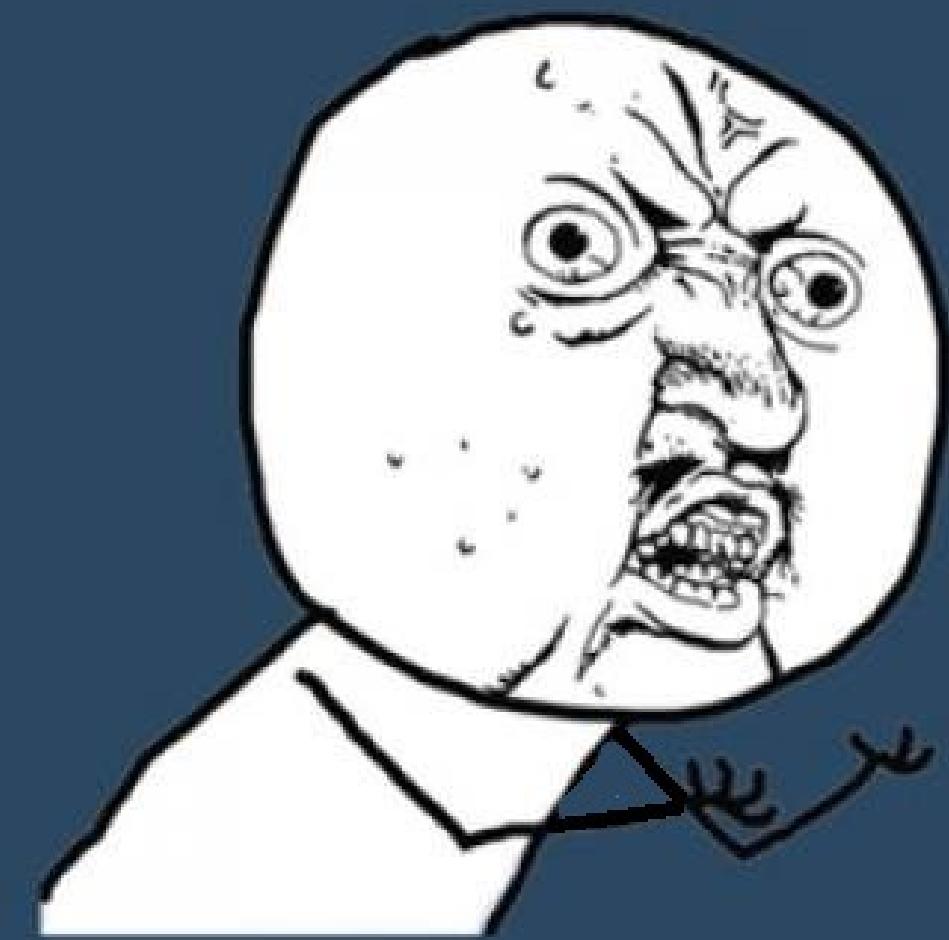
CI/CD for Experiments



Jenkins

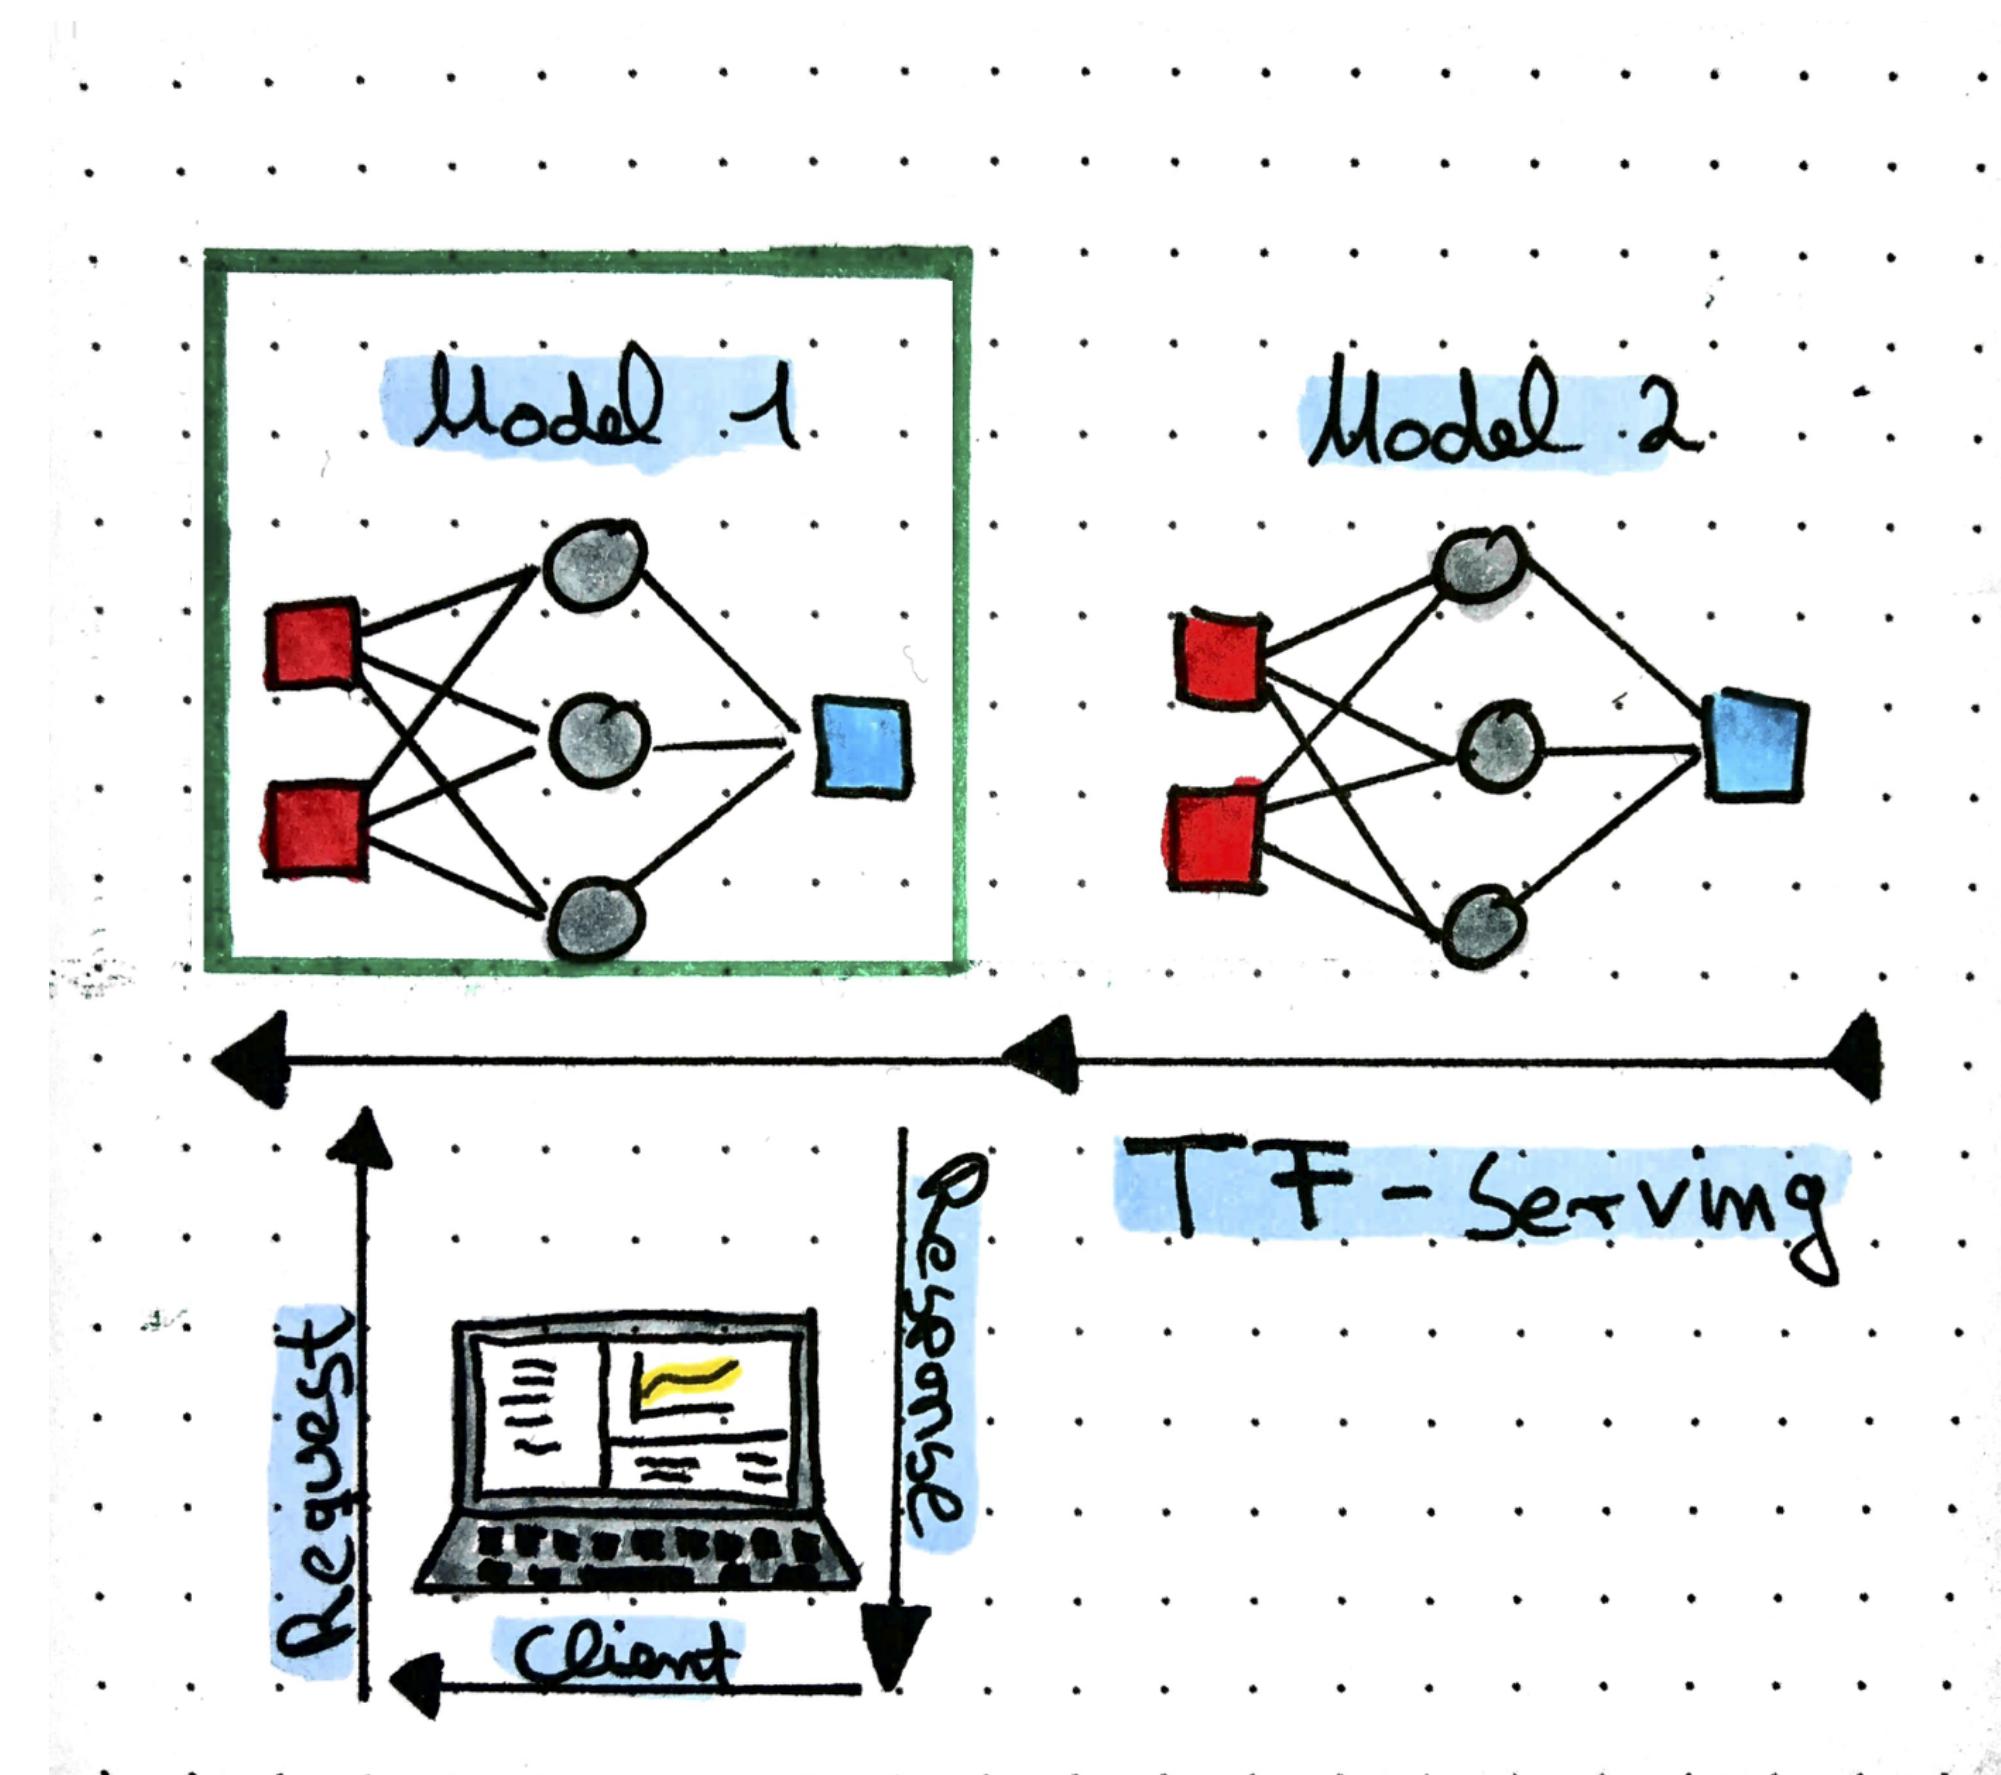
- Jenkins is an automation server
- Used for CI and CD
- Helps us to automate builds and deployments

WHY YOU NO



DO YOURSELF

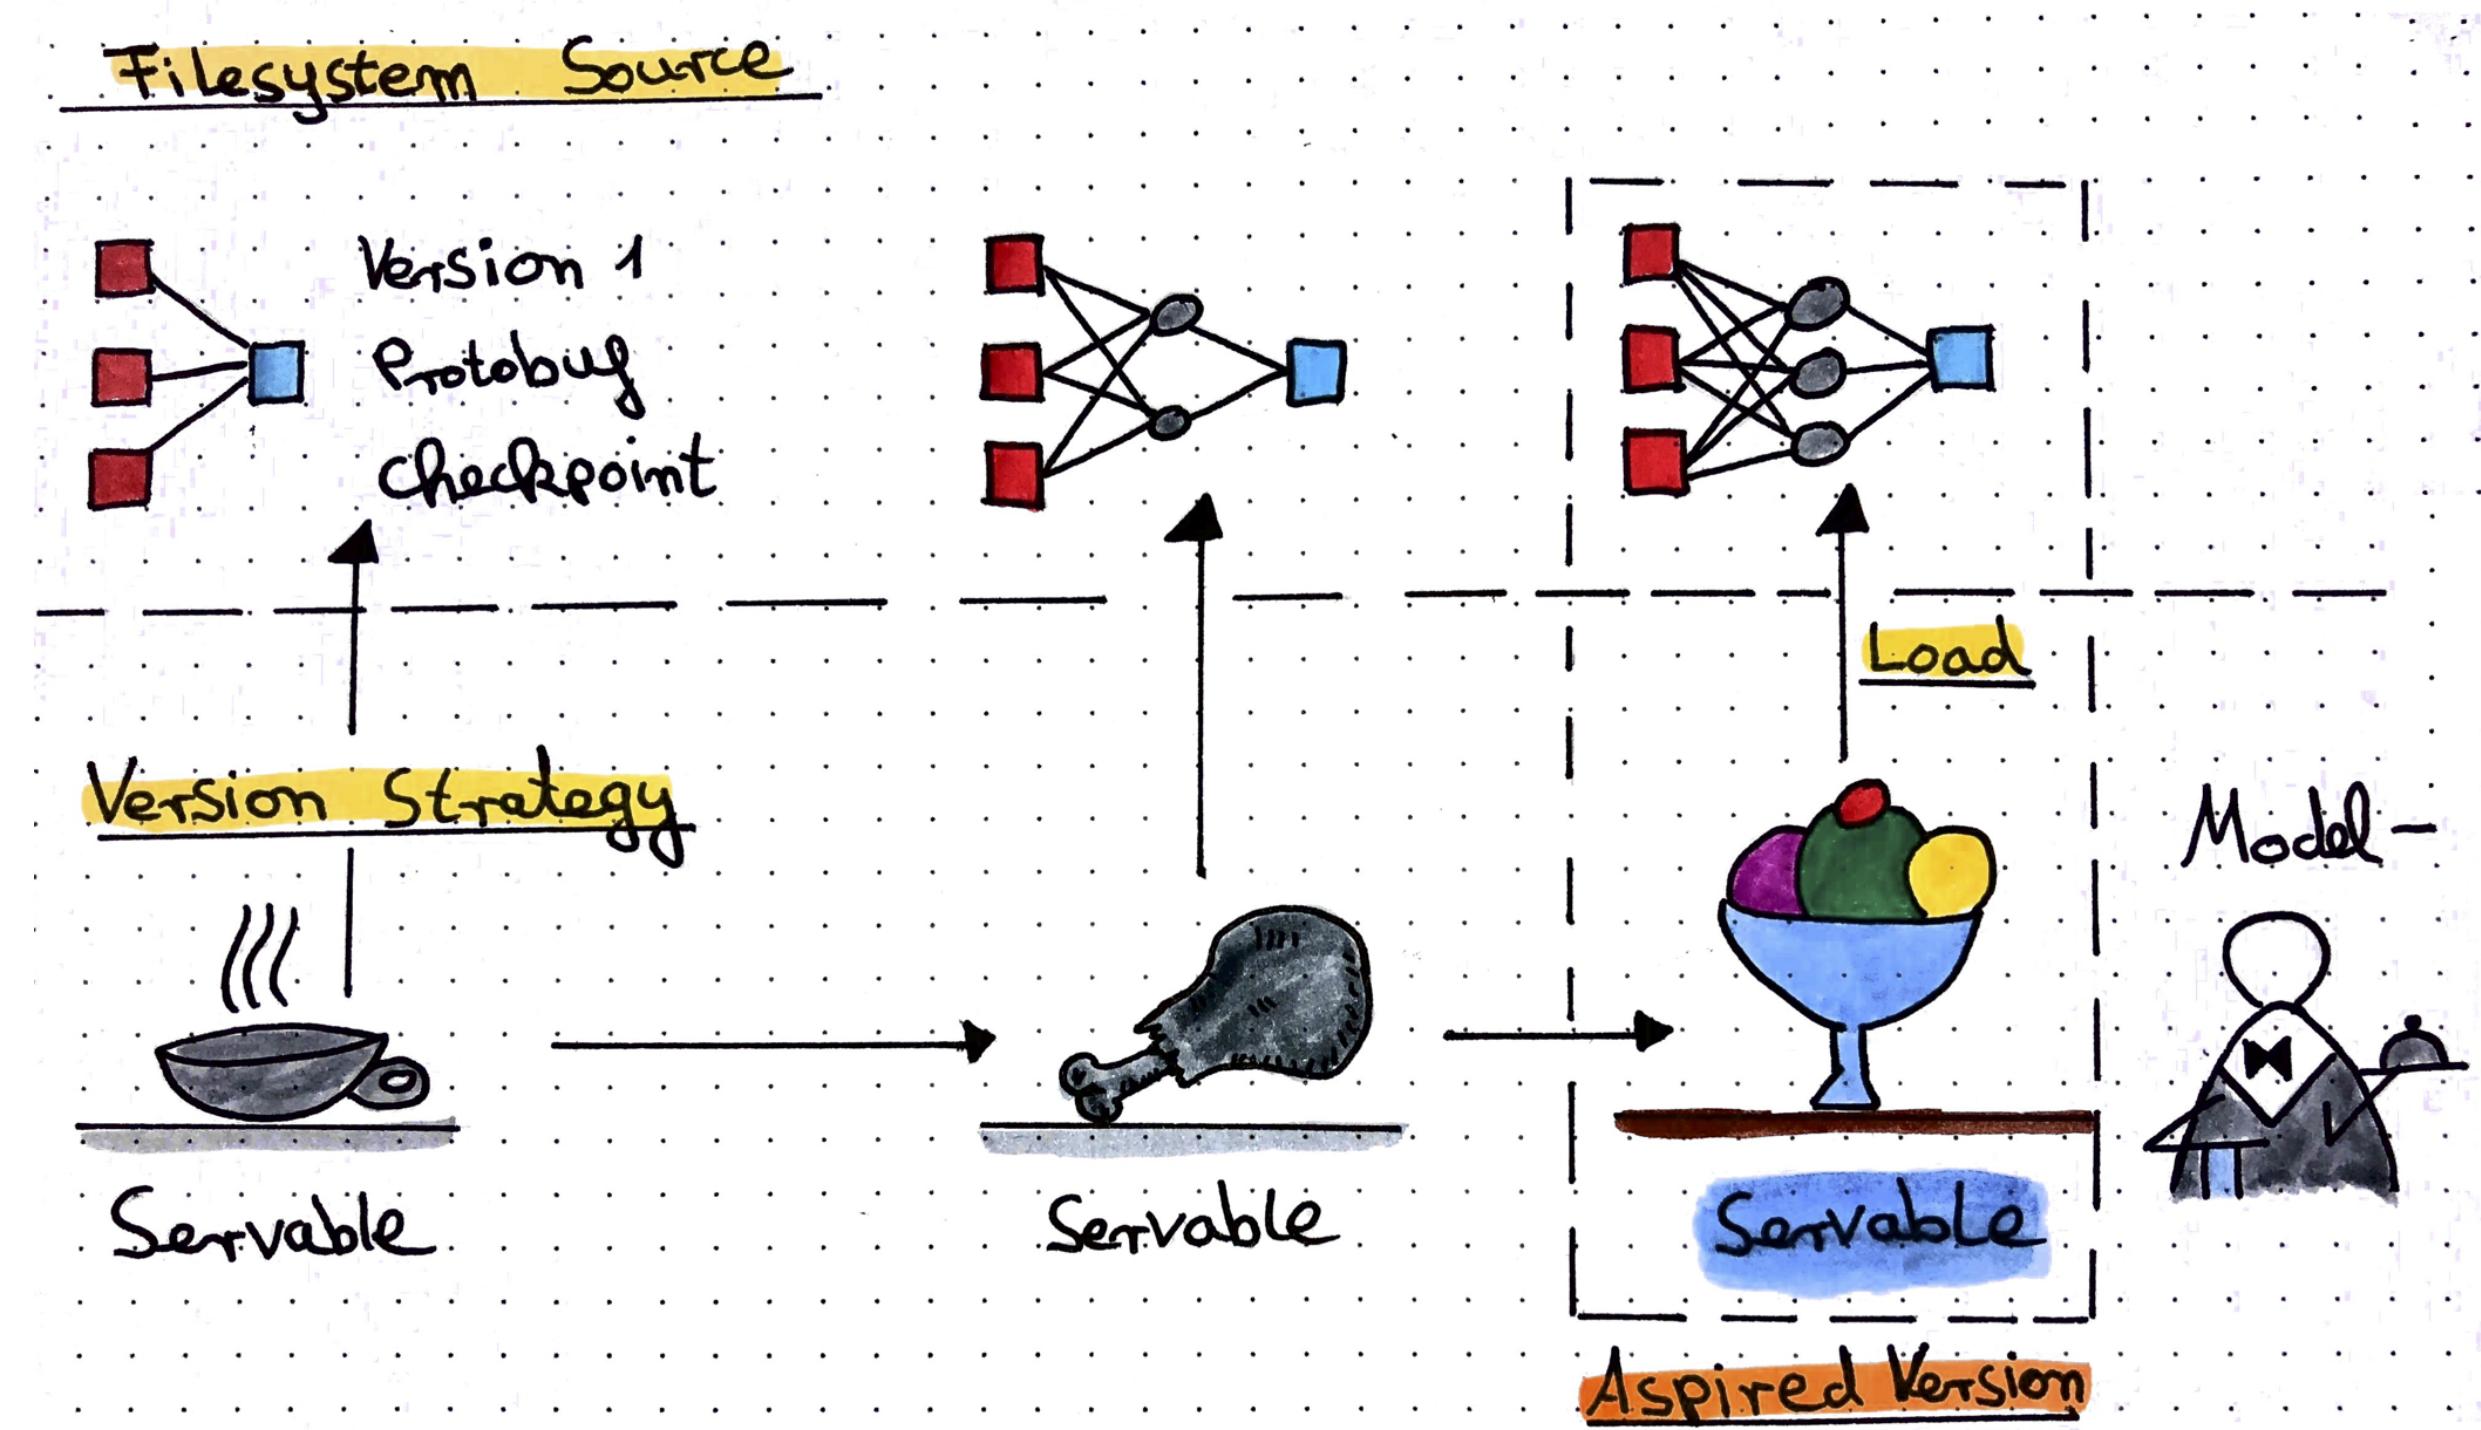
TensorFlow Serving



TensorFlow Serving

- ModelServer as Server-runtime for ML models
- Can natively handle TensorFlow graphs
- Highly flexible
- Designed for production use

TensorFlow Serving in Detail



TensorFlow Serving Components

- Protobuf file for graph, checkpoint files for weights (if not frozen)
- Version strategy defines how many models are loaded into memory
- Clients request inferences via gRPC or REST (yes REST!)
- TF-Serving consists of pluggable components, namely sources, loaders and version strategies

WHY YOU NO



DO YOURSELF

Integrate pre-processing into model

- The pre-processing and the model are tightly coupled
- Datagenerator can be pickled to make training reproducible
- But how can the client be sure to do the "right" thing?

Possible ways:

- The client just "knows"
- Wrap service around model and expose an API
- Built the pre-processing right into the model

```
In [ ]: def pre_process(x):
    resized = tf.image.resize_images(x, size=[100, 100])
    max_color = tf.constant(255, dtype=tf.float32)
    rescaled = tf.divide(resized, max_color)
    return rescaled

additional = Lambda(pre_process, output_shape=(100, 100, 3), name="preprocessing")
```

```
In [ ]: prod_model = Sequential()
prod_model.add(InputLayer(input_shape=(None, None, 3)))
prod_model.add(additional)
for layer in model.layers:
    prod_model.add(layer)
```